

NeRF与3DGS在动态三维重建中的应用

姚遥

南京大学 NJU-3DV

动态三维重建



静态三维重建

2D Observations → 3D



动态三维重建

2D Observations → 4D

动态三维重建的关键挑战：

- 高效连续的四维空间表达方式
- 图像观测不足，重建求解欠定
- ...



动态三维重建

- 基于四维空间 (4D Representation) 的动态三维重建方法
- 基于形变场 (Deformation Field) 的动态三维重建方法
- 基于逐帧重建 (Per-frame Reconstruction) 的动态三维重建方法
- 基于模型先验 (Model Priors) 的动态三维重建方法

4D Representations

- **NeRF Related:** DyNeRF
- **Triplane Related:** K–Planes, HexPlane, Tensor4D
- **3DGS Related:** 4DGS, 4D–Rotor GS

Neural 3D Video Synthesis from Multi-view Video

CVPR 2022 (oral)

Tianye Li^{1,3,*} Mira Slavcheva^{1,*} Michael Zollhoefer¹

Simon Green¹ Christoph Lassner¹ Changil Kim² Tanner Schmidt¹

Steven Lovegrove¹ Michael Goesele¹ Richard Newcombe¹ Zhaoyang Lv¹

¹Reality Labs Research

²Meta

³University of Southern California



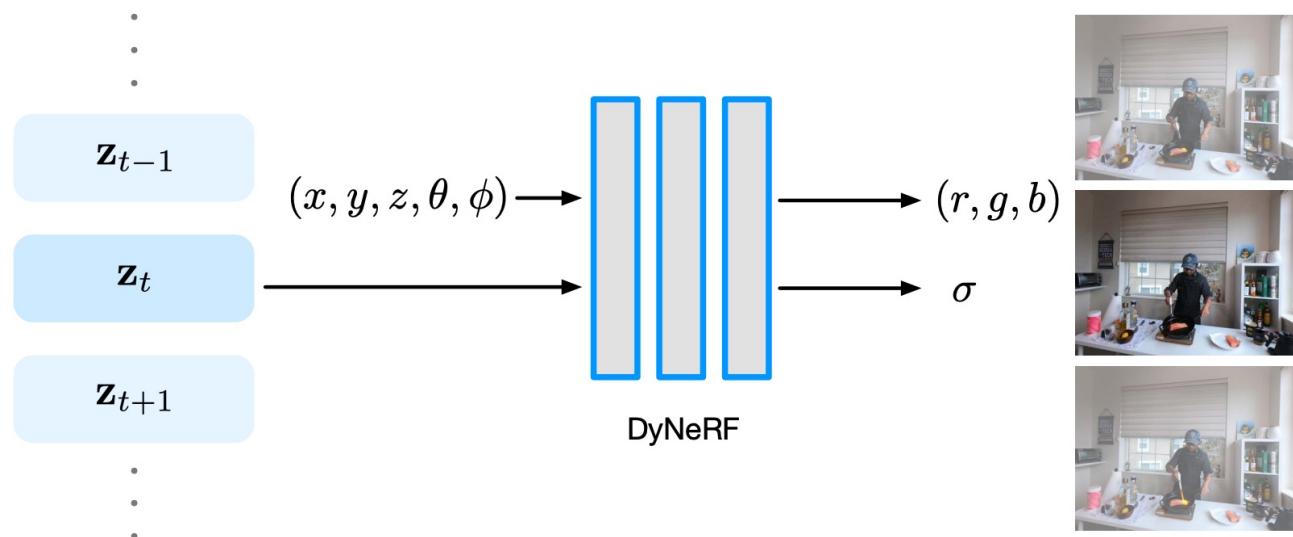


Table 1. Quantitative comparison of our proposed method to baselines of existing methods and radiance field baselines trained at 200K iterations on a 10-second sequence.

Method	PSNR \uparrow	MSE \downarrow	DSSIM \downarrow	LPIPS \downarrow	FLIP \downarrow
MVS	19.1213	0.01226	0.1116	0.2599	0.2542
NeuralVolumes	22.7975	0.00525	0.0618	0.2951	0.2049
LLFF	23.2388	0.00475	0.0762	0.2346	0.1867
NeRF-T	28.4487	0.00144	0.0228	0.1000	0.1415
DyNeRF †	28.4994	0.00143	0.0231	0.0985	0.1455
DyNeRF	29.5808	0.00110	0.0197	0.0832	0.1347

K–Planes: Explicit Radiance Fields in Space, Time, and Appearance

Sara Fridovich-Keil*

UC Berkeley

sfk@berkeley.edu

Giacomo Meanti*
Istituto Italiano di Tecnologia

giacomo.meanti@iit.it

Frederik Rahbæk Warburg
Technical University of Denmark

frwa@dtu.dk

Benjamin Recht
UC Berkeley

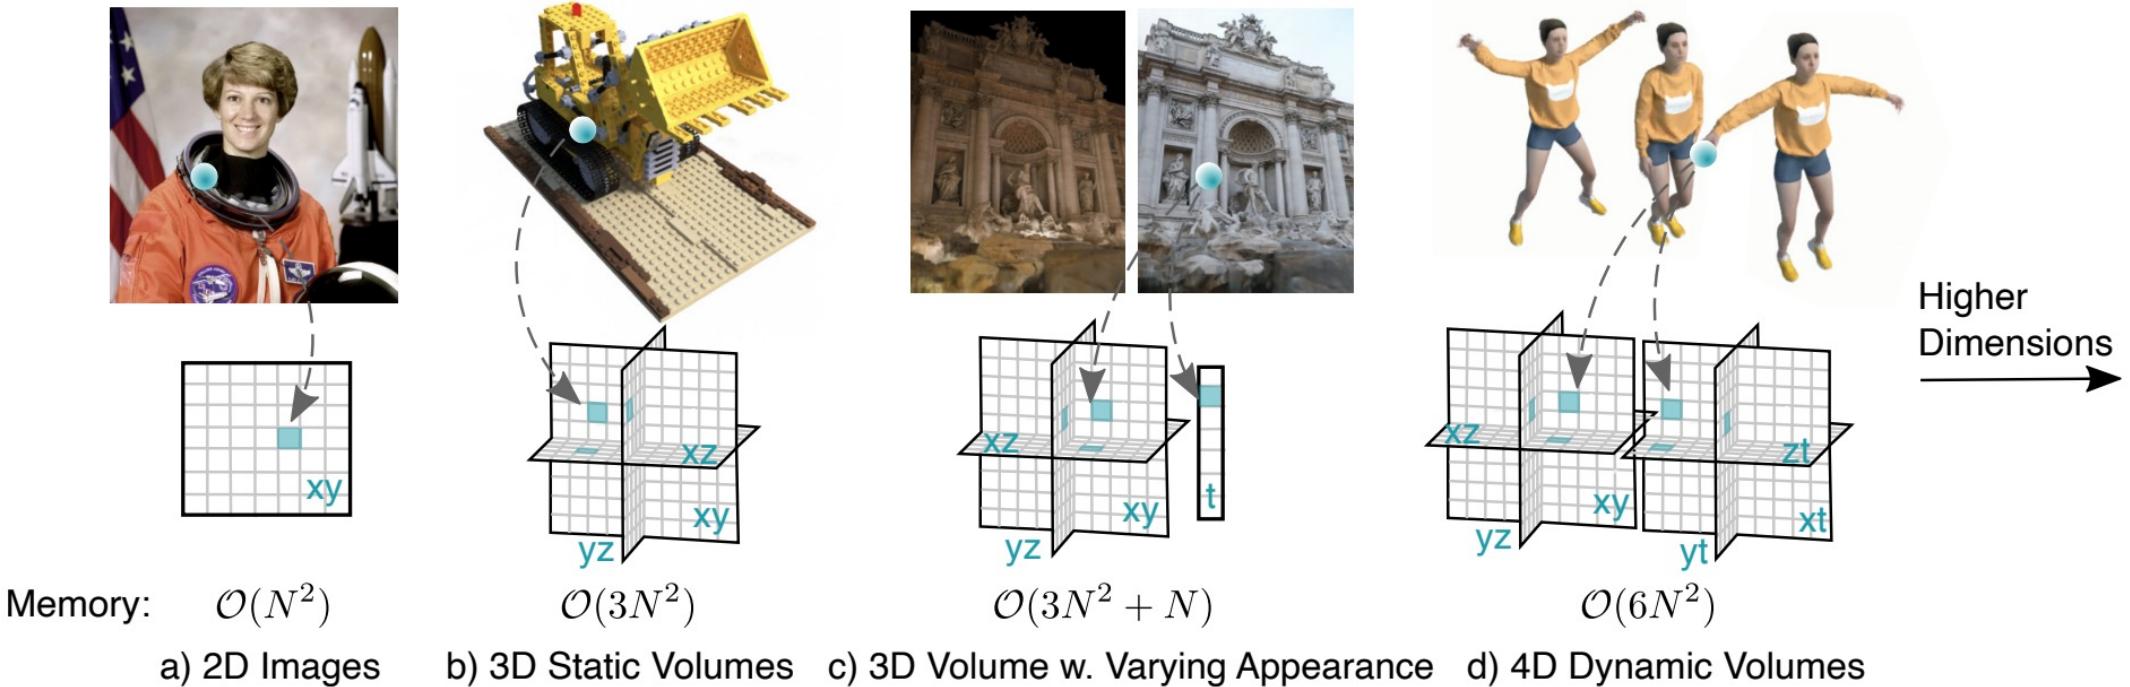
brecht@berkeley.edu

Angjoo Kanazawa
UC Berkeley

kanazawa@berkeley.edu



4D Representations | 2D Planes | K–Planes



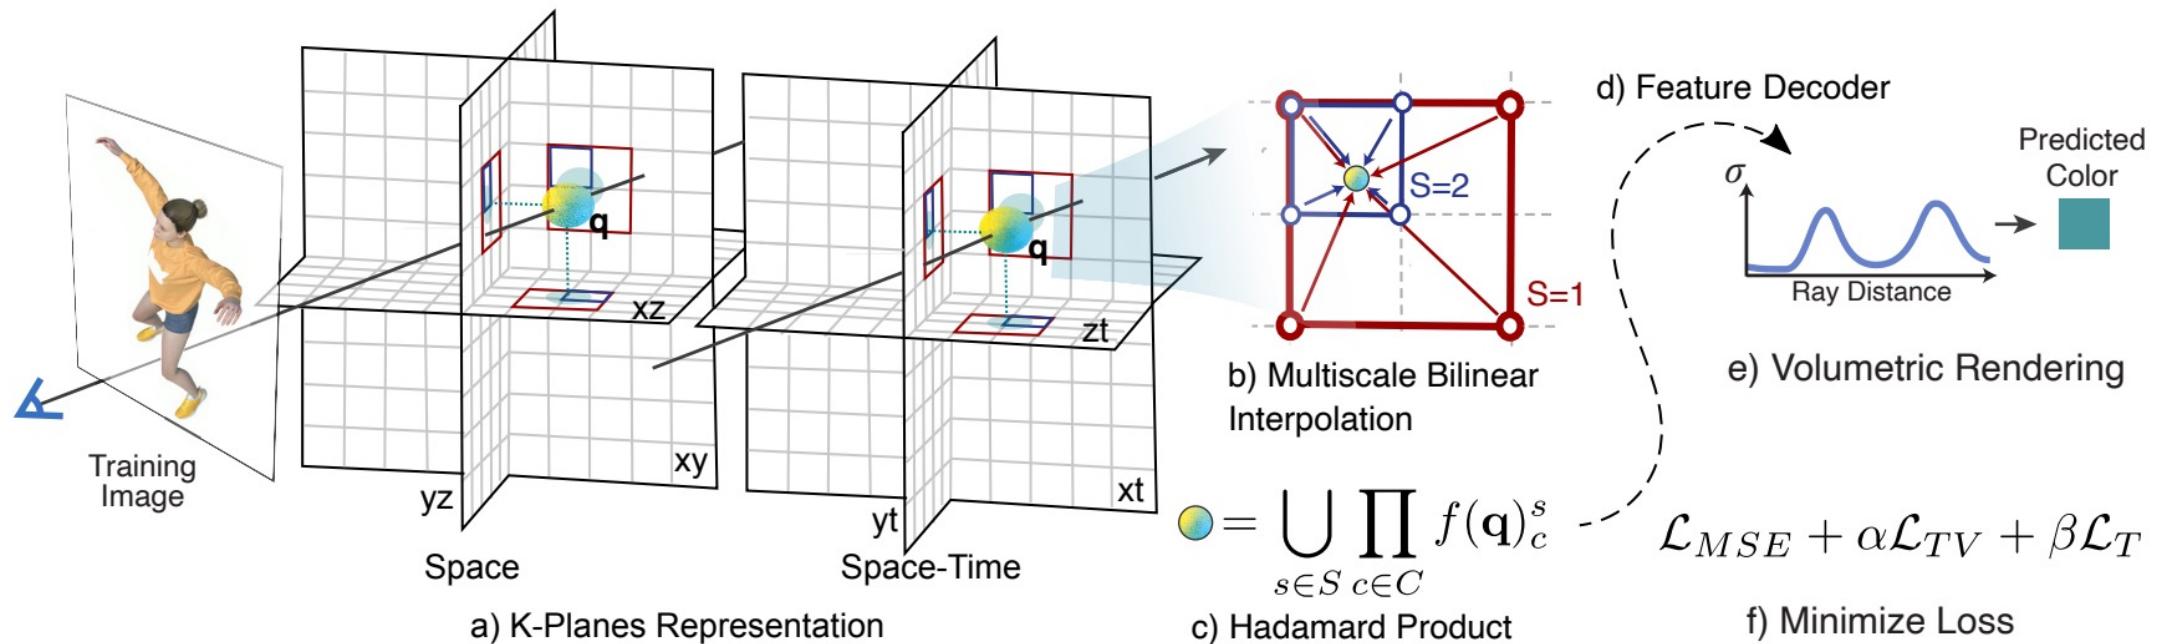
Memory: $\mathcal{O}(N^2)$

a) 2D Images

$\mathcal{O}(3N^2)$

b) 3D Static Volumes c) 3D Volume w. Varying Appearance d) 4D Dynamic Volumes

4D Representations | 2D Planes | K–Planes

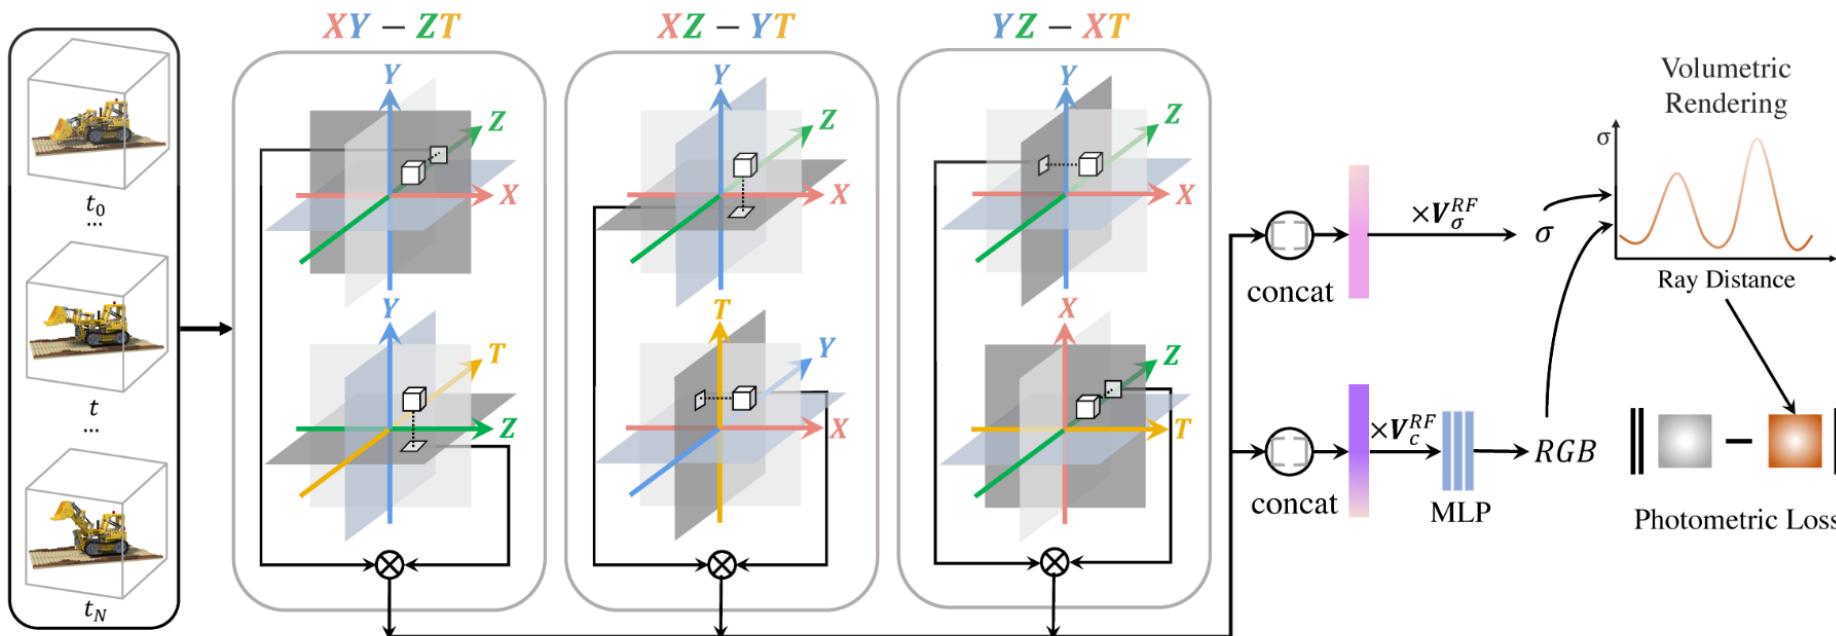


HexPlane: A Fast Representation for Dynamic Scenes

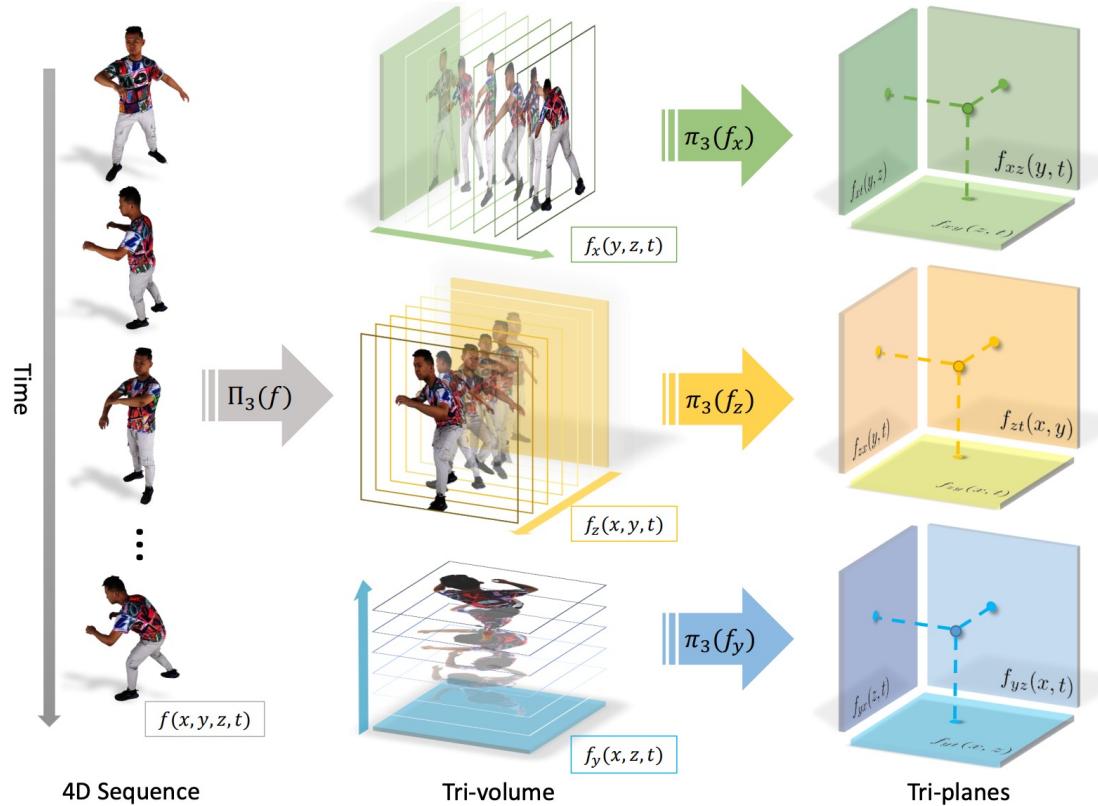
Ang Cao, Justin Johnson

University of Michigan

CVPR 2023



4D Representations | 2D Planes | Tensor4D



- **4D space as 3D volumes**

$$\Pi_3(f(x, y, z, t)) = \{f_z(x, y, t), f_y(x, z, t), f_x(y, z, t)\}$$

- **3D volume as 2D maps**

$$\pi_3(f_z) = \{f_{zt}(x, y), f_{zy}(x, t), f_{zx}(y, t)\}$$

$$\pi_3(f_y) = \{f_{yt}(x, z), f_{yz}(x, t), f_{yx}(z, t)\}$$

$$\pi_3(f_x) = \{f_{xt}(y, z), f_{xz}(y, t), f_{xy}(z, t)\}$$

- **Deeper correlation with t axis**

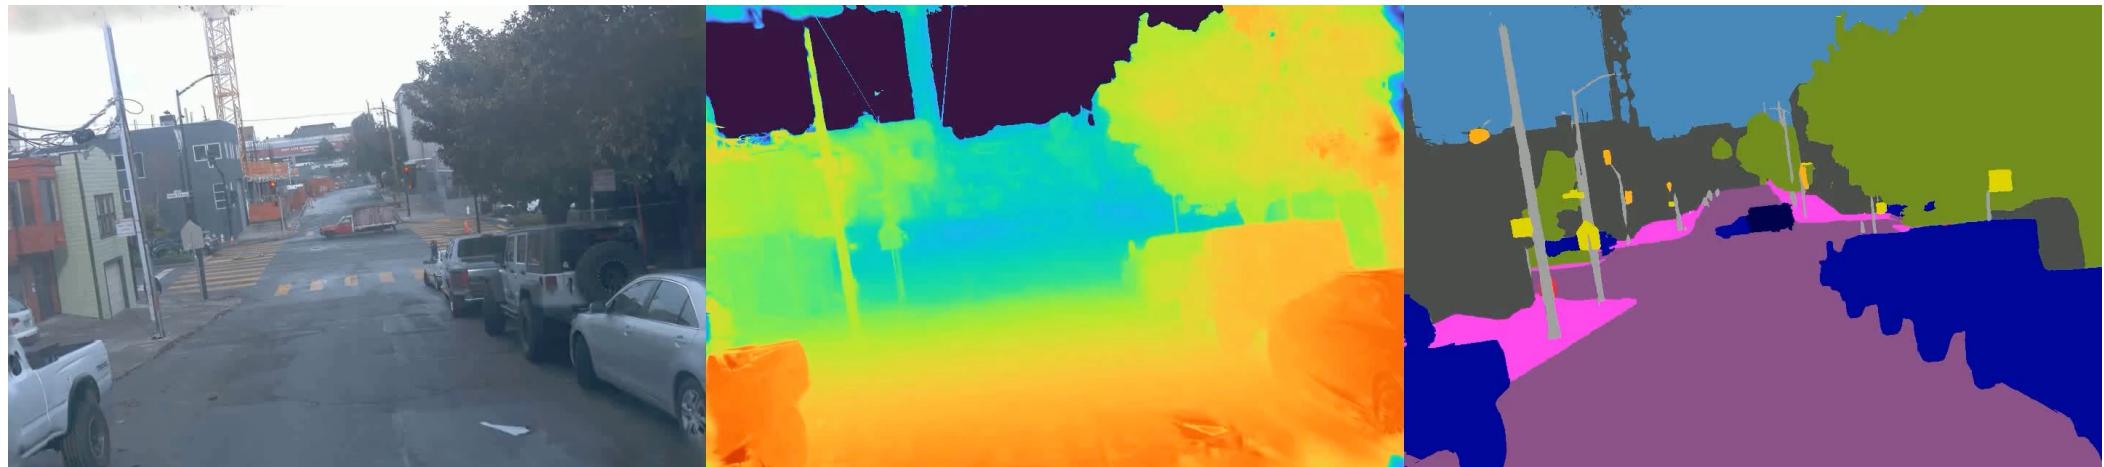
$$(xt, yt, zt) \longrightarrow (xt, yt, zt, xyt, yzt, xzt, xyzt)$$

Real-time Photorealistic Dynamic Scene Representation and Rendering with 4D Gaussian Splatting

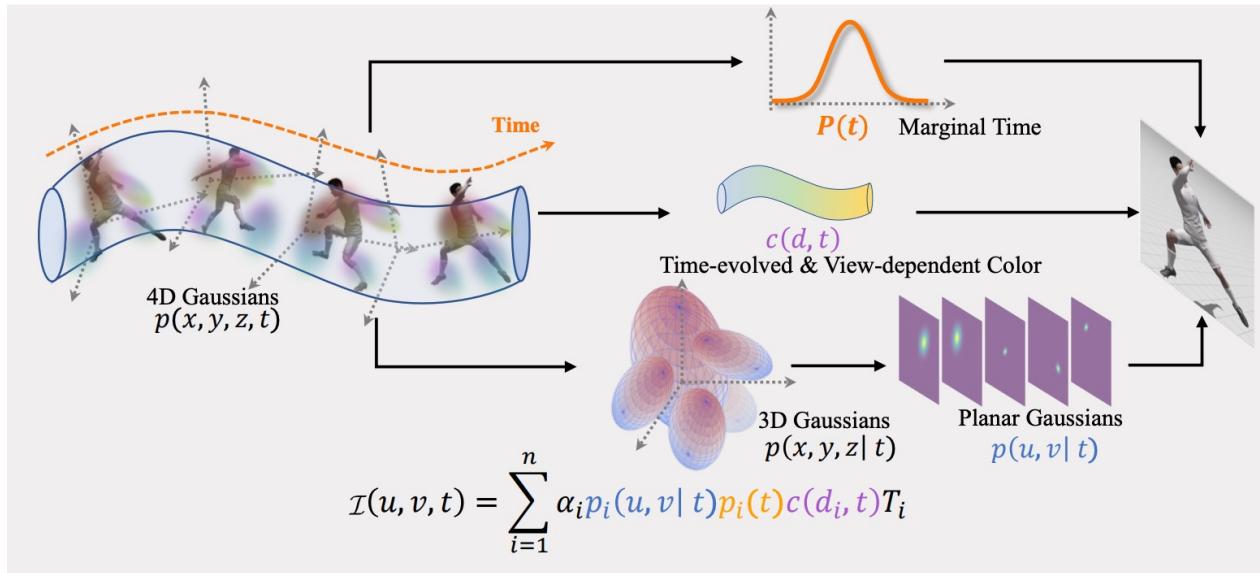
ICLR 2024

Zeyu Yang, Hongye Yang, Zijie Pan, Li Zhang✉

Fudan University



4D Representations | Gaussians | 4DGS



Method	PSNR ↑	DSSIM ↓	LPIPS ↓	FPS ↑
<i>- Plenoptic Video (real, multi-view)</i>				
Neural Volumes (Lombardi et al., 2019) ¹	22.80	0.062	0.295	-
LLFF (Mildenhall et al., 2019) ¹	23.24	0.076	0.235	-
DyNeRF (Li et al., 2022b) ¹	29.58	0.020	0.099	0.015
HexPlane (Cao & Johnson, 2023)	31.70	0.014	0.075	0.56 ³
K-Planes-explicit (Fridovich-Keil et al., 2023)	30.88	0.020	-	0.23 ³
K-Planes-hybrid (Fridovich-Keil et al., 2023)	31.63	0.018	-	-
MixVoxels-L (Wang et al., 2023)	30.80	0.020	0.126	16.7
StreamRF (Li et al., 2022a) ¹	29.58	-	-	8.3
NeRFPlayer (Song et al., 2023)	30.69	0.035 ²	0.111	0.045
HyperReel (Attal et al., 2023)	31.10	0.037 ²	0.096	2.00
4DGS (Wu et al., 2023) ⁴	31.02	0.030	0.150	36
4DGS (Ours)	32.01	0.014	0.055	114

- Defining a Gaussian point in 4D space

4D-Rotor Gaussian Splatting: Towards Efficient Novel View Synthesis for Dynamic Scenes

Yuanxing Duan^{*}
Peking University
China
mjdyx@pku.edu.cn

Yuhang He
Peking University
China
2100014725@stu.pku.edu.cn

Fangyin Wei^{*}
Princeton University
USA
fwei@princeton.edu

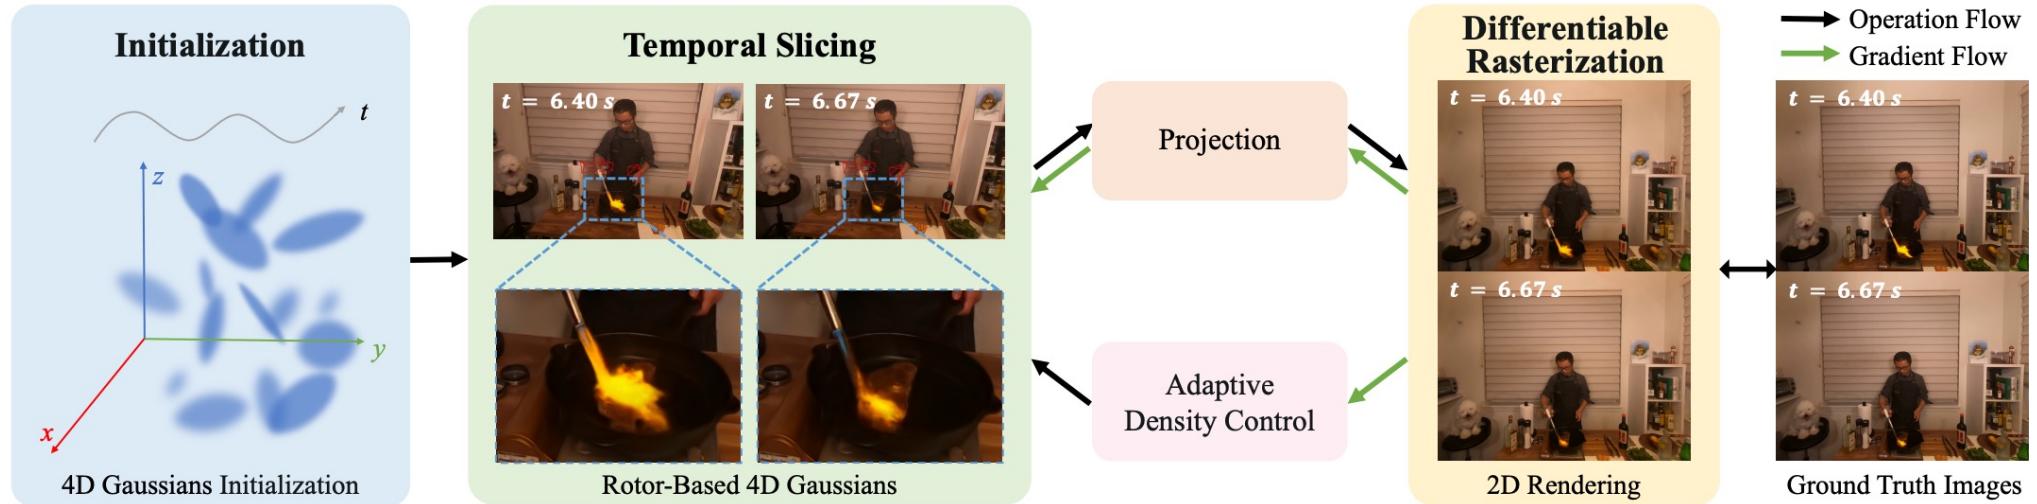
Wenzheng Chen[†]
Peking University
China
NVIDIA
Canada
wenzhengchen@pku.edu.cn

Qiyu Dai
Peking University
China
State Key Laboratory of General AI
China
qiyudai@pku.edu.cn

Baoquan Chen[†]
Peking University
China
State Key Laboratory of General AI
China
baoquan@pku.edu.cn

- **Input:** multi-view video
- **Output:** 4D Gaussian

4D Representations | Gaussians | 4D–Rotor GS



Method	PSNR↑	SSIM↑	LPIPS↓	Train↓	FPS↑
D-NeRF [Pumarola et al. 2021]	29.17	0.95	0.07	24 h	0.13
TiNeuVox [Fang et al. 2022]	32.87	0.97	0.04	28 min	1.60
K-Planes [Fridovich-Keil et al. 2023]	31.07	0.97	0.02	54 min	1.20
Deformable3DGS [Yang et al. 2023]	39.31	0.99	0.01	26 min	85.45
Deformable4DGS [Wu et al. 2023]	32.99	0.97	0.05	13 min	104.00
RealTime4DGS [Yang et al. 2024]	32.71	0.97	0.03	10 min	289.07
Ours	34.26	0.97	0.03	5 min	1257.63

$$G_{4D}(\mathbf{x}) = e^{-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu}_{4D})^T \boldsymbol{\Sigma}_{4D}^{-1} (\mathbf{x}-\boldsymbol{\mu}_{4D})}.$$

$$\boldsymbol{\Sigma}_{4D} = \mathbf{R}_{4D} \mathbf{S}_{4D} \mathbf{S}_{4D}^T \mathbf{R}_{4D}^T.$$

$$\boldsymbol{\Sigma}_{4D} = \begin{pmatrix} \mathbf{U} & \mathbf{V} \\ \mathbf{V}^T & \mathbf{W} \end{pmatrix} \text{ and } \boldsymbol{\Sigma}_{4D}^{-1} = \begin{pmatrix} \mathbf{A} & \mathbf{M} \\ \mathbf{M}^T & \mathbf{Z} \end{pmatrix},$$

- **NeRF Related:** D–NeRF, Nerfies, Dynamic NeRF
- **Explicit NeRF Related:** MixVoxel, NeRFPlayer
- **3DGS Related:** 4D–GS, Deformable 3DGS, Gaussian–Flow, Spacetime GS, SC–GS

D-NeRF: Neural Radiance Fields for Dynamic Scenes

Albert Pumarola¹

Enric Corona¹

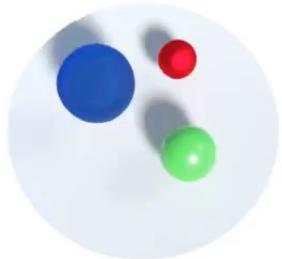
Gerard Pons-Moll^{2,3}

Francesc Moreno-Noguer¹

¹Institut de Robòtica i Informàtica Industrial, CSIC-UPC

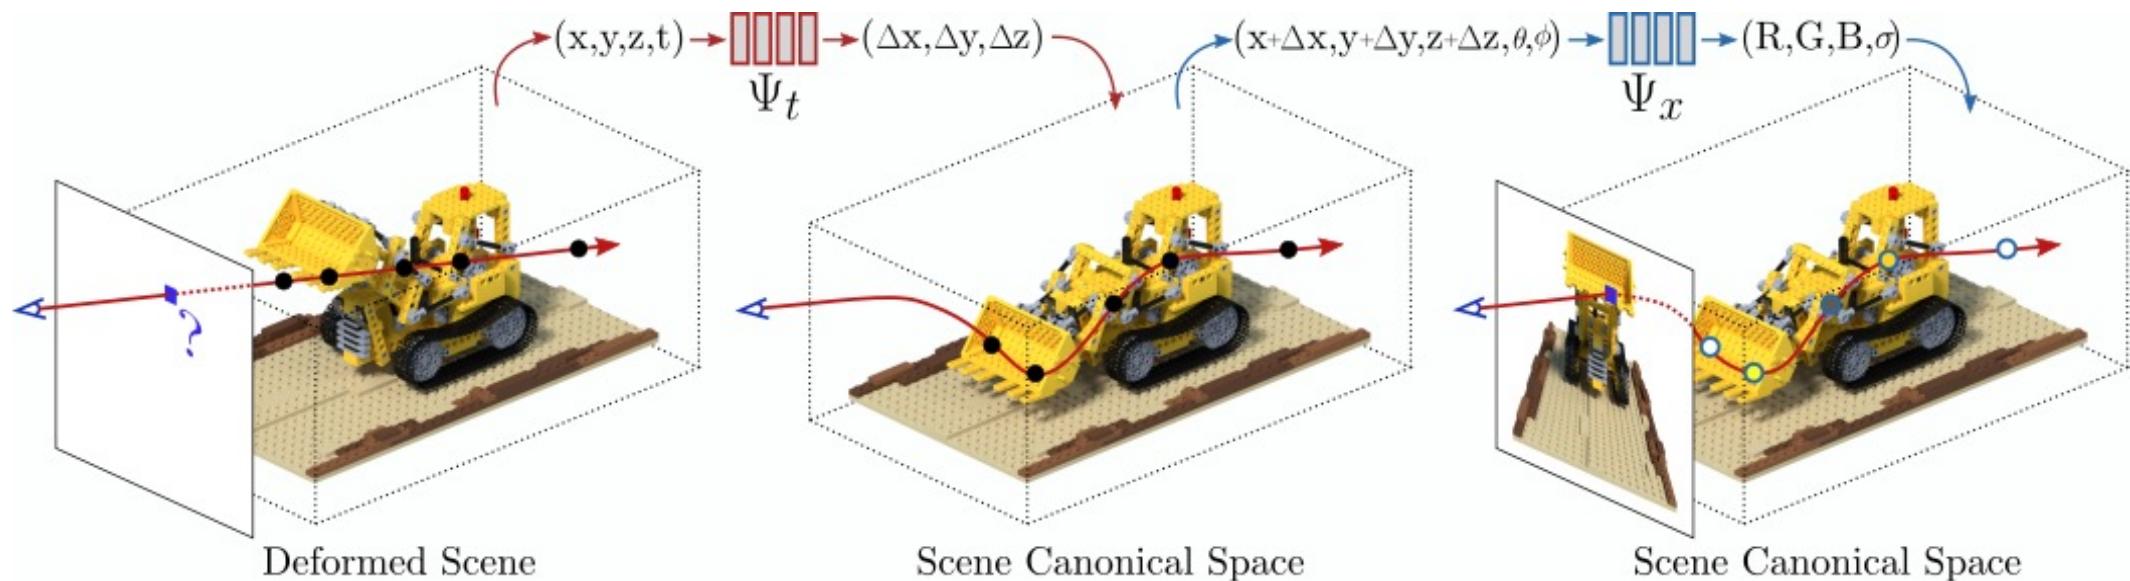
²University of Tübingen

³Max Planck Institute for Informatics



- **Input:** monocular video
- **Output:** radiance & deformation fields

Deformation Fields | NeRFs | D–NeRF



Deformation Field: $\Psi_t : (\mathbf{x}, t) \rightarrow \Delta \mathbf{x}$

Radiance Field: $\Psi_x : (\mathbf{x} + \Delta \mathbf{x}, \mathbf{d}) \rightarrow (\mathbf{c}, \sigma)$

Nerfies: Deformable Neural Radiance Fields

Keunhong Park^{1*}

Utkarsh Sinha²

Jonathan T. Barron²

Sofien Bouaziz²

Dan B Goldman²

Steven M. Seitz^{1,2}

Ricardo Martin-Brualla²

¹University of Washington

²Google Research



(a) Capture Process



(b) Input



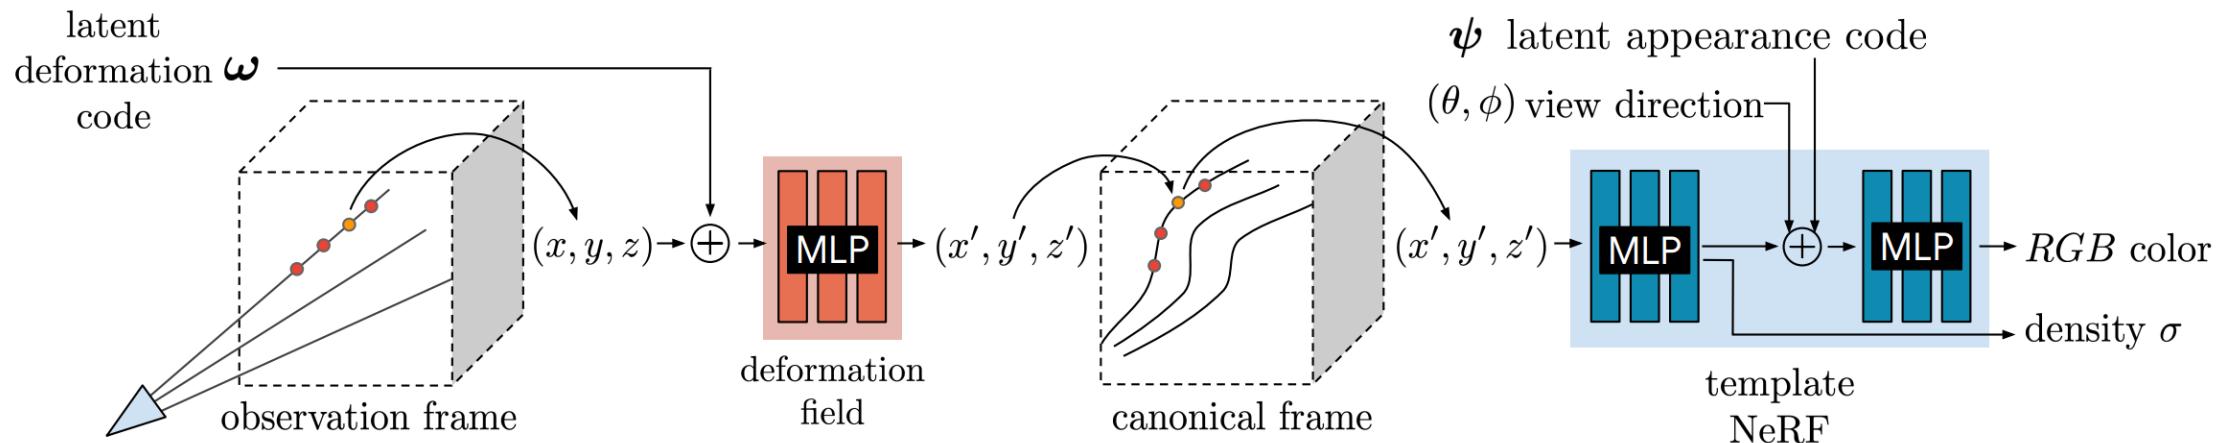
(c) Nerfie



(d) Nerfie Depth

- **Input:** multi-view images of non-rigid scenes
- **Output:** radiance & deformation fields

Deformation Fields | NeRFs | Nerfies



Deformation Field: $T : (\mathbf{x}, \omega_i) \rightarrow \mathbf{x}'$
(with per-frame deformation code)

Radiance Field: $F : (\mathbf{x}, \mathbf{d}, \psi_i) \rightarrow (\mathbf{c}, \sigma)$
(with per-frame appearance code)

Dynamic View Synthesis from Dynamic Monocular Video

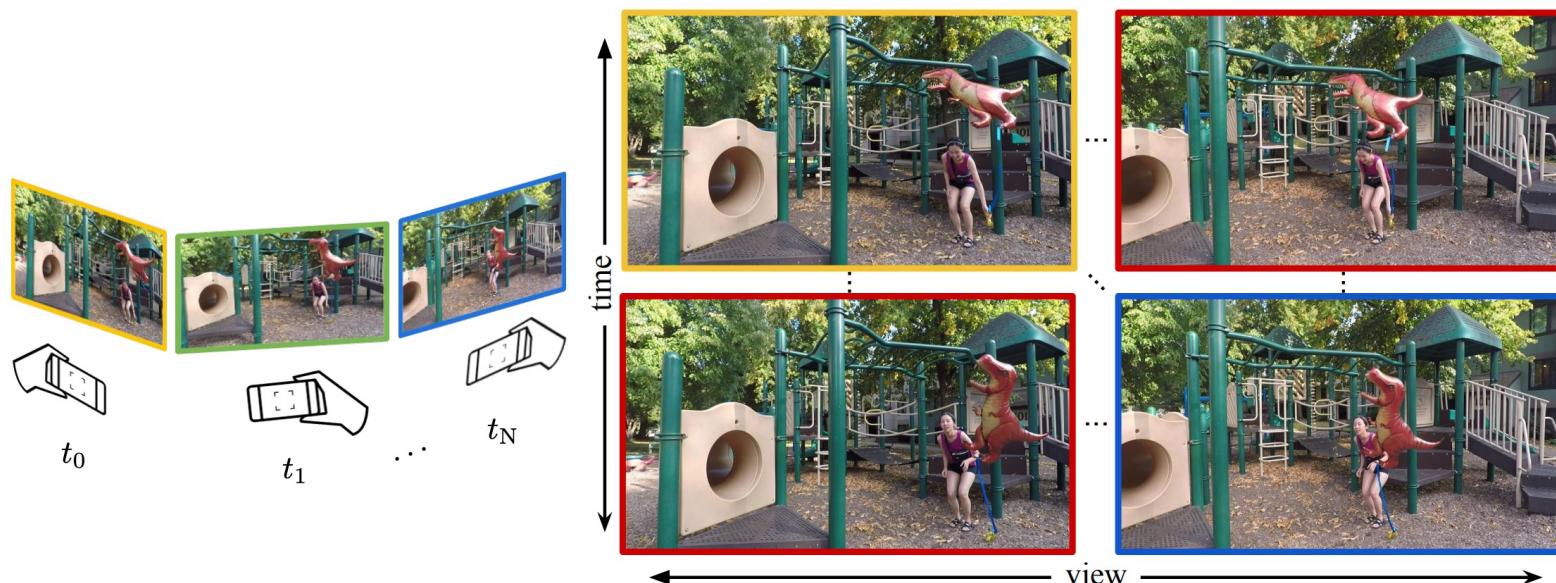
Chen Gao
Virginia Tech
`chengao@vt.edu`

Ayush Saraf
Facebook
`ayush29feb@fb.com`

Johannes Kopf
Facebook
`jkopf@fb.com`

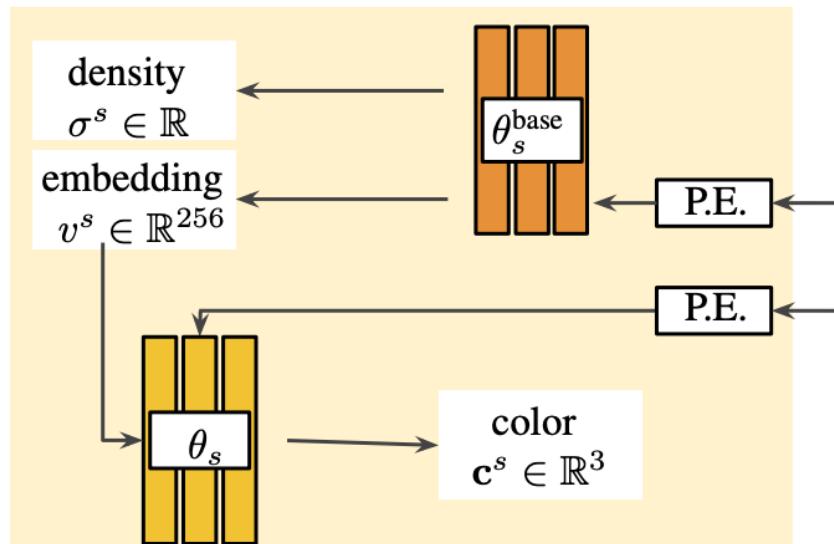
Jia-Bin Huang
Virginia Tech
`jbhuang@vt.edu`

<https://free-view-video.github.io>

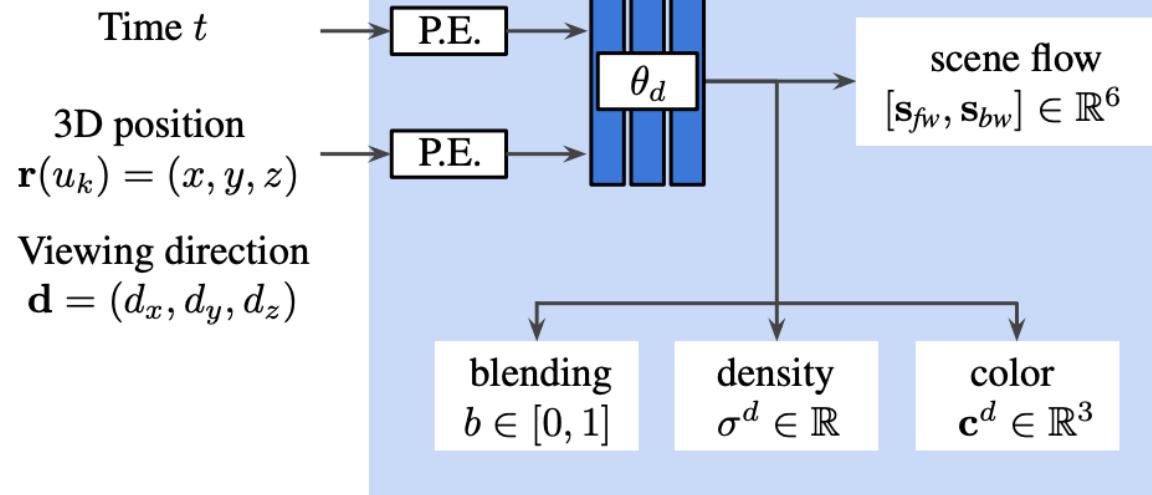


- **Input:** monocular video
- **Output:** static NeRF + dynamic NeRF with scene flow

Deformation Fields | NeRFs | Dynamic NeRF



(a) Static NeRF (Section 3.2)



(b) Dynamic NeRF (Section 3.3)

Scene Flow and Radiance Field: $(\mathbf{s}_{fw}, \mathbf{s}_{bw}, \sigma_t^d, \mathbf{c}_t^d, b) = \text{MLP}_{\theta_d}(\mathbf{r}(u_k), t)$
(for dynamic foreground areas)

Static Radiance Field: $(\sigma^s, \mathbf{c}^s) = \text{MLP}_{\theta}(\mathbf{r}(u_k))$
(for static background areas)

Mixed Neural Voxels for Fast Multi-view Video Synthesis

Feng Wang¹ Sinan Tan¹ Xinghang Li¹ Zeyue Tian² Yafei Song³ Huaping Liu¹ *

¹Beijing National Research Center for Information Science and Technology(BNRist),
Department of Computer Science and Technology, Tsinghua University

²Hong Kong University of Science and Technology

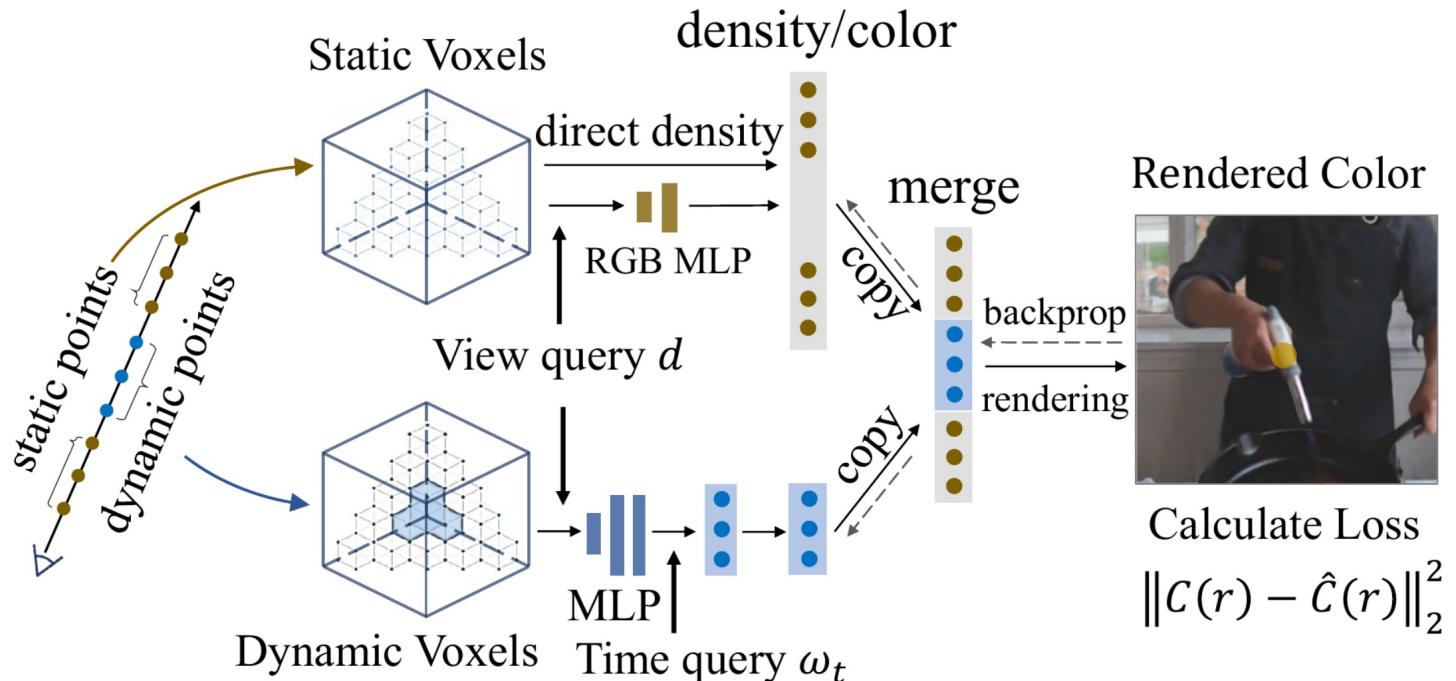
³ XR Lab, DAMO Academy, Alibaba Group

wang-f20@mails.tsinghua.edu.cn, hpliu@tsinghua.edu.cn



- **Input:** monocular/multi–view video
- **Output:** static voxels + dynamic voxels

Deformation Fields | Voxel | MixVoxel



Static Radiance Field:
(for static background areas)

$$\sigma(x, y, z) = S_{x,y,z}^\sigma, \quad c(x, y, z, \mathbf{d}) = \mathcal{C}_\theta(S_{x,y,z}^c, \mathbf{d}).$$

Dynamic Radiance Field:
(for dynamic foreground areas)

$$\begin{aligned}\sigma(x, y, z, t) &= \omega_t^\sigma \cdot \mathcal{T}_{\theta_1}^\sigma(\mathcal{G}_{x,y,z}^\sigma), \\ c(x, y, z, \mathbf{d}, t) &= \omega_t^c \cdot \mathcal{T}_{\theta_2}^c(\mathcal{G}_{x,y,z}^c, \mathbf{d}).\end{aligned}$$

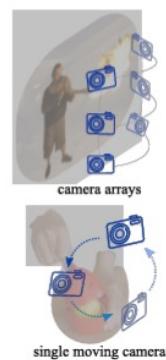
NeRFPlayer: A Streamable Dynamic Scene Representation with Decomposed Neural Radiance Fields

Liangchen Song¹ Anpei Chen^{2,4} Zhong Li³ Zhang Chen³ Lele Chen³
Junsong Yuan¹ Yi Xu³ Andreas Geiger⁴

¹University at Buffalo ²ETH Zürich

³OPPO US Research Center, InnoPeak Tech ⁴University of Tübingen

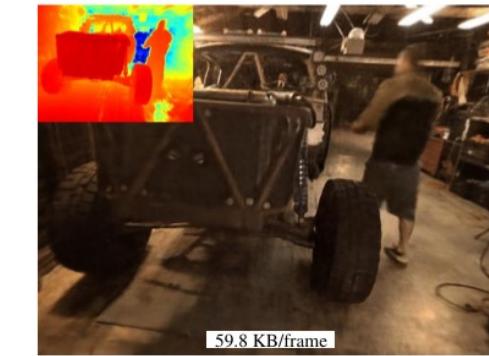
<https://bit.ly/nerfplayer>



(a) Inputs



(b) Real-time rendering



59.8 KB/frame

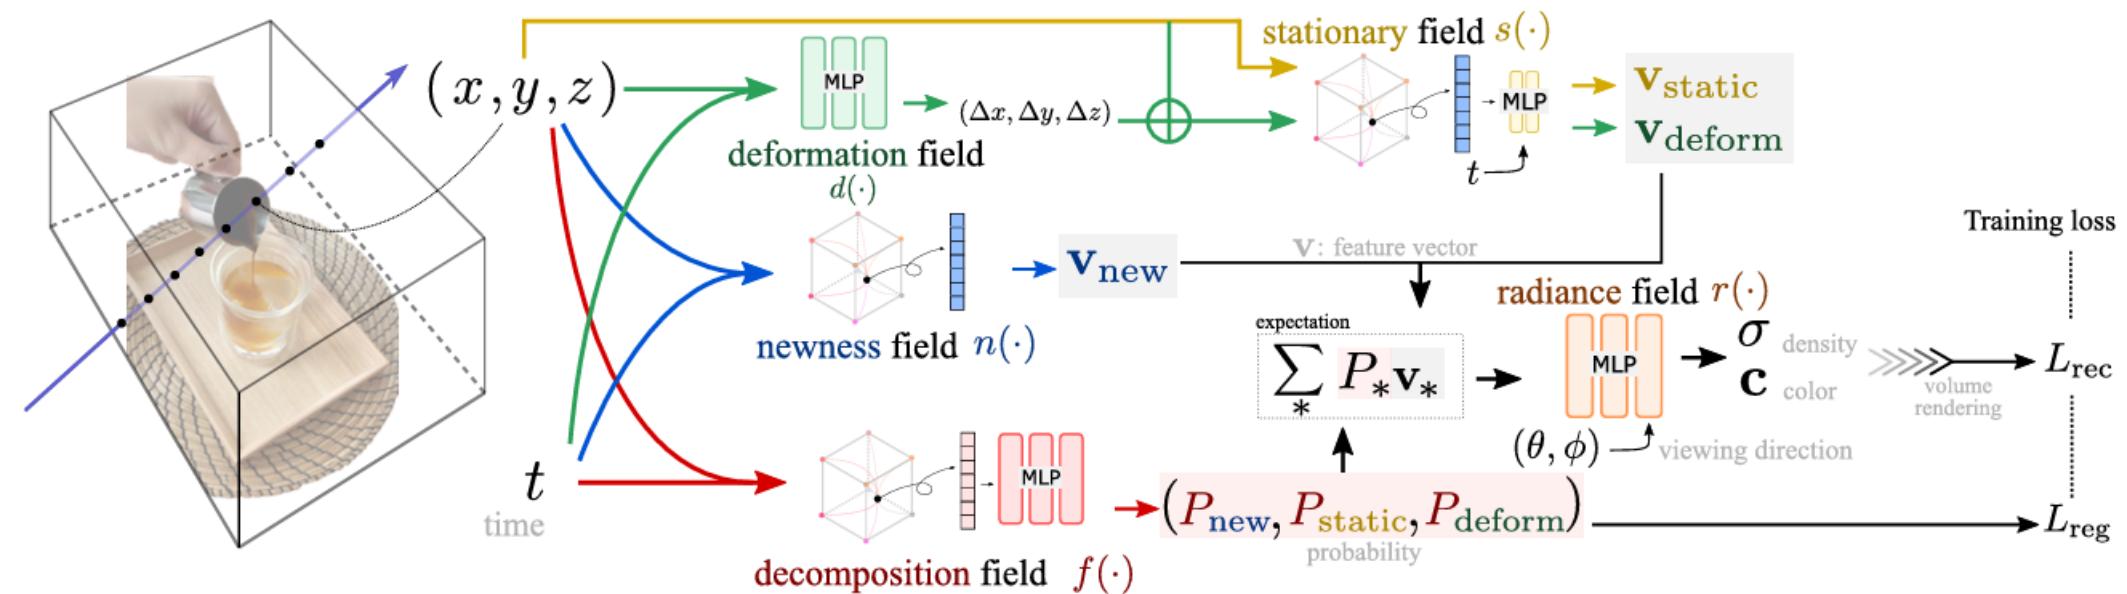


78.8 KB/frame

(c) Low bitrate streaming

- **Input:** monocular/multi–view video
- **Output:** static NeRF + deform MLP + new MLP

Deformation Fields | Instant-NPG | NeRFPlayer



Deformation Field: $d(\cdot) : (\mathbf{p}, t) \mapsto (\Delta \mathbf{p})$.

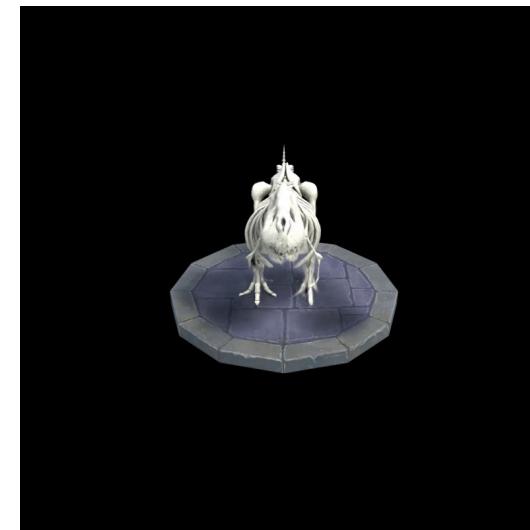
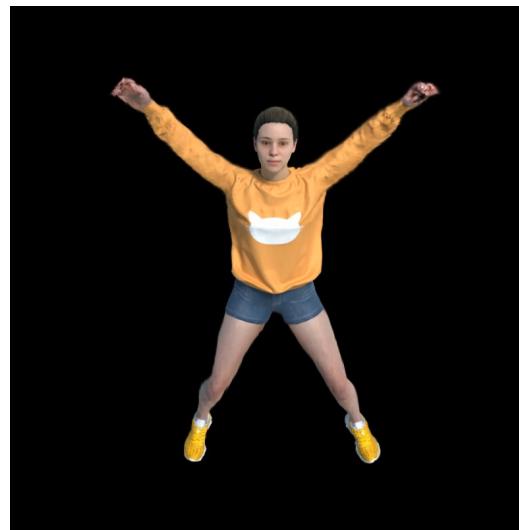
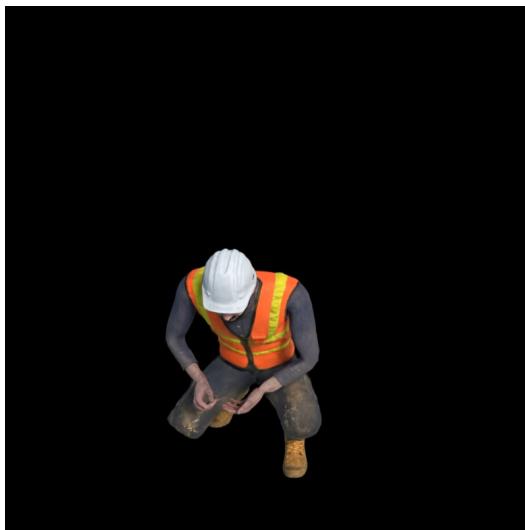
$f(\cdot) : (\mathbf{p}, t) \mapsto (P_{\text{static}}, P_{\text{deform}}, P_{\text{new}})$,

$\mathbf{v} = \sum_* P_* \mathbf{v}_*, * \in \{\text{static, deform, new}\}$.

Deformable 3D Gaussians for High-Fidelity Monocular Dynamic Scene Reconstruction

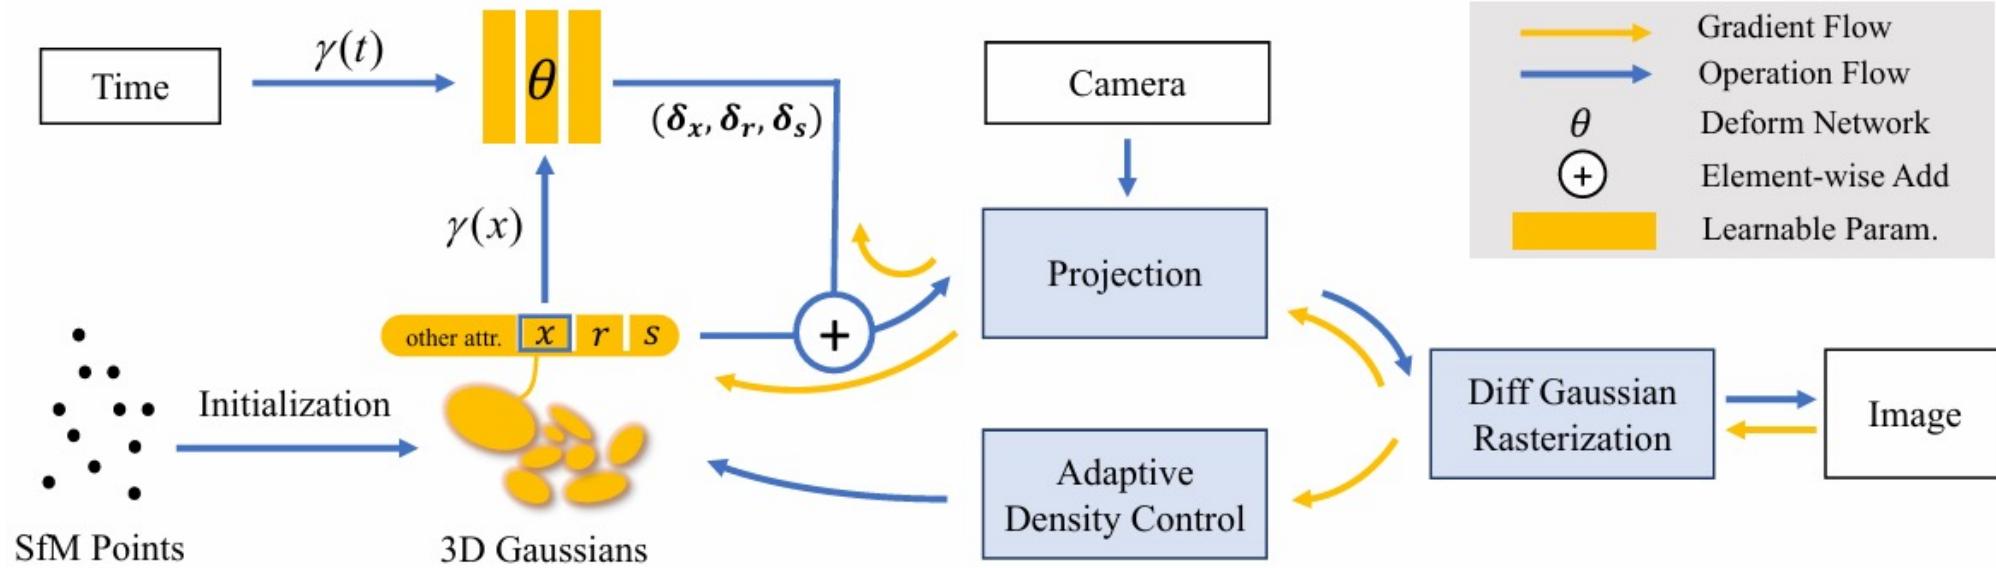
Ziyi Yang^{1,2} Xinyu Gao¹ Wen Zhou² Shaohui Jiao² Yuqing Zhang¹ Xiaogang Jin^{1†}

¹State Key Laboratory of CAD&CG, Zhejiang University ²ByteDance Inc.



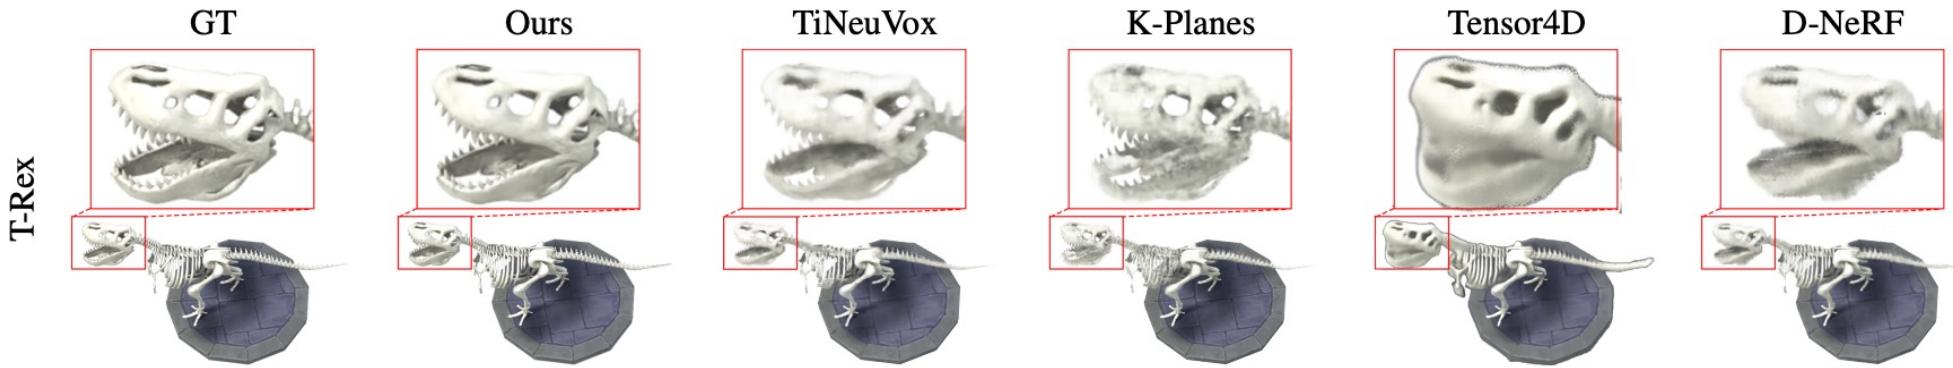
- **Input:** multi-view/monocular video
- **Output:** dynamic gaussian

Deformation Fields | Gaussians | Deformable 3DGs



$$\text{Deformation Field: } (\delta x, \delta r, \delta s) = \mathcal{F}_\theta(\gamma(\text{sg}(x)), \gamma(t))$$

Deformation Fields | Gaussians | Deformable 3DGs



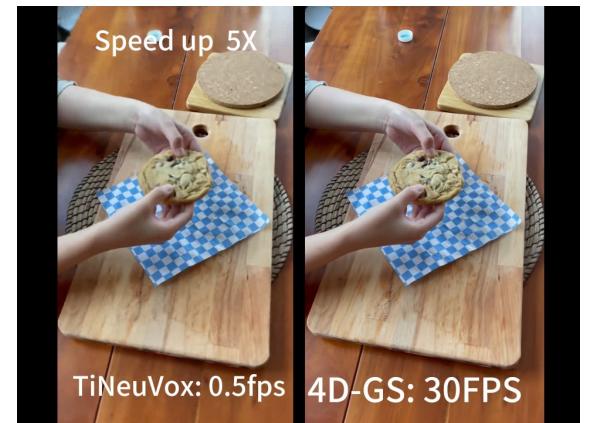
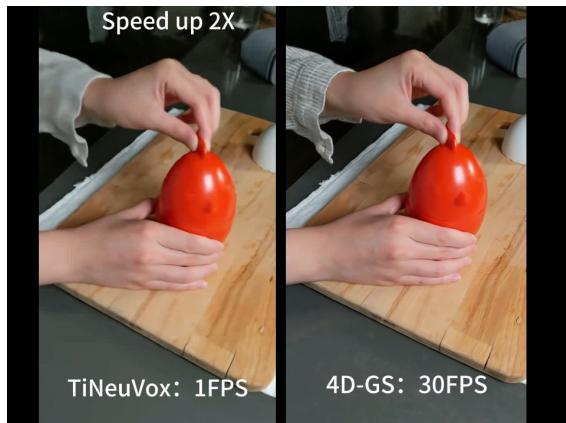
Method	Hell Warrior				Mutant				Hook			Bouncing Balls			
	PSNR↑	SSIM↑	LPIPS↓	PSNR↑	SSIM↑	LPIPS↓	PSNR↑	SSIM↑	LPIPS↓	PSNR↑	SSIM↑	LPIPS↓	PSNR↑	SSIM↑	LPIPS↓
3D-GS	29.89	0.9155	0.1056	24.53	0.9336	0.0580	21.71	0.8876	0.1034	23.20	0.9591	0.0600			
D-NeRF	24.06	0.9440	0.0707	30.31	0.9672	0.0392	29.02	0.9595	0.0546	38.17	0.9891	0.0323			
TiNeuVox	27.10	0.9638	0.0768	31.87	0.9607	0.0474	30.61	0.9599	0.0592	40.23	0.9926	0.0416			
Tensor4D	31.26	0.9254	0.0735	29.11	0.9451	0.0601	28.63	0.9433	0.0636	24.47	0.9622	0.0437			
K-Planes	24.58	0.9520	0.0824	32.50	0.9713	0.0362	28.12	0.9489	0.0662	40.05	0.9934	0.0322			
Ours	41.54	0.9873	0.0234	42.63	0.9951	0.0052	37.42	0.9867	0.0144	41.01	0.9953	0.0093			
Method	Lego				T-Rex				Stand Up			Jumping Jacks			
	PSNR↑	SSIM↑	LPIPS↓	PSNR↑	SSIM↑	LPIPS↓	PSNR↑	SSIM↑	LPIPS↓	PSNR↑	SSIM↑	LPIPS↓	PSNR↑	SSIM↑	LPIPS↓
3D-GS	22.10	0.9384	0.0607	21.93	0.9539	0.0487	21.91	0.9301	0.0785	20.64	0.9297	0.0828			
D-NeRF	25.56	0.9363	0.0821	30.61	0.9671	0.0535	33.13	0.9781	0.0355	32.70	0.9779	0.0388			
TiNeuVox	26.64	0.9258	0.0877	31.25	0.9666	0.0478	34.61	0.9797	0.0326	33.49	0.9771	0.0408			
Tensor4D	23.24	0.9183	0.0721	23.86	0.9351	0.0544	30.56	0.9581	0.0363	24.20	0.9253	0.0667			
K-Planes	28.91	0.9695	0.0331	30.43	0.9737	0.0343	33.10	0.9793	0.0310	31.11	0.9708	0.0468			
Ours	33.07	0.9794	0.0183	38.10	0.9933	0.0098	44.62	0.9951	0.0063	37.72	0.9897	0.0126			

4D Gaussian Splatting for Real-Time Dynamic Scene Rendering

Guanjun Wu^{1*}, Taoran Yi^{2*}, Jiemin Fang^{3†}, Lingxi Xie³, Xiaopeng Zhang³,
Wei Wei¹, Wenyu Liu², Qi Tian³, Xinggang Wang^{2‡}

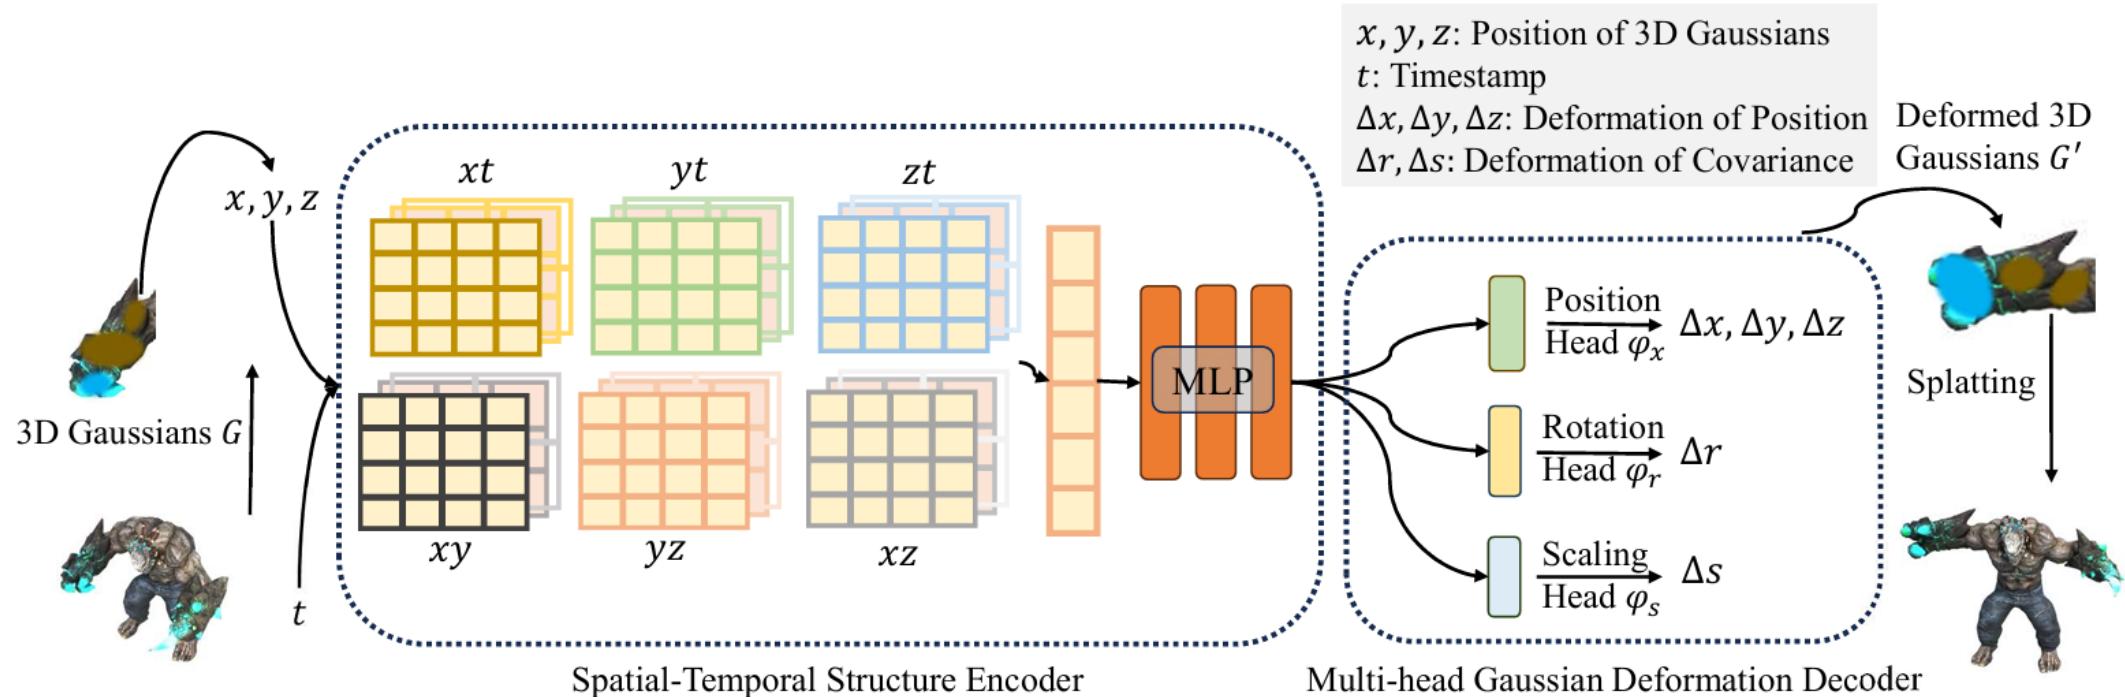
¹School of CS, Huazhong University of Science and Technology

²School of EIC, Huazhong University of Science and Technology ³Huawei Inc.



- **Input:** multi-view/monocular video
- **Output:** dynamic gaussian

Deformation Fields | Gaussians | 4D-GS



Deformation Field:

$$f_h = \bigcup_l \prod \text{interp}(R_l(i, j)),$$

$$(i, j) \in \{(x, y), (x, z), (y, z), (x, t), (y, t), (z, t)\}.$$

$$f_d = \phi_d(f_h). \quad \Delta \mathcal{X} = \phi_x(f_d) \quad \Delta r = \phi_r(f_d) \quad \Delta s = \phi_s(f_d)$$

Gaussian–Flow: 4D Reconstruction with Dynamic 3D Gaussian Particle

Youtian Lin¹

¹Nanjing University

Zuozhuo Dai²

²Alibaba Group

Siyu Zhu³

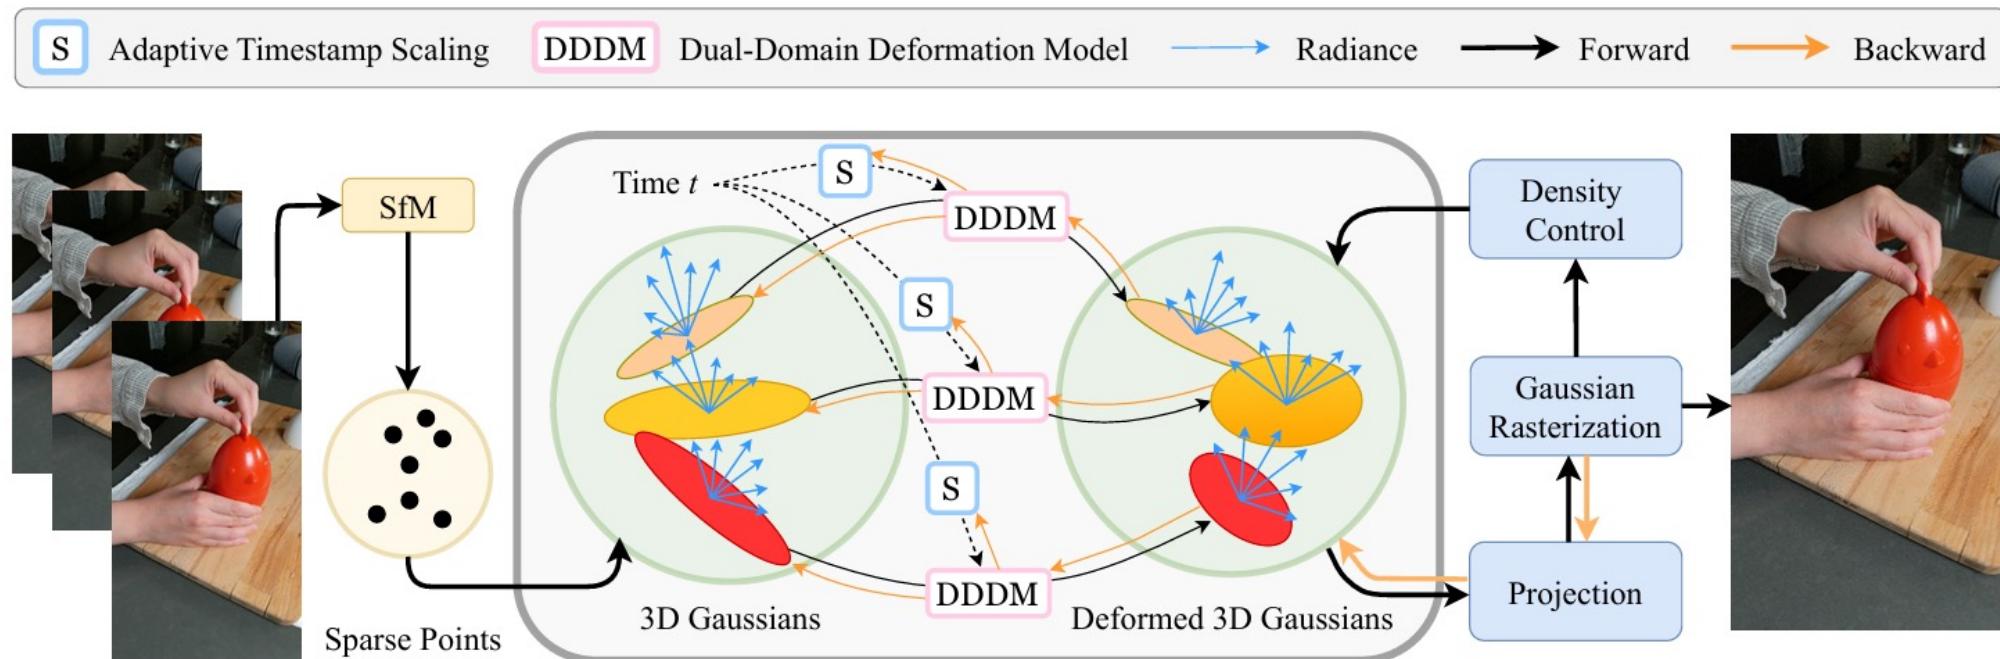
³Fudan University

Yao Yao^{1✉}



- **Input:** multi–view/monocular video
- **Output:** dynamic gaussian

Deformation Fields | Gaussians | Gaussian-Flow



Deformation Field: $D(t) = P_N(t) + F_L(t)$ $P_N(t) = \sum_{n=0}^N a_n t^n$, $F_L(t) = \sum_{l=1}^L (f_{sin}^l \cos(lt) + f_{cos}^l \sin(lt))$.

Regularizations: $\mathcal{L}_t = \|D(t) - D(t + \epsilon)\|_2$. $\mathcal{L}_s = \sum_{j \in \mathcal{N}_i} \|D(t)_i - D(t)_j\|_2$

- Explicit motion modeling
- As fast as static 3DGS

Deformation Fields | Gaussians | Gaussian–Flow

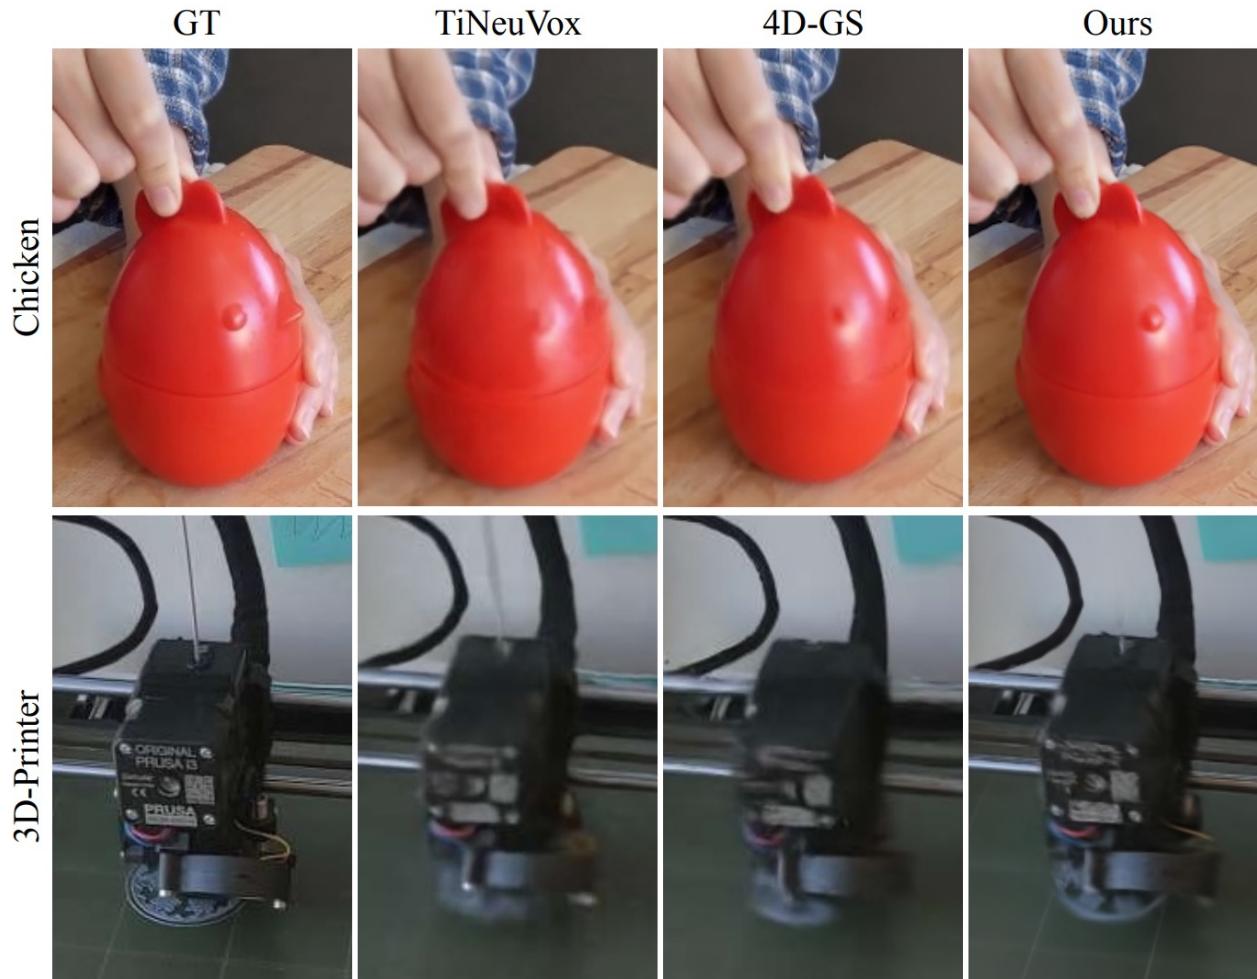


Table 4. Quantitative comparison on Plenoptic Video dataset [18]. Our training speed is $5\times$ faster of magnitude faster than previous leading approaches, . Also, we achieved the highest PSNR score among all methods.

Method	Train Time ↓	PSNR ↑	SSIM ↑
LLFF [24]	-	23.2	-
DyNeRF [18]	1344 hours	29.6	0.96
NeRFPlayer [32]	5.5 hours	30.7	-
K-Planes [8]	1.8 hours	31.6	0.96
4D-GS [35]	2 hours	31.0	0.94
Ours (30K)	22.5 min	30.5	0.97
Ours (60K)	41.8 min	32.0	0.97

Spacetime Gaussian Feature Splatting for Real-Time Dynamic View Synthesis

Zhan Li^{1,2*}

¹ OPPO US Research Center

lizhan@pdx.edu

Zhang Chen^{1†}

zhang.chen@oppo.com

Zhong Li^{1†}

² Portland State University

Yi Xu¹

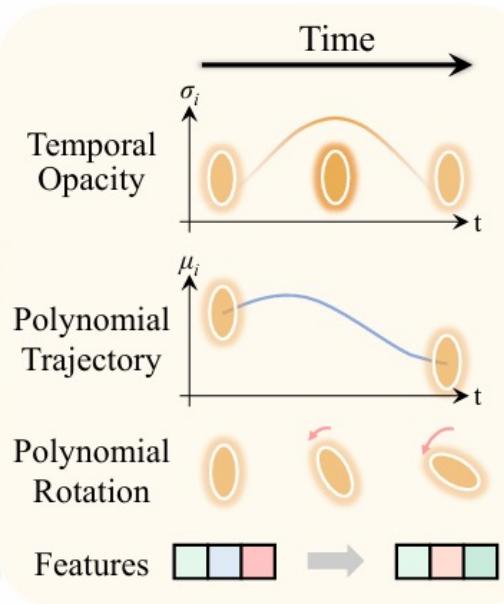
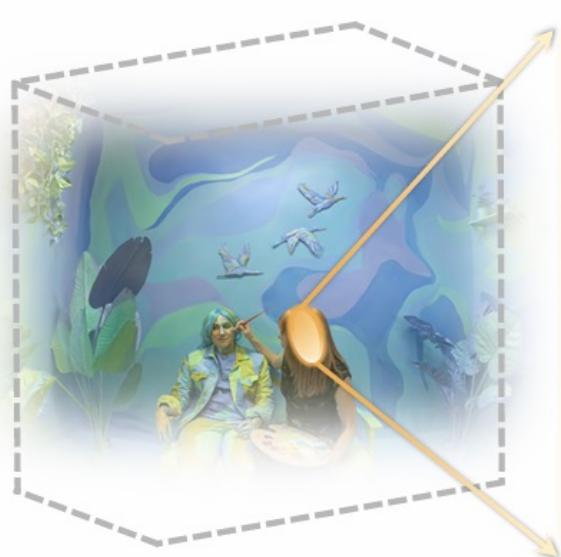
zhong.li@oppo.com

yi.xu@oppo.com

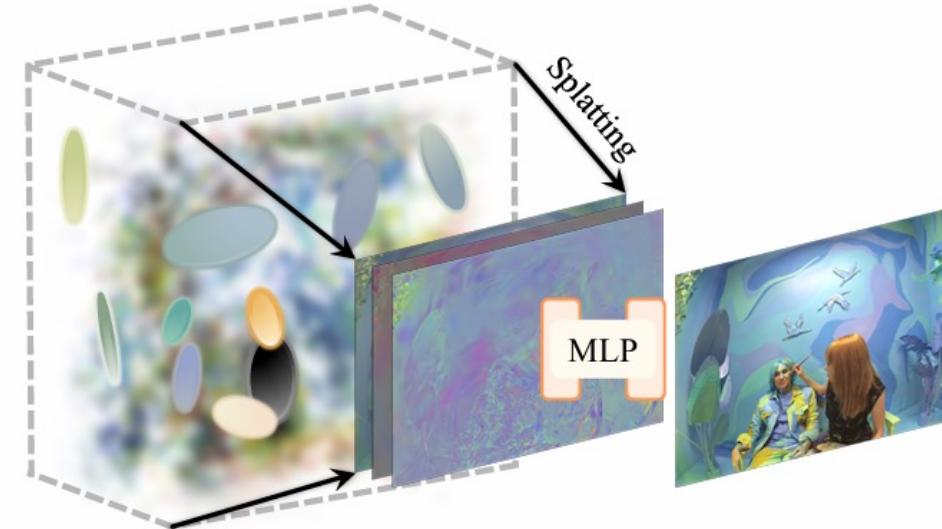


- **Input:** multi-view video
- **Output:** per-frame gaussian

Deformation Fields | Gaussians | Spacetime Gaussian



(a) Spacetime Gaussians



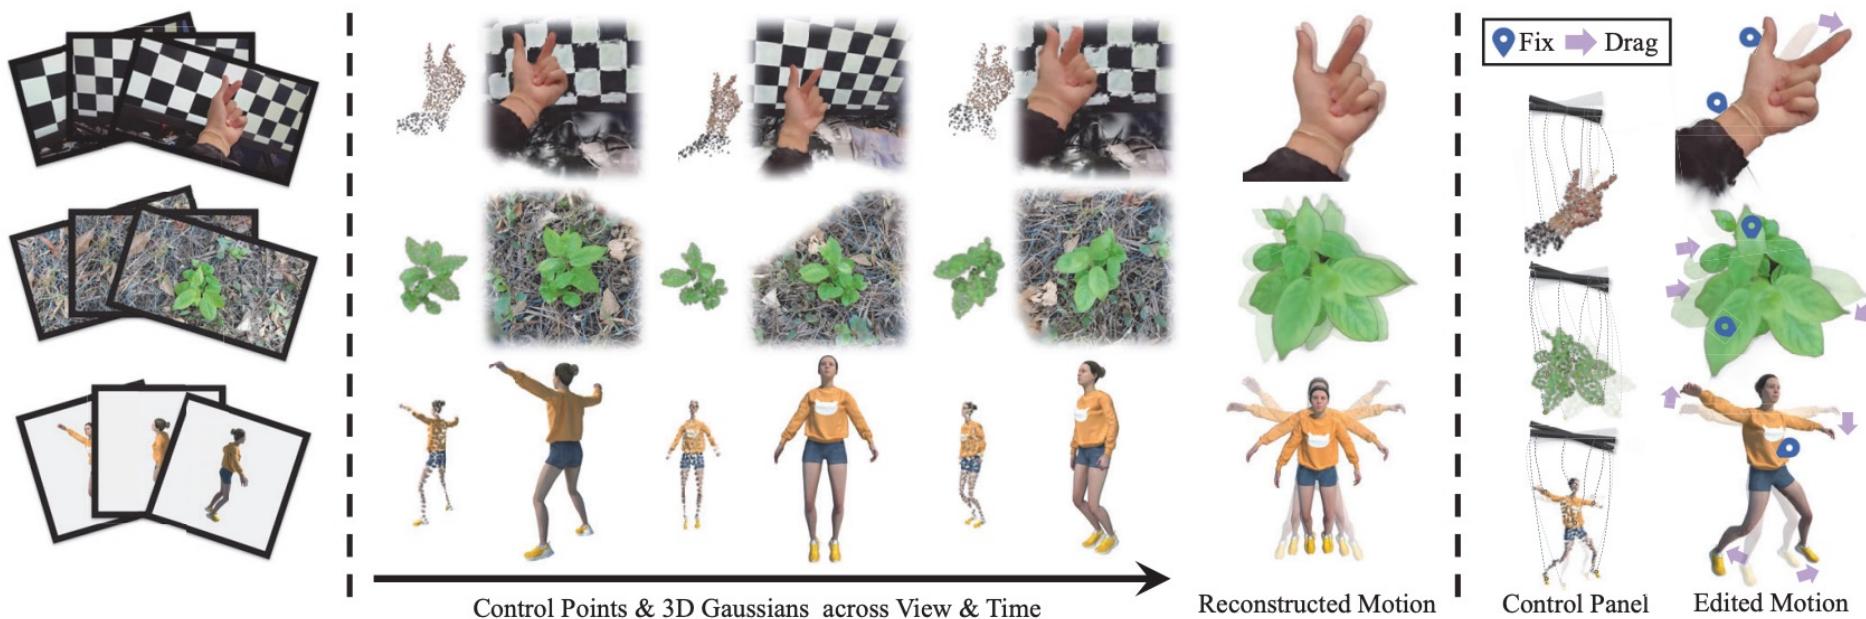
(b) Feature Splatting and Rendering

- Motion trajectory modeling
- Feature map rendering
- RGB image decoding

SC-GS: Sparse-Controlled Gaussian Splatting for Editable Dynamic Scenes

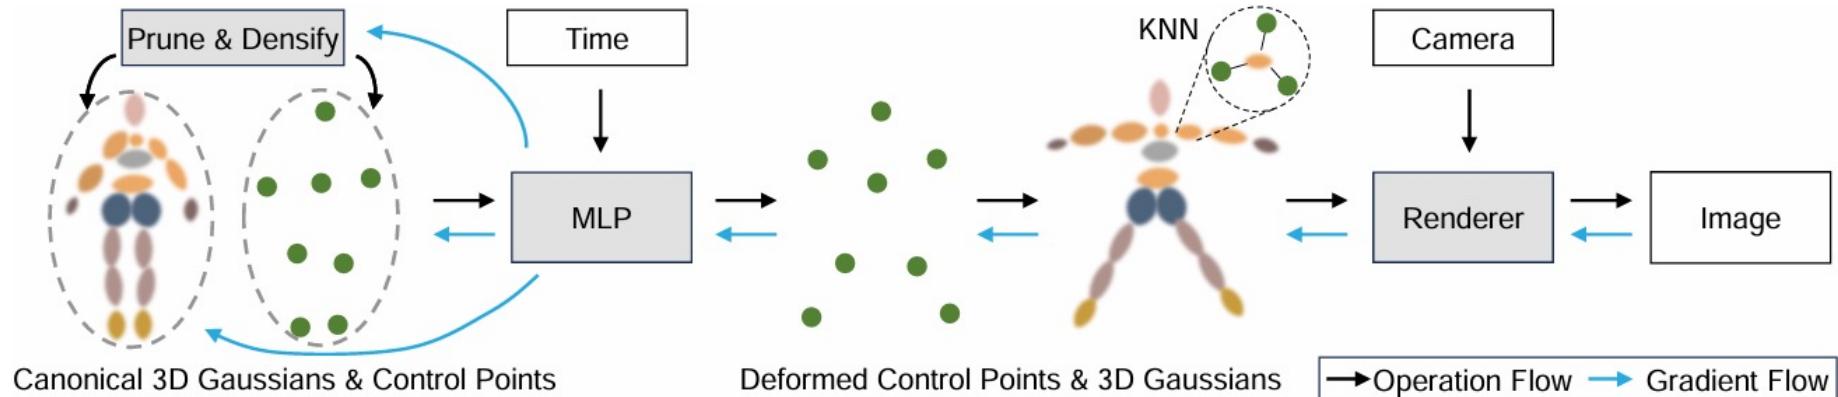
Yi-Hua Huang^{1*} Yang-Tian Sun^{1*} Ziyi Yang^{3*} Xiaoyang Lyu¹ Yan-Pei Cao^{2†} Xiaojuan Qi^{1†}

¹ The University of Hong Kong ² VAST ³ Zhejiang University



- **Input:** monocular video
- **Output:** dynamic gaussian

Deformation Fields | Gaussians | SC-GS



- Sparse control points for modeling scene motion
- Rigid constraints during optimization
- Allow motion editing benefiting the learned control points

Deformation Field: $\Psi : (p_i, t) \rightarrow (R_i^t, T_i^t).$ $\mu_j^t = \sum_{k \in \mathcal{N}_j} w_{jk} (R_k^t(\mu_j - p_k) + p_k + T_k^t),$

$$q_j^t = (\sum_{k \in \mathcal{N}_j} w_{jk} r_k^t) \otimes q_j,$$

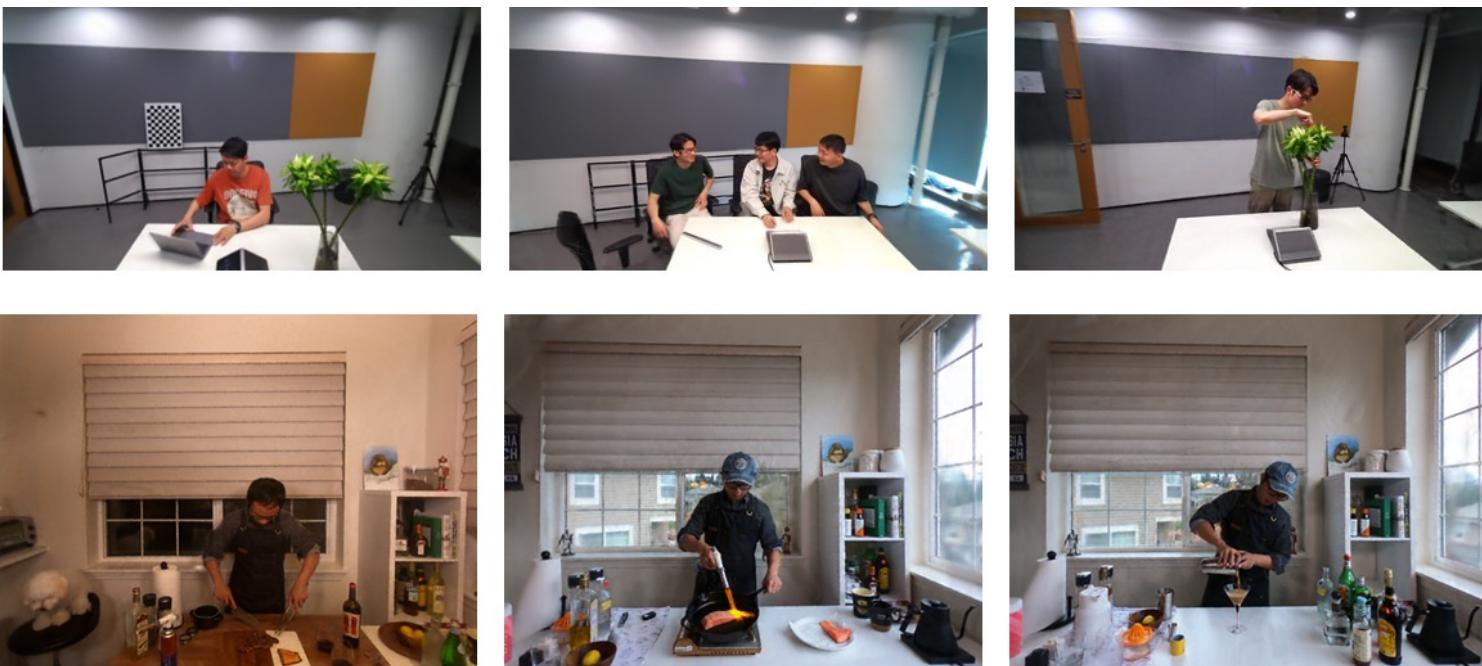
Per-frame Reconstruction

- **NeRF Related:** Streaming NeRF, 4K4D
- **3DGS Related:** Dynamic 3DGS, 3DGStream
- **Feed Forward:** ENeRF, GPS–Gaussian

Streaming Radiance Fields for 3D Video Synthesis

Lingzhi Li, Zhen Shen, Zhongshu Wang, Li Shen, Ping Tan

Alibaba Group



- **Input:** multi-view video
- **Output:** compressed per-frame Voxels

Table 1: **Comparison with related works on the N3DV and Meet Room datasets.** We report *per frame* PSNR, training and inference time and storage cost in average for a clear comparison. All the results are recorded with our 3090 GPU except the results of N3DV are referred to the numbers in the original paper. Compared to all the baselines, our method can achieve remarkable speedup on training time with competitive performance on the other metrics.

Methods	N3DV				Meet Room			
	PSNR (dB)	Train (minutes)	Inference (seconds)	Storage (MB)	PSNR (dB)	Train (minutes)	Inference (seconds)	Storage (MB)
JaxNeRF*[13]	28.53	485	67	14	27.11	473	40	14
LLFF*[18]	23.23	8	0.004	3192	22.88	3	0.003	1500
N3DV [10]	29.58	260	>67	0.1	-	-	-	-
Plenoxels*[11]	28.68	23	0.12	4106	27.15	14	0.1	1015
Ours	28.26	0.25	0.12	17.7/31.4	26.72	0.17	0.1	5.7/9.0

* Denote training from scratch per frame.

Streaming Voxel Reconstruction: $\mathbf{V}^i = \mathbf{V}^{i-1} + \delta_i = \mathbf{V}^0 + \sum_{j=1}^i \delta_j.$

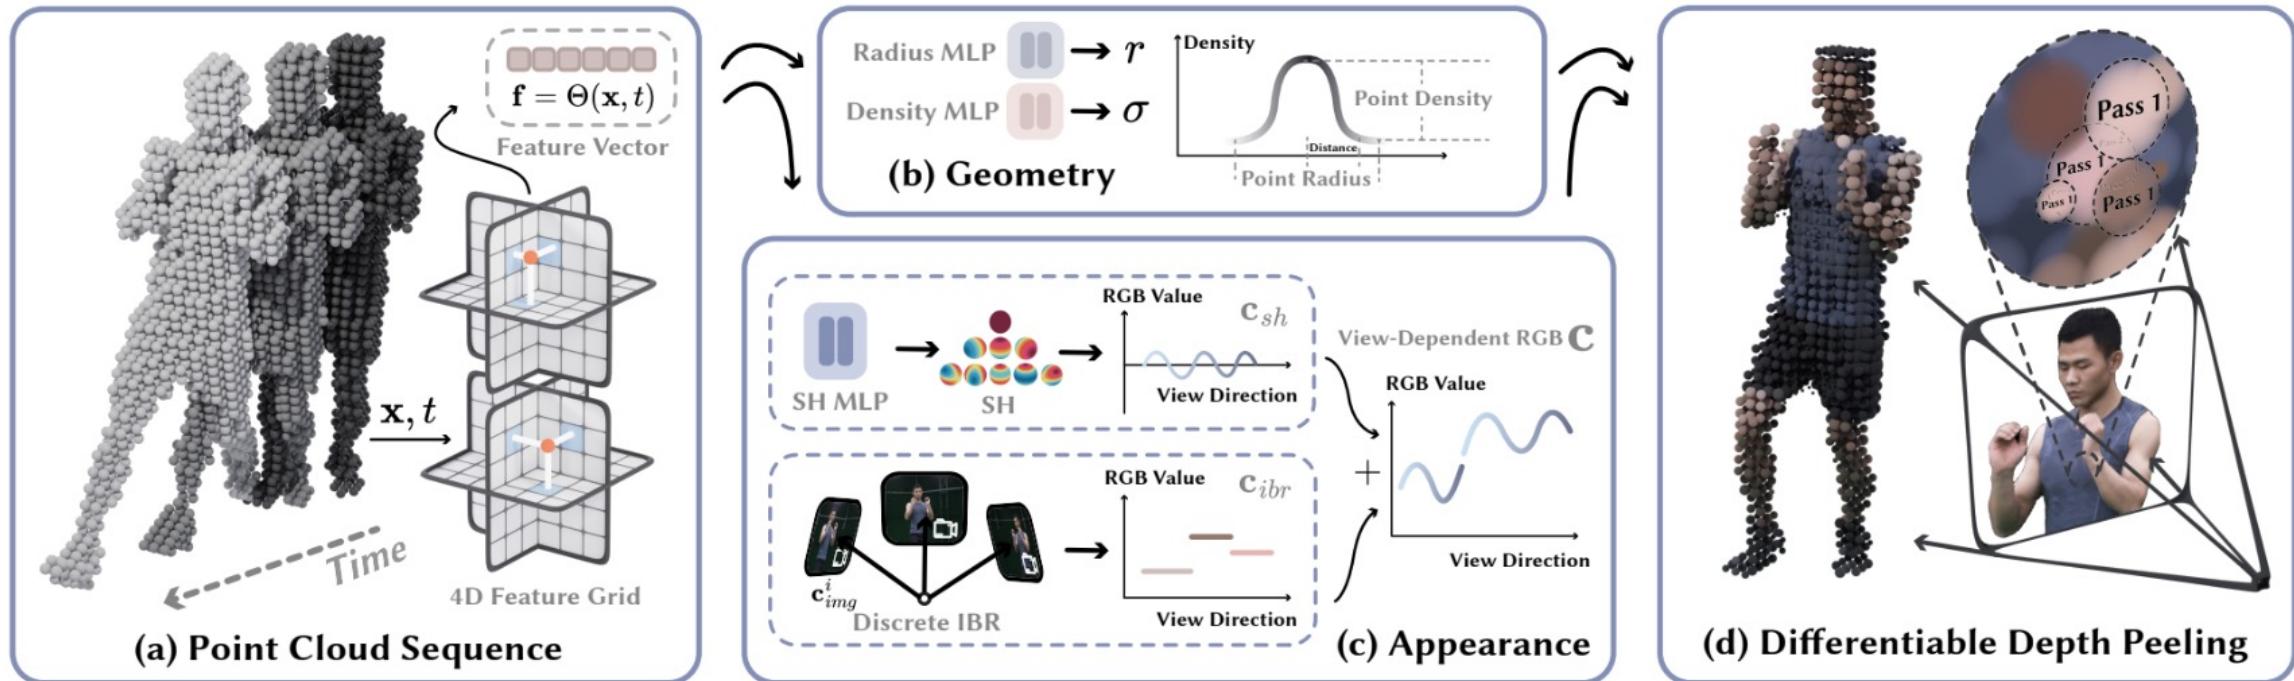
4K4D: Real-Time 4D View Synthesis at 4K Resolution

Zhen Xu¹ Sida Peng¹ Haotong Lin¹ Guangzhao He¹
Jiaming Sun² Yujun Shen³ Hujun Bao¹ Xiaowei Zhou¹



- **Input:** multi-view video
- **Output:** per-frame point cloud representation

Per-frame | Points | 4K4D



- Point cloud from space carving
- Geometry and appearance info decoded from feature grid
- Differentiable depth peeling

Dynamic 3D Gaussians: Tracking by Persistent Dynamic View Synthesis

Jonathon Luiten^{1,2} Georgios Kopanas³ Bastian Leibe² Deva Ramanan¹

¹Carnegie Mellon University, USA ²RWTH Aachen University, Germany ³Inria & Universite Cote d'Azur, France

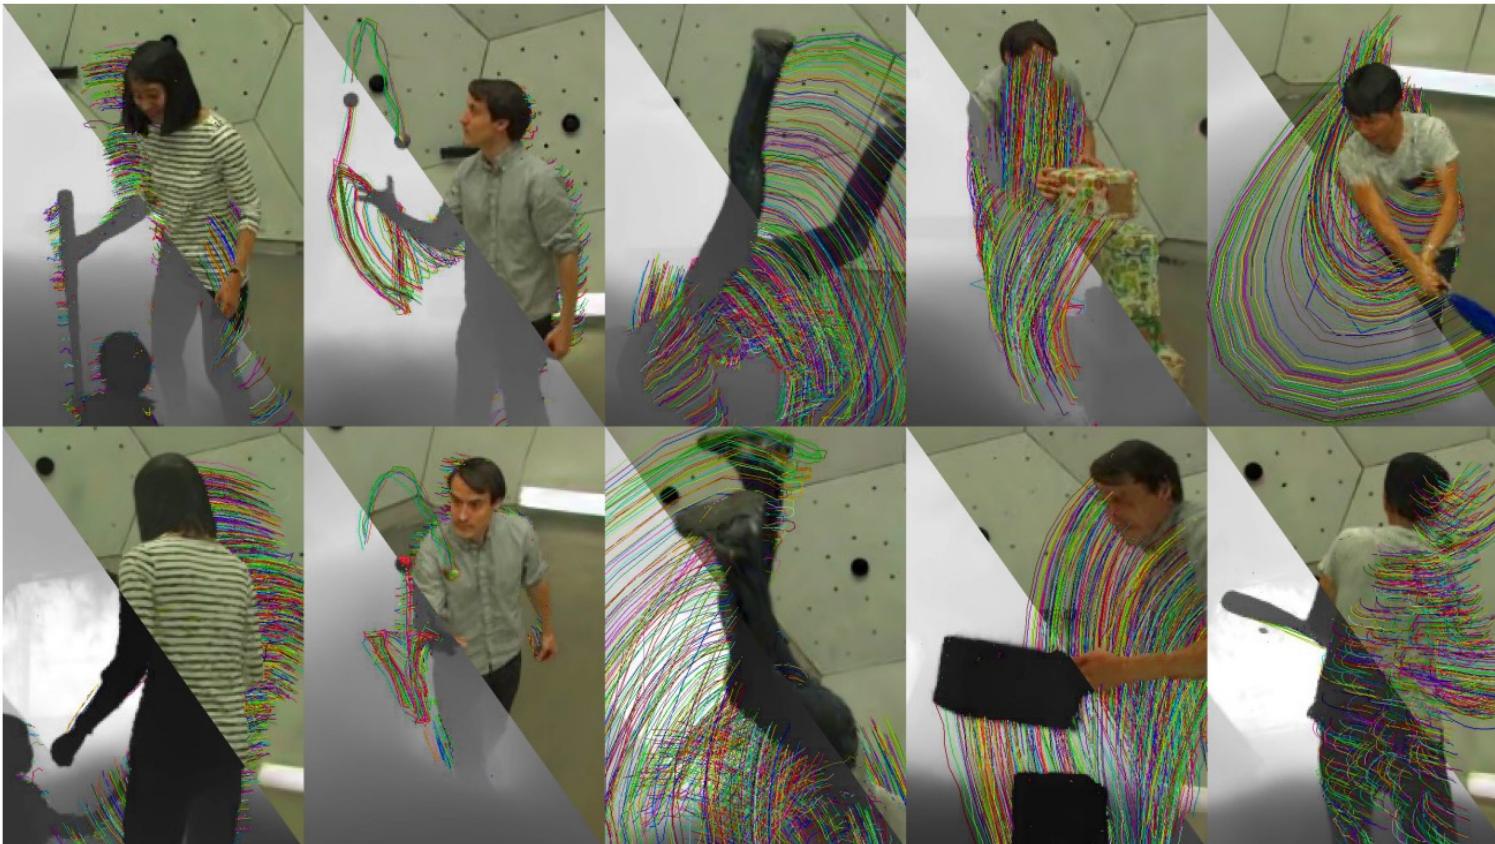




Figure 2. **Gaussian Centers.** Point-cloud of colored centers in contiguous timesteps, showing how they model scene geometry and move over time. [Videos]

- Multi–view video input
- Per–frame optimization based on previous frame
- Local rigidity loss

Method	PSNR↑	SSIM↑	LPIPS↓
TiNeuVox-S [9]	26.64	0.92	0.14
TiNeuVox [9]	27.28	0.91	0.13
InstantNGP [26]	24.69	0.91	0.12
Particle-NeRF [1]	27.47	0.94	0.08
Ours	39.49	0.99	0.02

Table 2. **Result on the Particle-NeRF dataset.** See text for details on the dataset, metrics, tasks and methods.

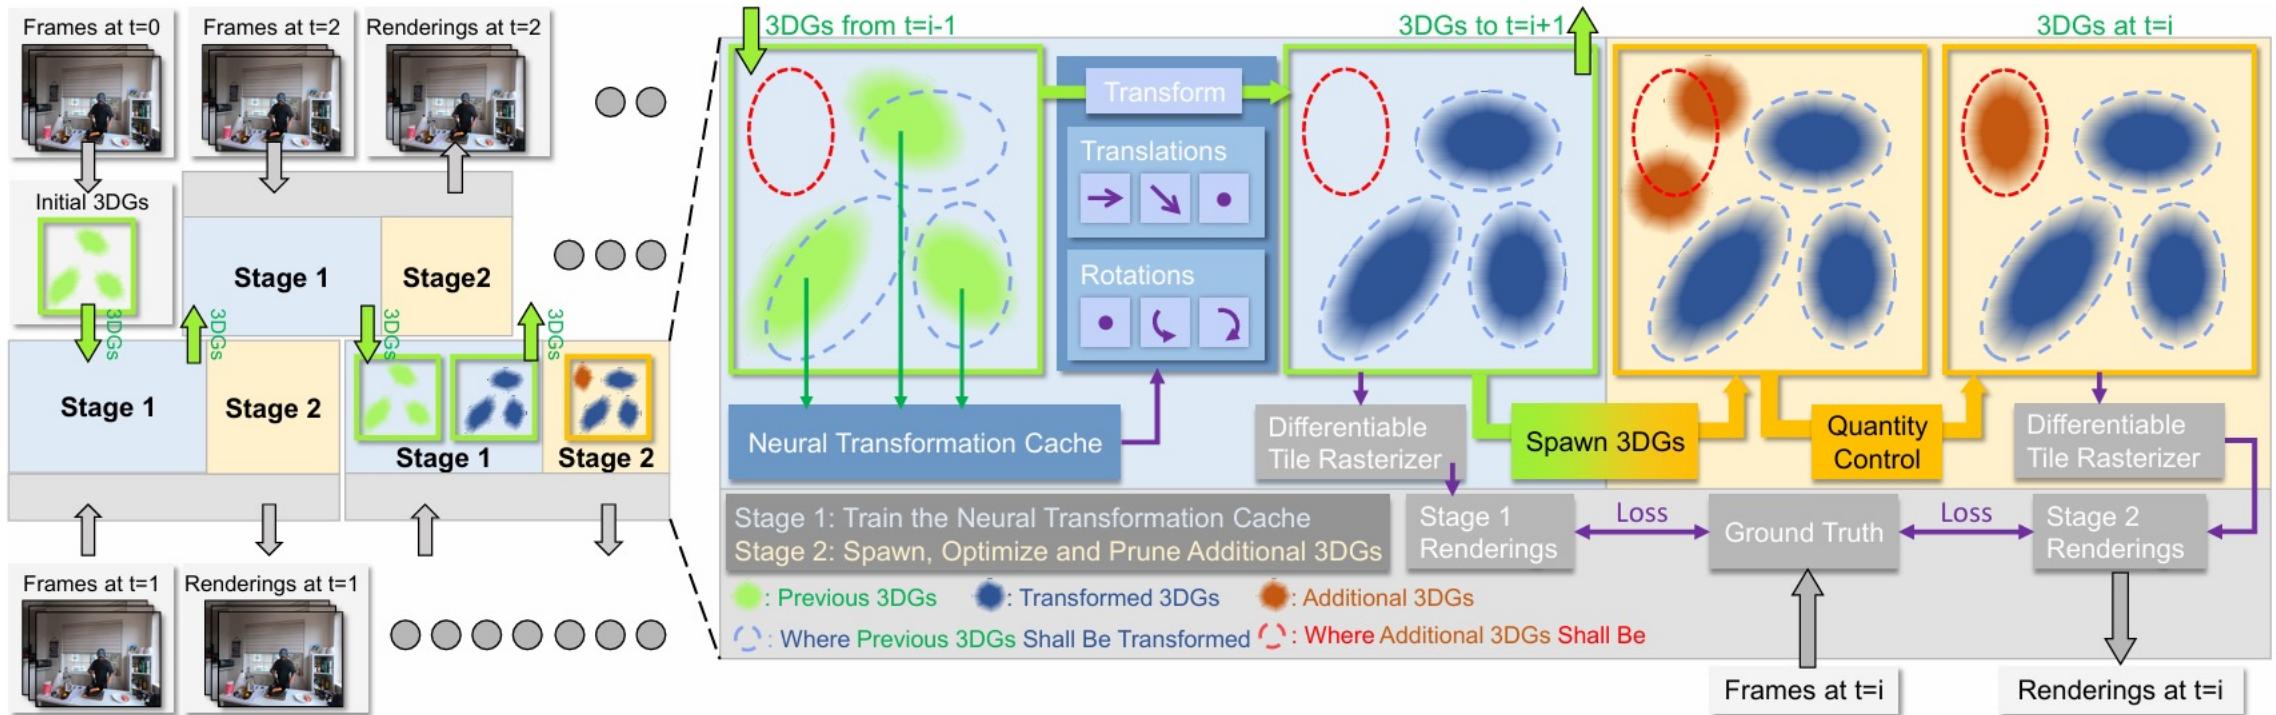
3DGStream: On-the-Fly Training of 3D Gaussians for Efficient Streaming of Photo-Realistic Free-Viewpoint Videos

Jiakai Sun, Han Jiao, Guangyuan Li, Zhanjie Zhang, Lei Zhao*, Wei Xing*
Zhejiang University



- **Input:** multi-view video
- **Output:** per-frame gaussian

Per-frame | Gaussians | 3DGStream



- Storing per-frame deformation rather than per-frame 3DGs
- Spawn, optimize, and prune additional 3DGs in the new frame

Efficient Neural Radiance Fields for Interactive Free-viewpoint Video

Haotong Lin*

Sida Peng*

haotongl@zju.edu.cn

pengsida@zju.edu.cn

State Key Laboratory of CAD&CG,
Zhejiang University
China

Zhen Xu

Yunzhi Yan

Qing Shuai

zhenx@zju.edu.cn

yanyz@zju.edu.cn

s_q@zju.edu.cn

Zhejiang University, China

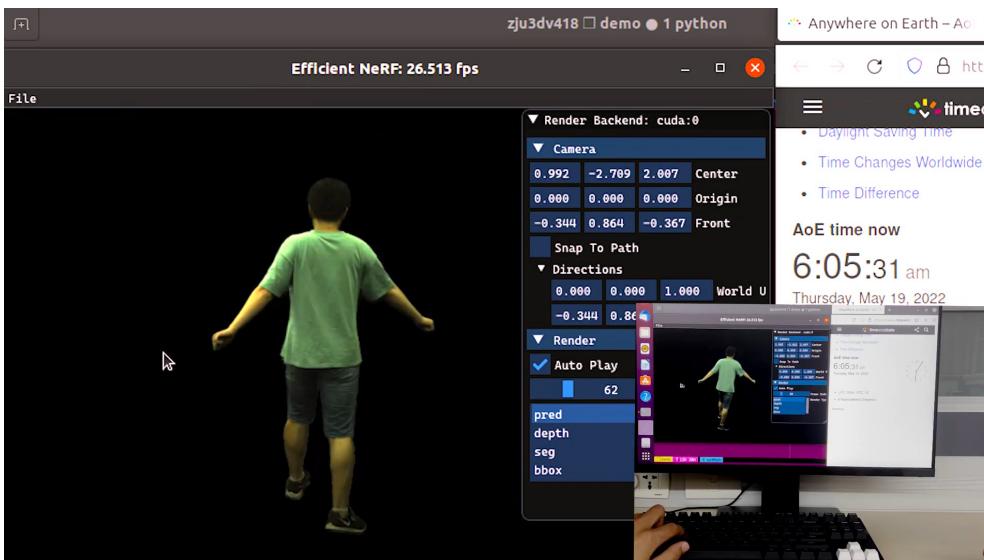
Hujun Bao

Xiaowei Zhou†

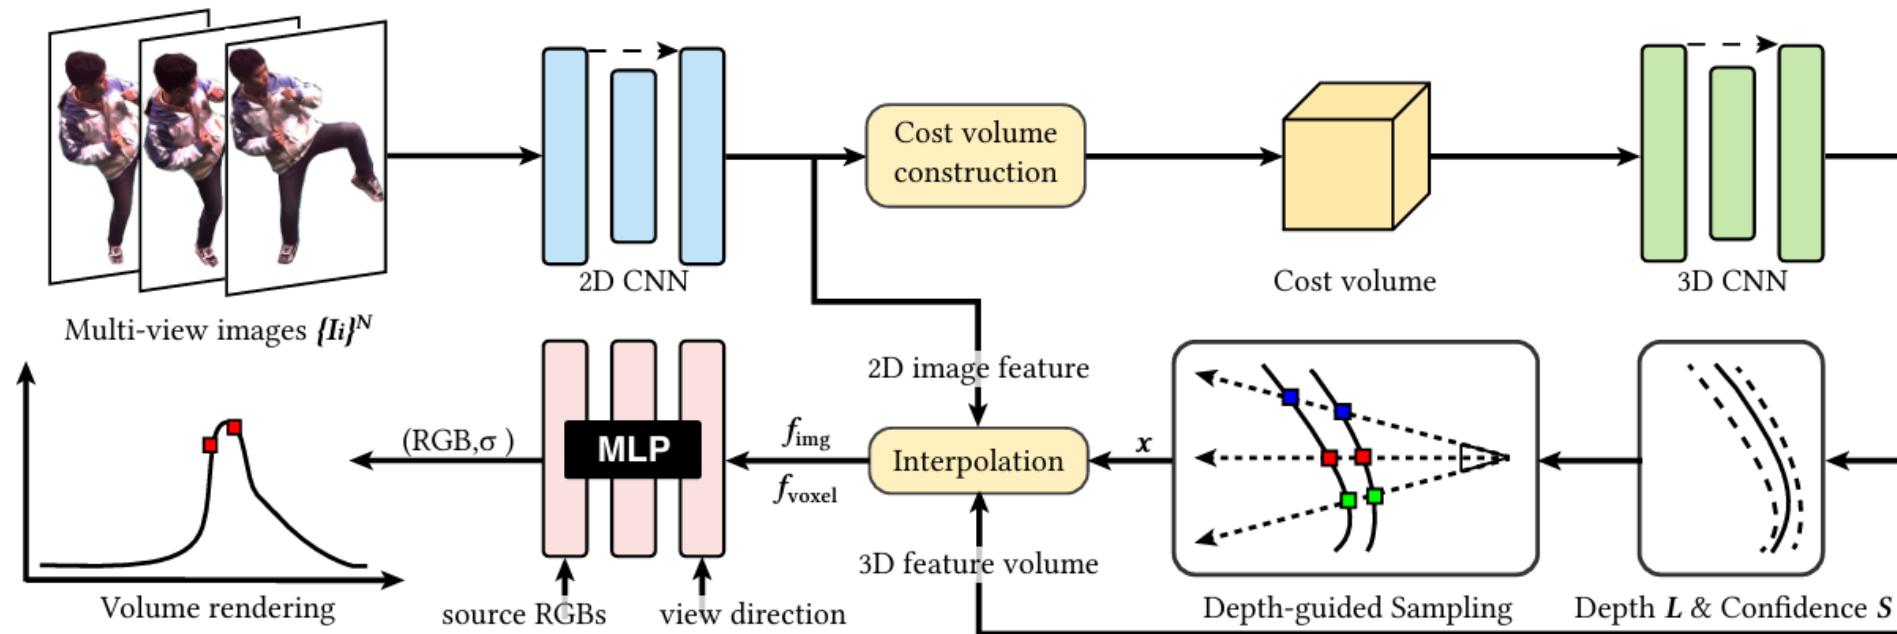
bao@cad.zju.edu.cn

xwzhou@zju.edu.cn

State Key Laboratory of CAD&CG,
Zhejiang University
China



Per-frame | Feed Forward | ENeRF



- Per-frame Generalized NeRF (e.g., MVSNet + NeRF)
- Fast and efficient sampling via MVS cost volume

GPS-Gaussian: Generalizable Pixel-wise 3D Gaussian Splatting for Real-time Human Novel View Synthesis

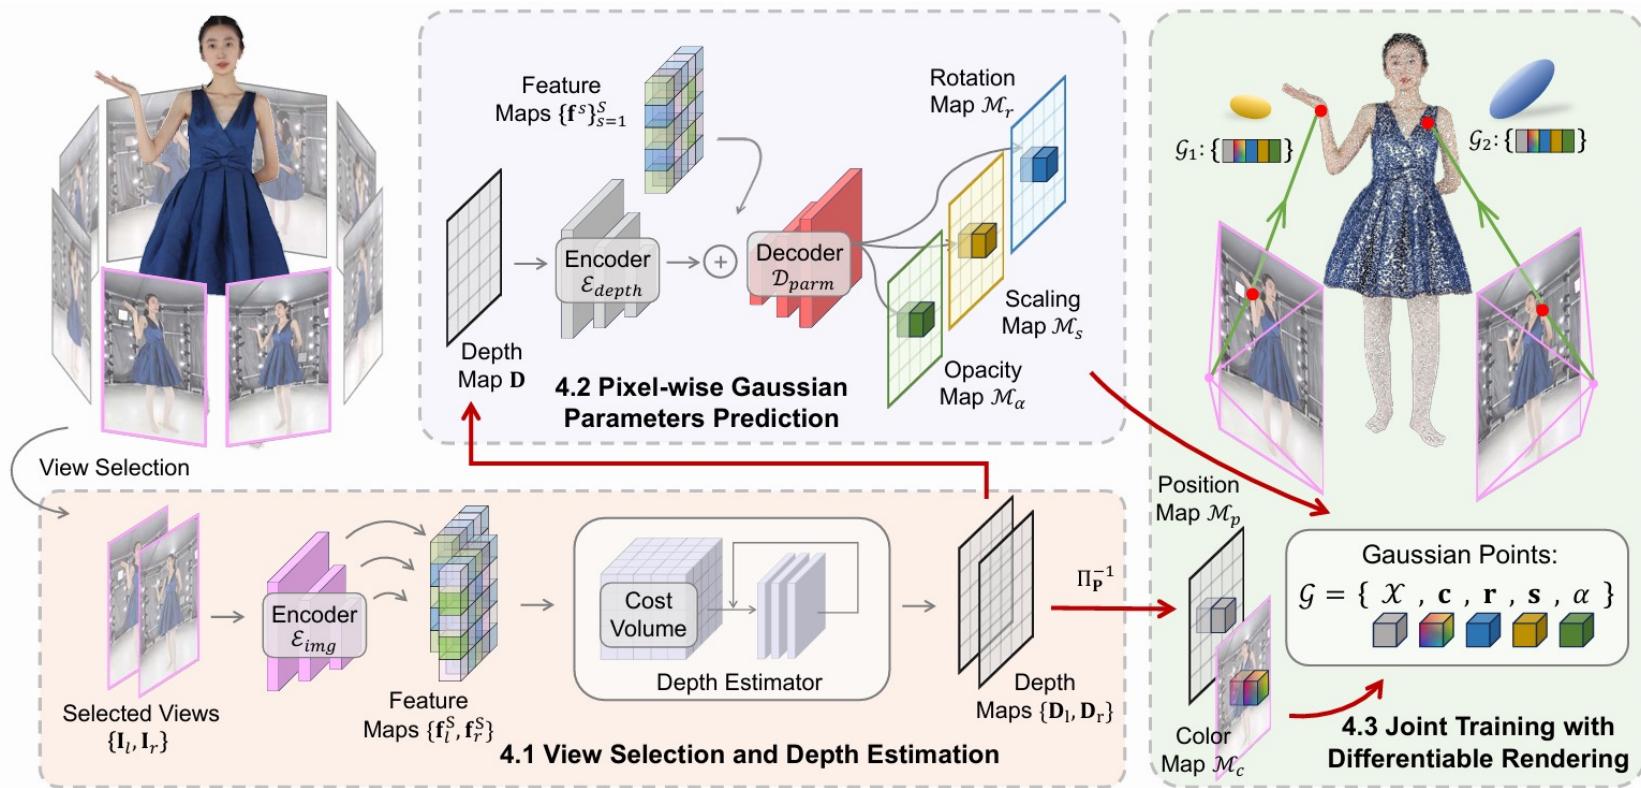
Shunyuan Zheng^{†,1}, Boyao Zhou², Ruizhi Shao², Boning Liu², Shengping Zhang^{*,1,3},
Liqiang Nie¹, Yebin Liu²

¹Harbin Institute of Technology ²Tsinghua University ³Peng Cheng Laboratory



- **Input:** multi-view video
- **Output:** per-frame gaussian

Per-frame | Feed Forward | GPS–Gaussian



- Per-frame Generalized 3DGS
- Learned MVS and features as 3DGS initialization
- Real-time 4D reconstruction

- **Depth and Flow:** Robust NeRF, DynlBaR, MoDGS, Shape of Motion
- **Dynamic Physics:** NVFi, PAC–NeRF, DeformGS
- **Parametric Human:** Animatable NeRF, GaussianAvatars, Animatable Gaussians
- **(TBD) Large Vision Model:** Cat3D/ReconX, Consistent4D/STAG4D/SV4D

Robust Dynamic Radiance Fields

Yu-Lun Liu^{2*} Chen Gao¹ Andreas Meuleman^{3*} Hung-Yu Tseng¹ Ayush Saraf¹

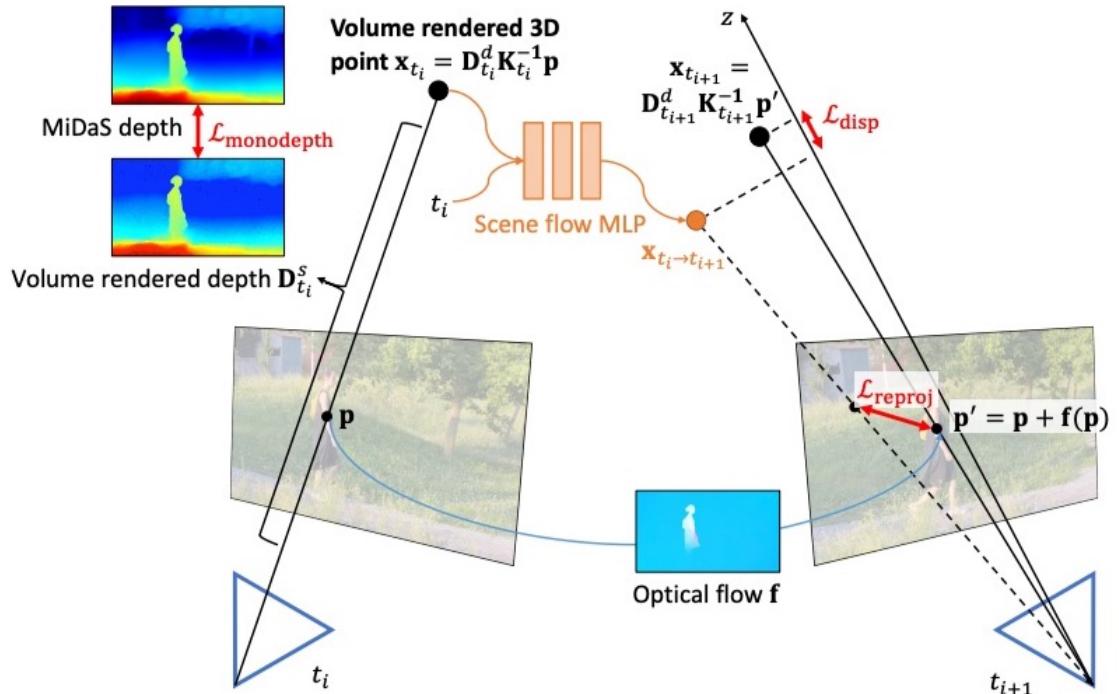
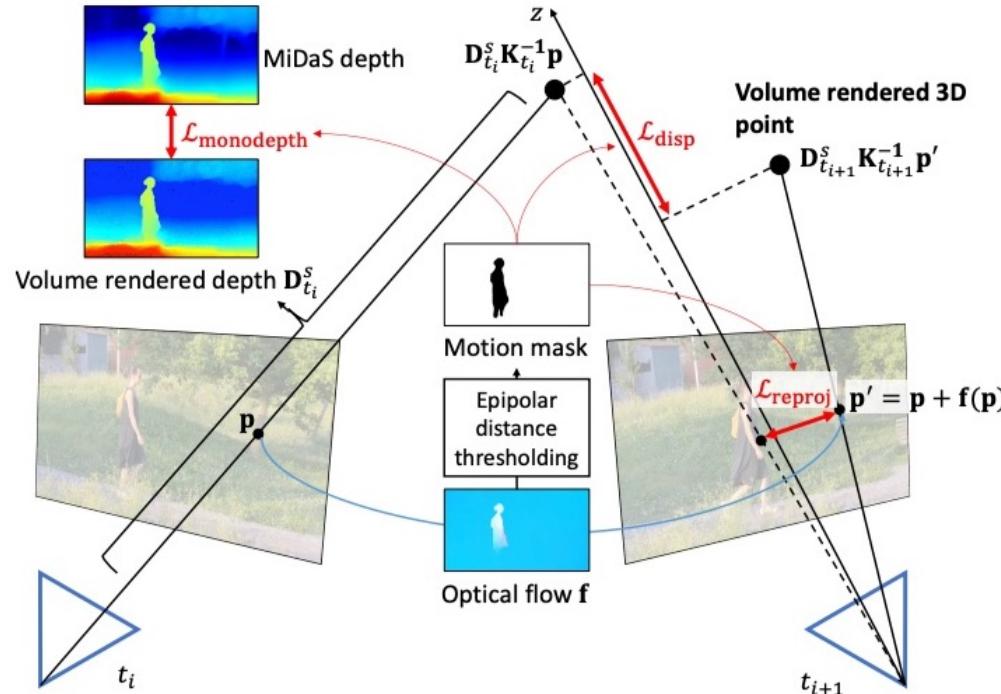
Changil Kim¹ Yung-Yu Chuang² Johannes Kopf¹ Jia-Bin Huang^{1,4}

¹Meta ²National Taiwan University ³KAIST ⁴University of Maryland, College Park

<https://robust-dynrf.github.io/>



Priors | Depth & Flow | RoDynRF



- MiDaS for Depth Supervision
- Guide the scene flow of reproject volume render point with RAFT

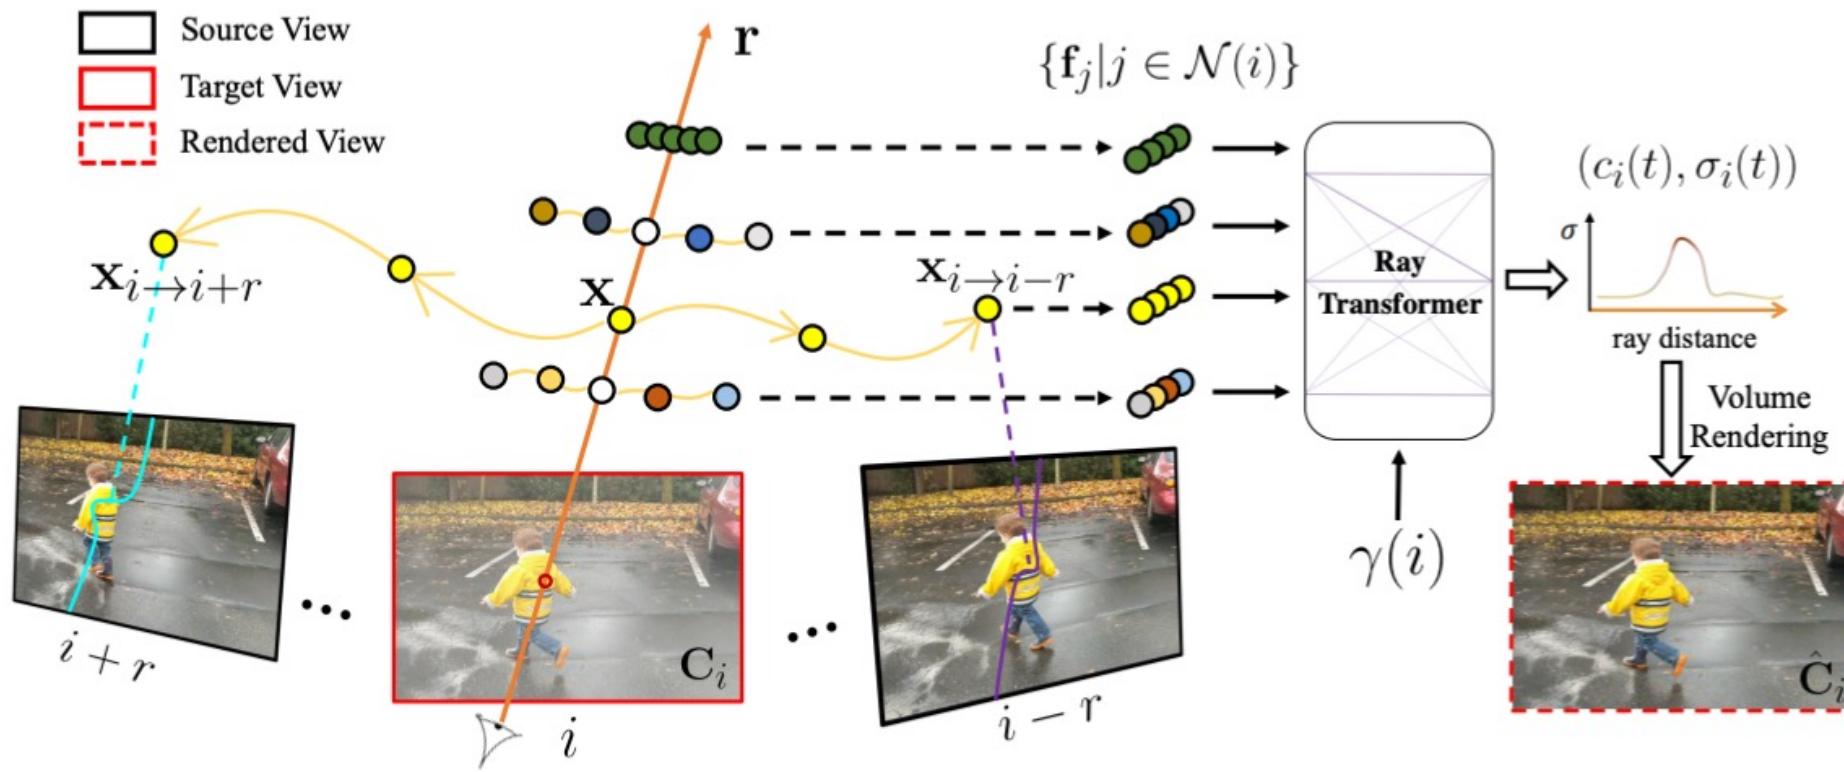
DynIBaR: Neural Dynamic Image-Based Rendering

Zhengqi Li¹, Qianqian Wang^{1,2}, Forrester Cole¹, Richard Tucker¹, Noah Snavely¹

¹Google Research ²Cornell Tech



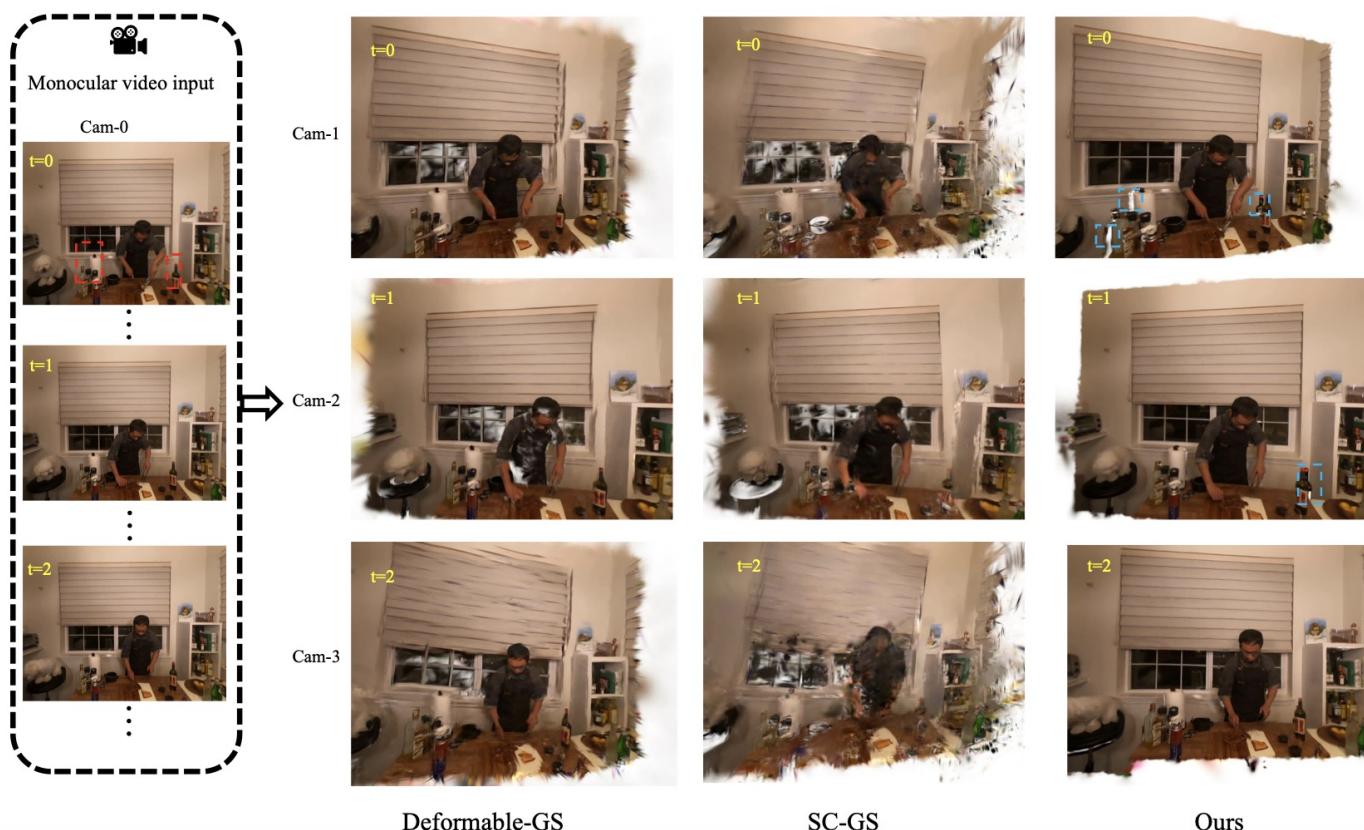
Priors | Depth & Flow | DynIBaR

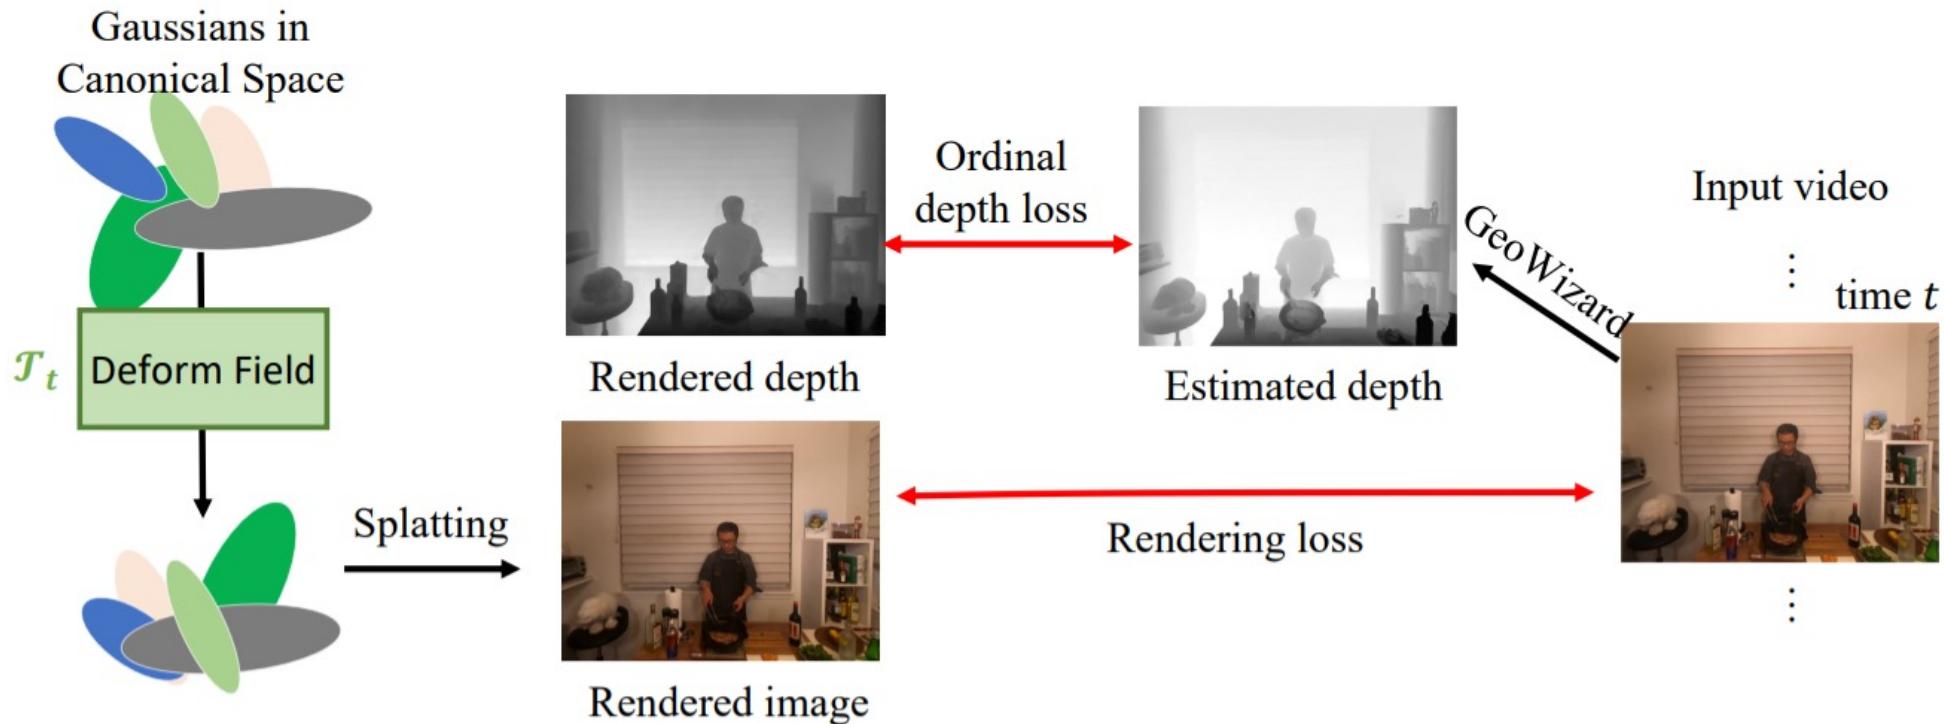


- Aggregate multi-view image features for dynamic NeRF rendering
- Leveraging optical flow consistency priors from RAFT

MoDGS: Dynamic Gaussian Splatting from Causually-captured Monocular Videos

Qingming Liu^{1*}, Yuan Liu^{2*}, Jiepeng Wang², Xianqiang Lv¹,
Peng Wang², Wenping Wang³, Junhui Hou^{1†},





- Computing 3D flow by lifting 2D flow (RAFT) to 3D with mono depth (GeoWizard)
- Gaussian (deformation) initialization with 3D flow
- Resolve depth inconsistency between frames

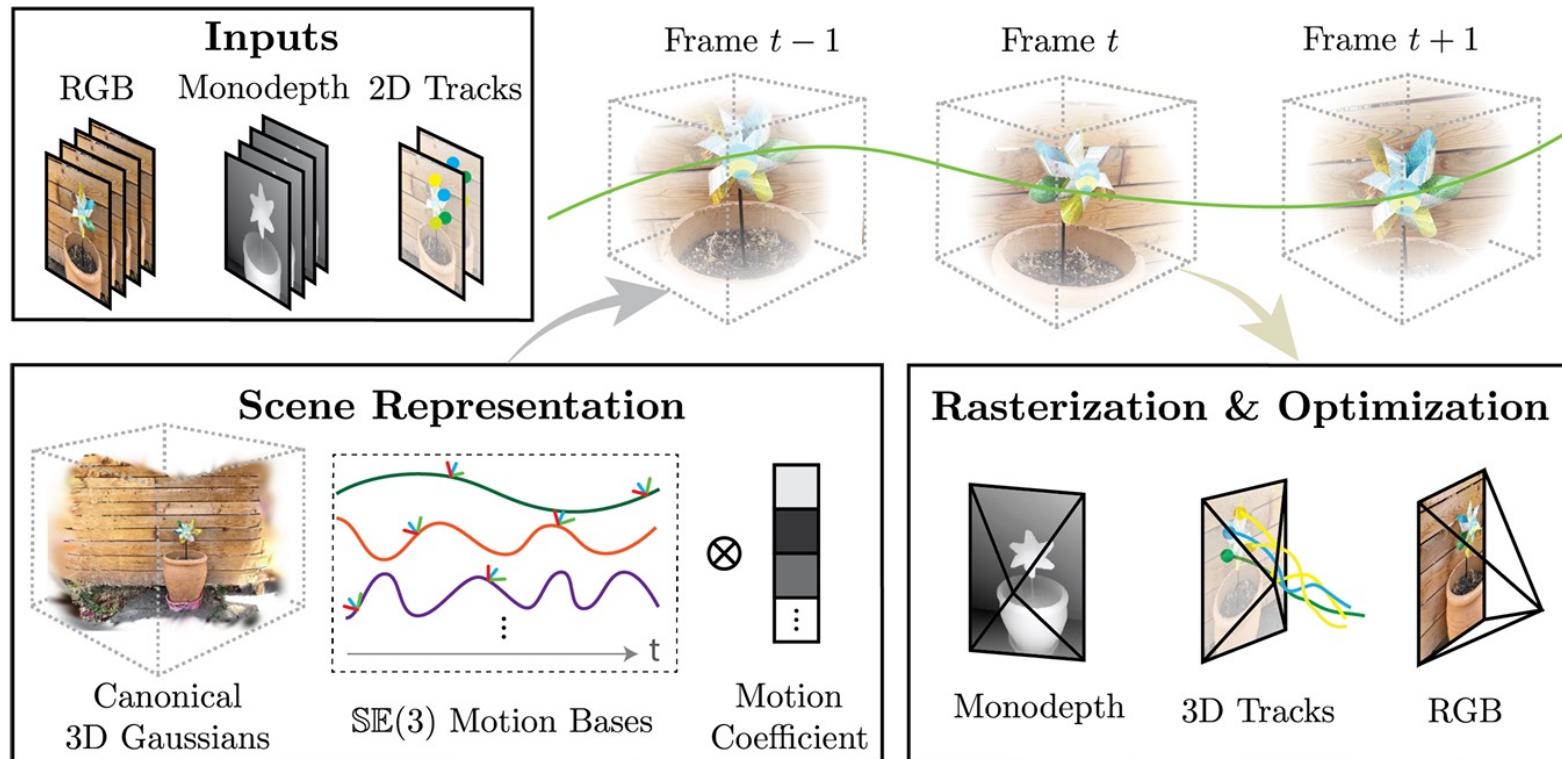
Shape of Motion: 4D Reconstruction from a Single Video

Qianqian Wang^{1,2*}, Vickie Ye^{1*}, Hang Gao^{1*},
Jake Austin¹, Zhengqi Li², and Angjoo Kanazawa¹

¹ UC Berkeley ² Google Research



Priors | Depth & Flow | Shape of Motion



- Explicit trajectory motion modeling
- Foreground mask by Track–Anything
- Depth prior by Depth–Anything
- Tracking prior by TAPIR

NVFi: Neural Velocity Fields for 3D Physics Learning from Dynamic Videos

Jinxi Li Ziyang Song Bo Yang

vLAR Group, The Hong Kong Polytechnic University



T-NeRF



D-NeRF



TiNeuVox

Extrapolation

Observed View



T-NeRF_{PINN}



HexPlane_{PINN}

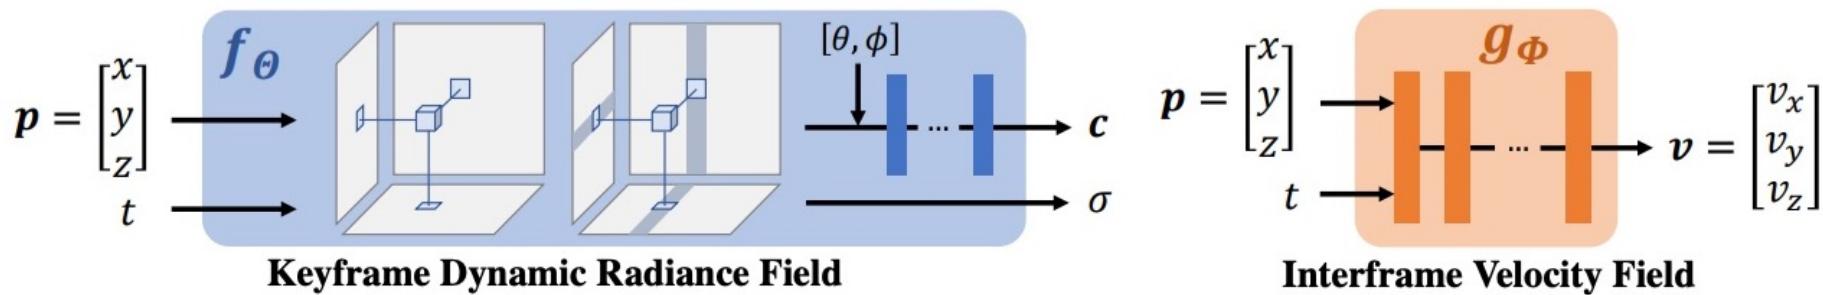


NVFi (Ours)

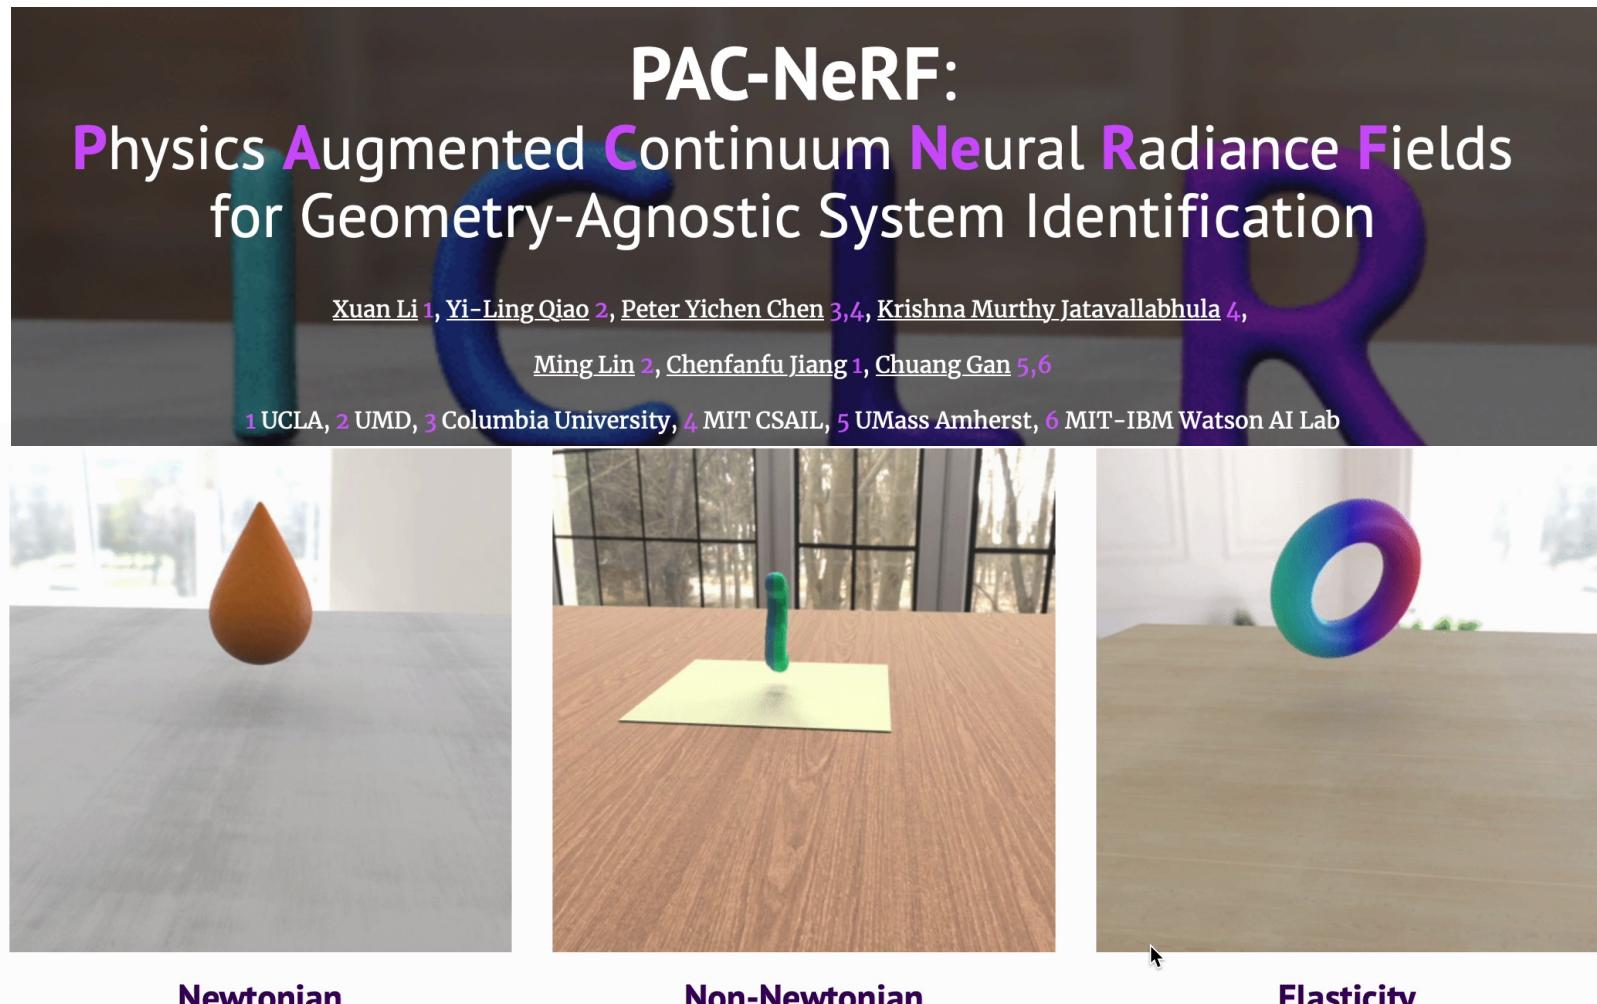


GT

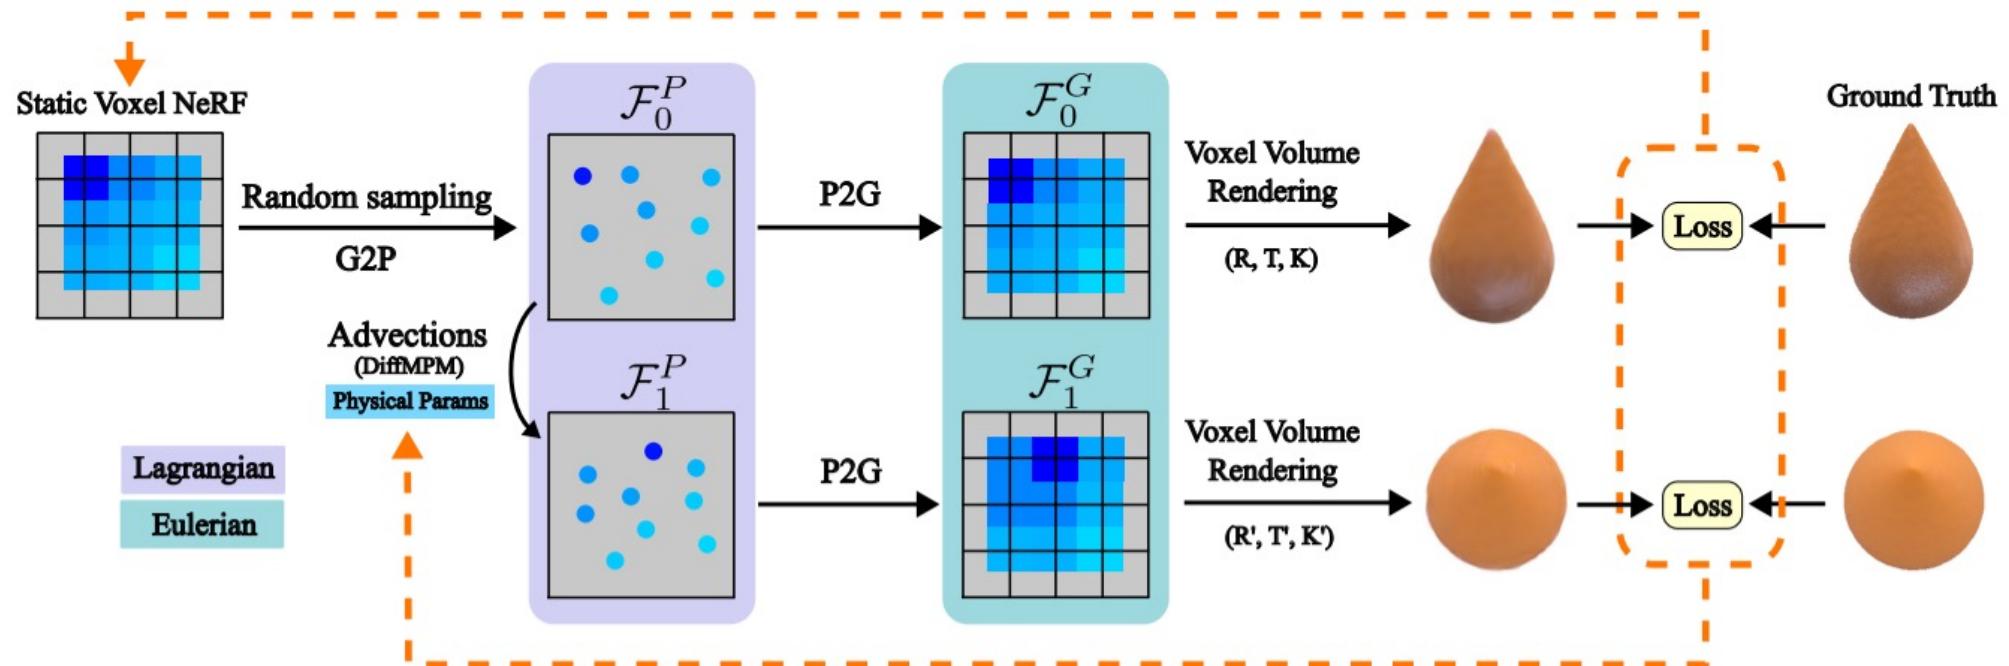
- **Input:** multi-view video
- **Output:** 4D NeRF & Physical Velocity



- Physics-informed radiance fields (velocity field modeling)
- Joint optimization of radiance fields and velocity field



- **Input:** multi-view video
- **Output:** static NeRF with physical properties



- Conservation laws of continuum mechanics
- Optimizing velocity field for continuum materials

DeformGS: Scene Flow in Highly Deformable Scenes for Deformable Object Manipulation

Bardienus P. Duisterhof¹, Mandi Zhao², Yunchao Yao¹, Jia-Wei Liu⁴,
Jenny Seidenschwarz^{1,5}, Mike Zheng Shou⁴, Deva Ramanan¹, Shuran Song²,
Stan Birchfield³, Bowen Wen³, and Jeffrey Ichnowski¹

¹ Carnegie Mellon University, The Robotics Institute {bduister,jeffi}@cmu.edu

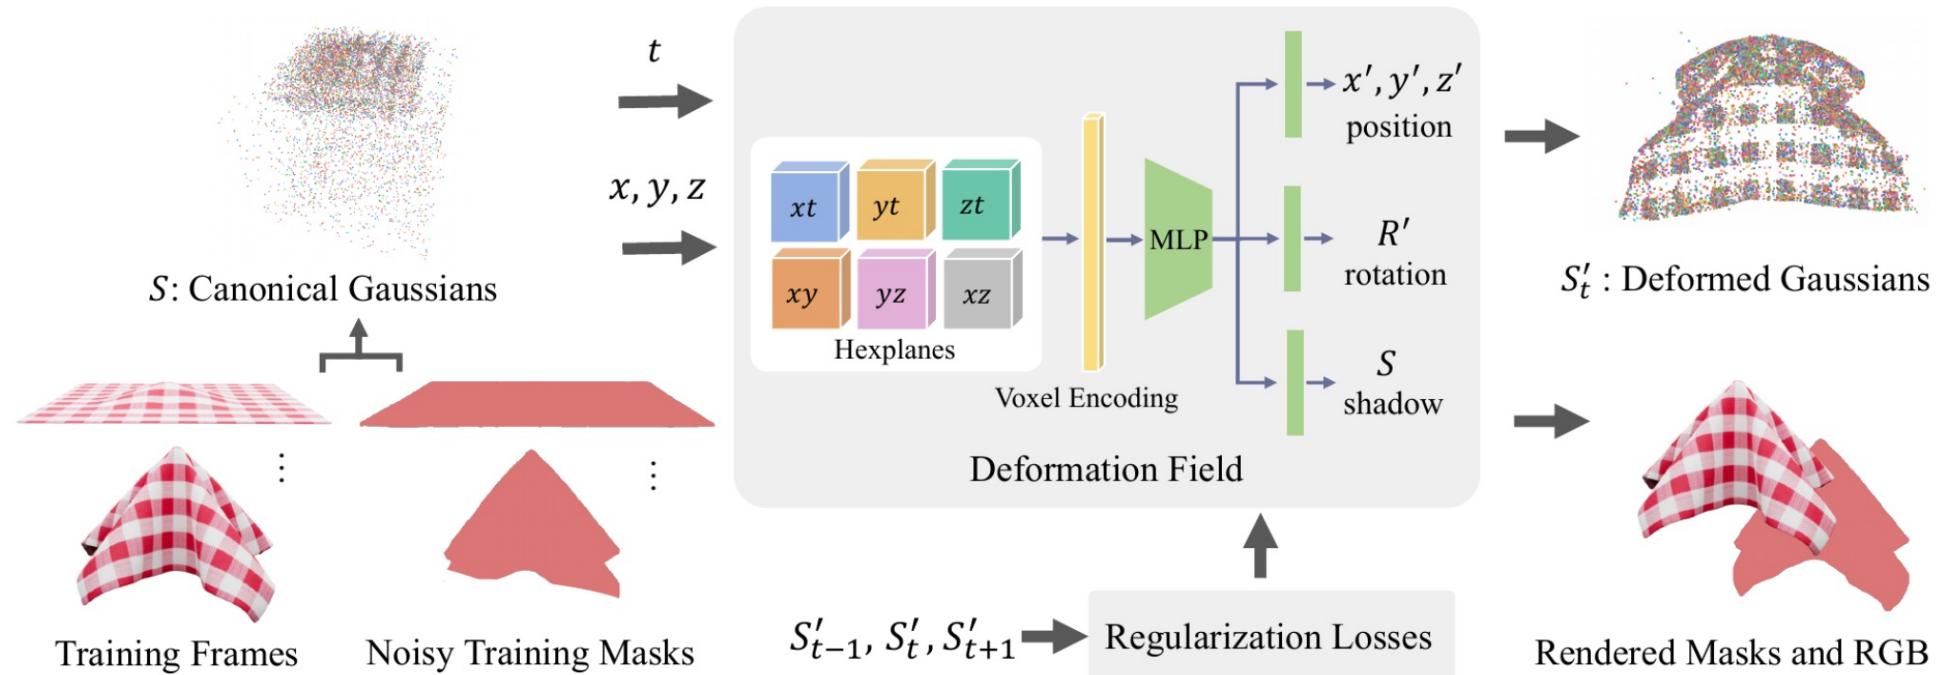
² Stanford University

³ NVIDIA

⁴ National University of Singapore

⁵ Technical University of Munich

- **Input:** multi-view video
- **Output:** dynamic gaussian



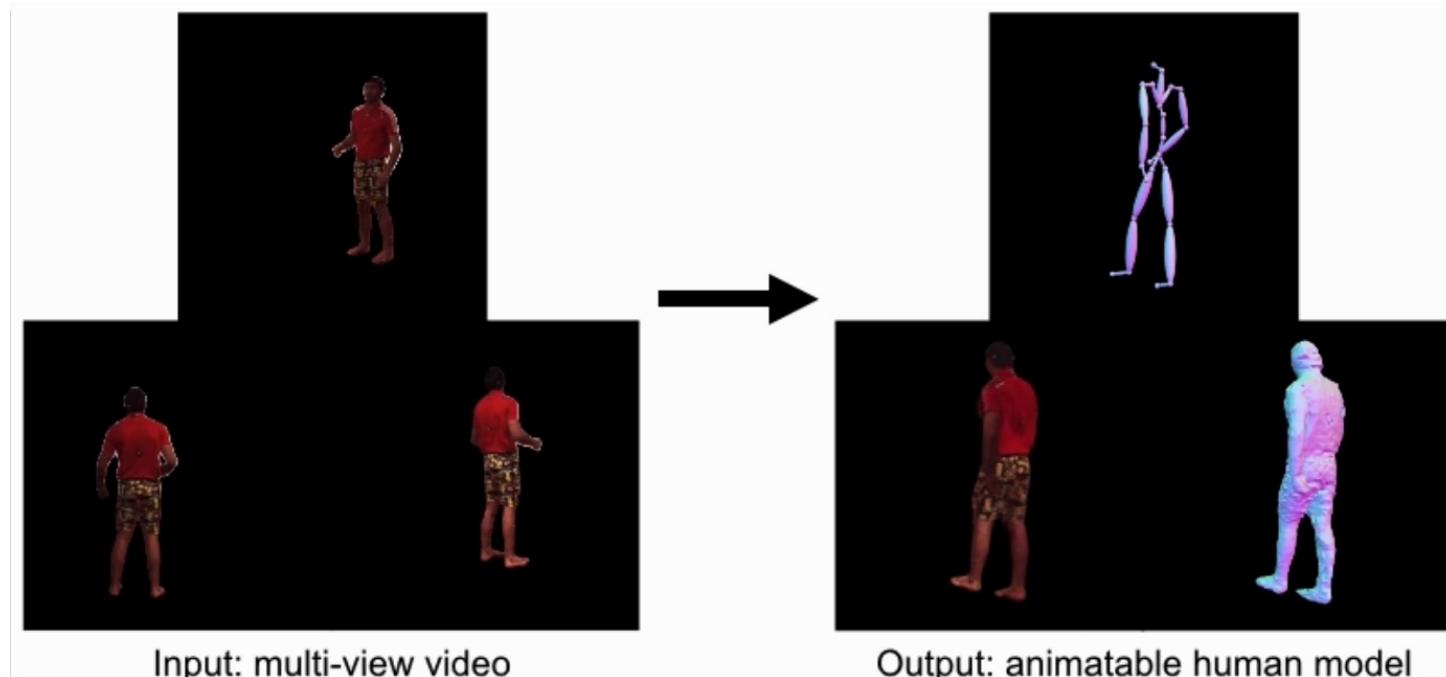
- Momentum conservation loss for gaussian position
- Local distance loss for adjacent gaussian

Animatable Neural Radiance Fields for Modeling Dynamic Human Bodies

Sida Peng^{1*} Junting Dong^{1*} Qianqian Wang² Shangzhan Zhang¹

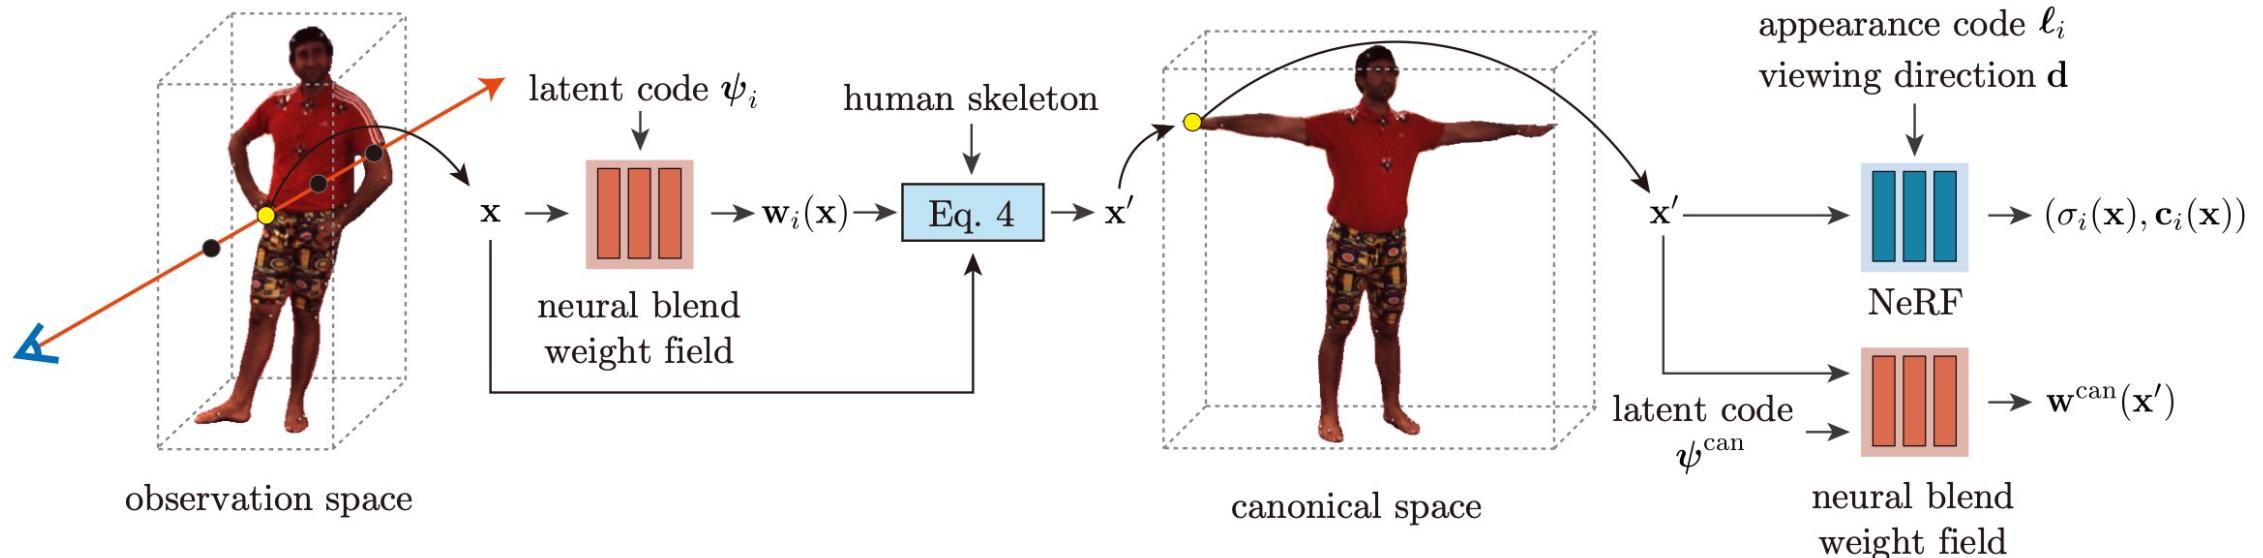
Qing Shuai¹ Xiaowei Zhou¹ Hujun Bao^{1†}

¹Zhejiang University ²Cornell University



- **Input:** multi–view human centric video
- **Output:** animatable human model

Priors | Human | Animatable NeRF



- Skeleton–driven deformation modeling
- Warping field from Linear Blend Skinning (LBS)
- Animatable radiance field

GaussianAvatars: Photorealistic Head Avatars with Rigged 3D Gaussians

Shenhan Qian¹ Tobias Kirschstein¹ Liam Schoneveld² Davide Davoli³ Simon Giebenhain¹ Matthias Nießner¹

¹Technical University of Munich



²Woven by Toyota

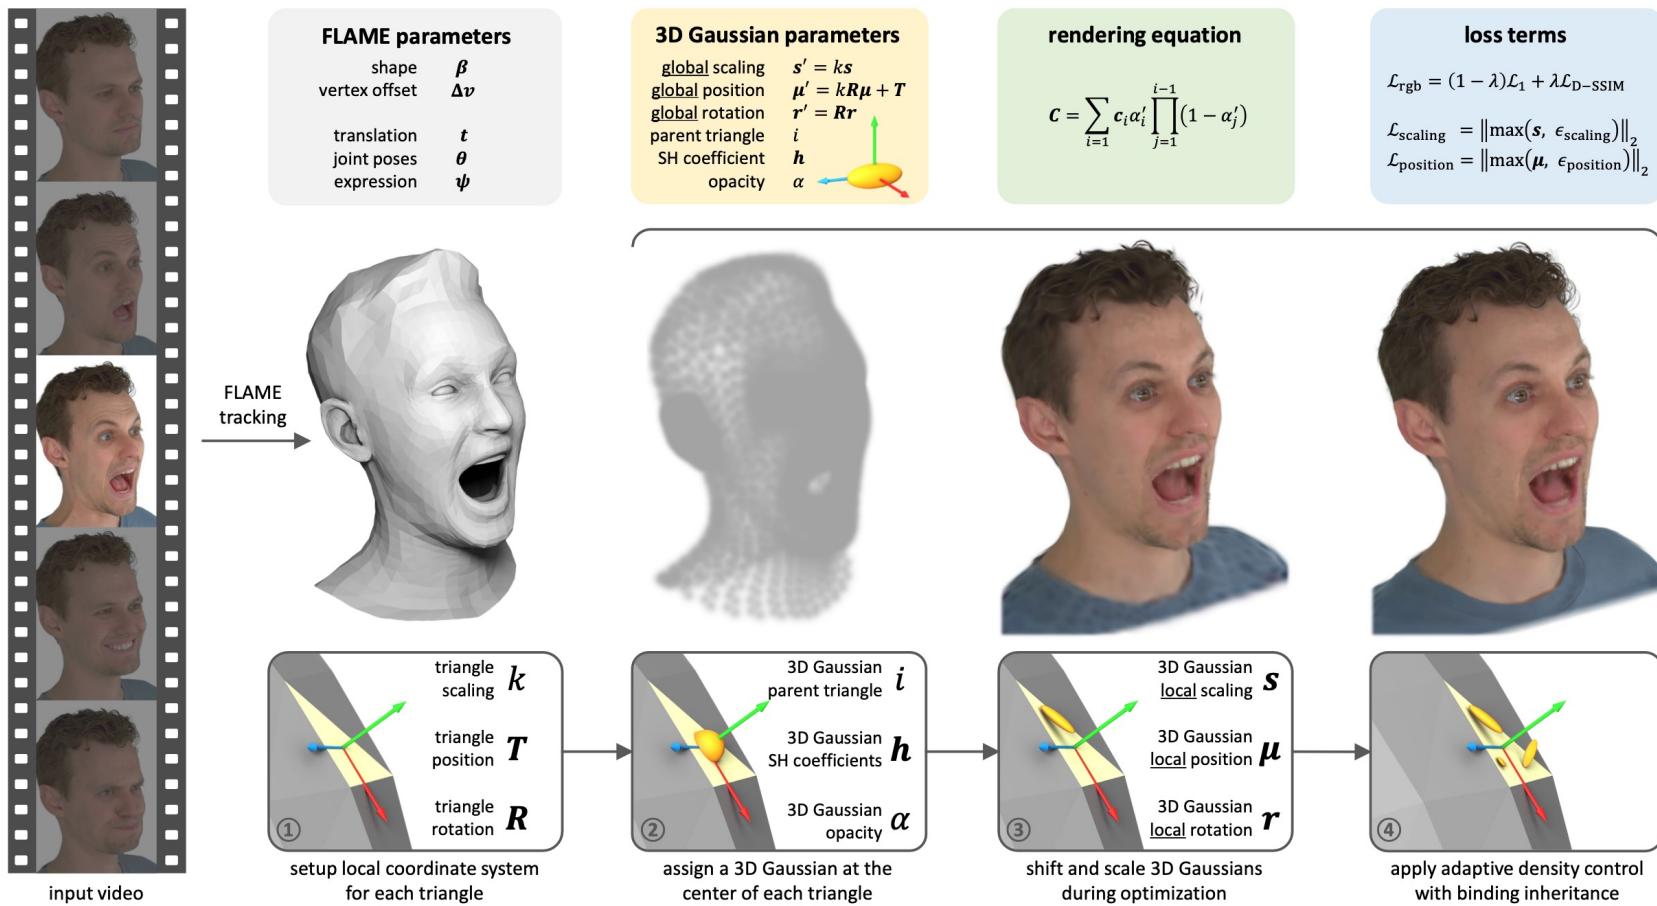


³Toyota Motor Europe NV/SA
associated partner by contracted services



- **Input:** multi-view human centric video
- **Output:** dynamic gaussian

Priors | Human | GaussianAvatars



- Flame parametric modeling
- Binding mesh with 3D gaussian

Animatable Gaussians:

Learning Pose-dependent Gaussian Maps for High-fidelity Human Avatar Modeling

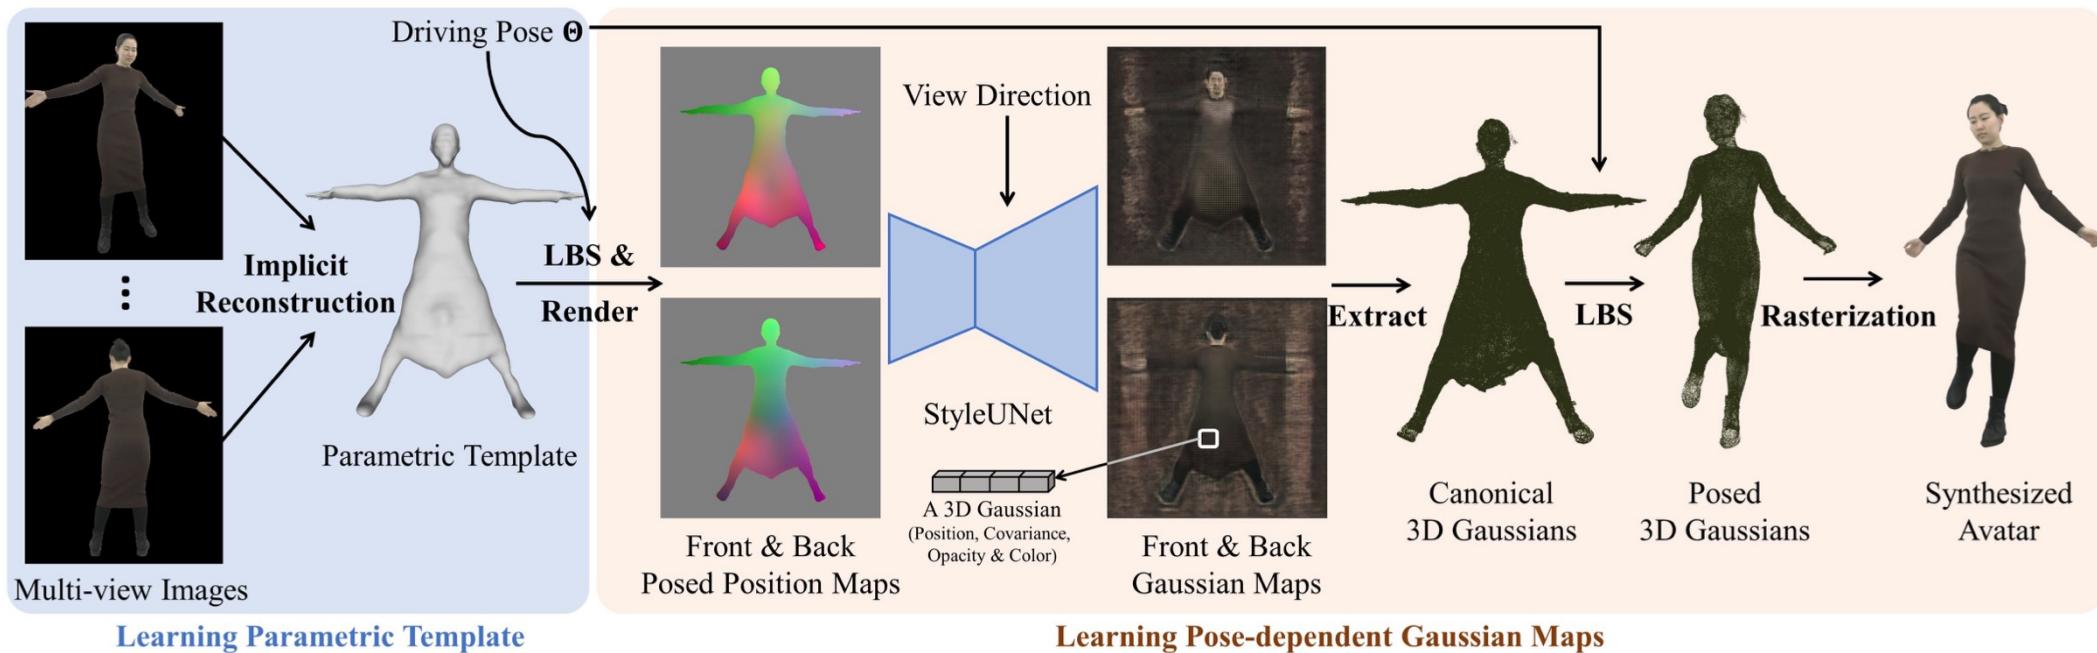
Zhe Li¹, Zerong Zheng², Lizhen Wang¹, Yebin Liu¹

¹Tsinghua University ²NNKosmos Technology



- **Input:** multi-view human centric video
- **Output:** dynamic gaussian

Priors | Human | Animatable Gaussians



- Replace MLPs with 2D CNNs
- LBS parametric modeling
- Parameterizes 3D Gaussians onto front & back Gaussian maps

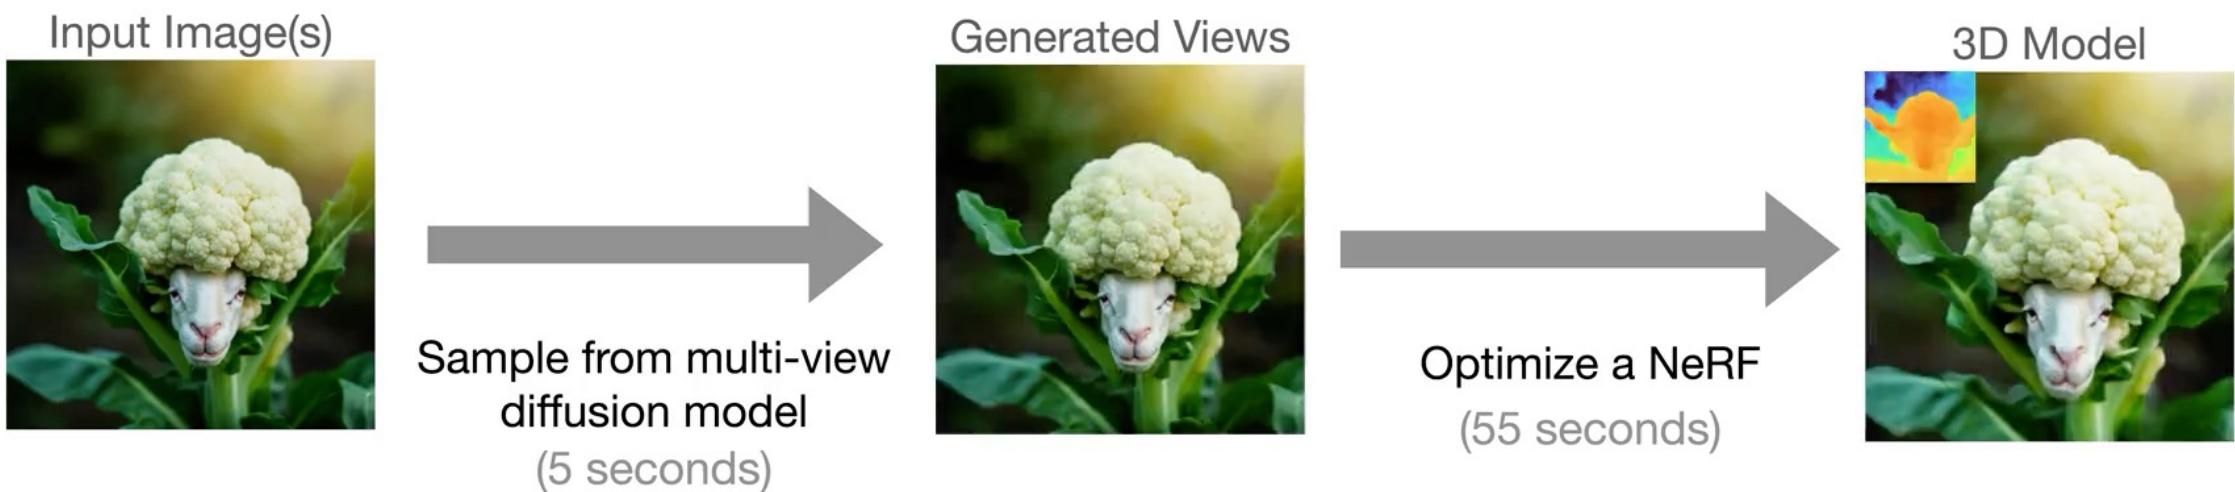
CAT3D: Create Anything in 3D with Multi-View Diffusion Models

Ruiqi Gao^{1*} Aleksander Holynski^{1*}

Philipp Henzler² Arthur Brussee¹ Ricardo Martin-Brualla²

Pratul P. Srinivasan¹ Jonathan T. Barron¹ Ben Poole^{1*}

¹Google DeepMind ²Google Research *equal contribution



ReconX: Reconstruct Any Scene from Sparse Views with Video Diffusion Model

Fangfu Liu^{1*}, Wenqiang Sun^{2*}, Hanyang Wang^{1*}, Yikai Wang¹, Haowen Sun¹,

Junliang Ye¹, Jun Zhang², Yueqi Duan¹

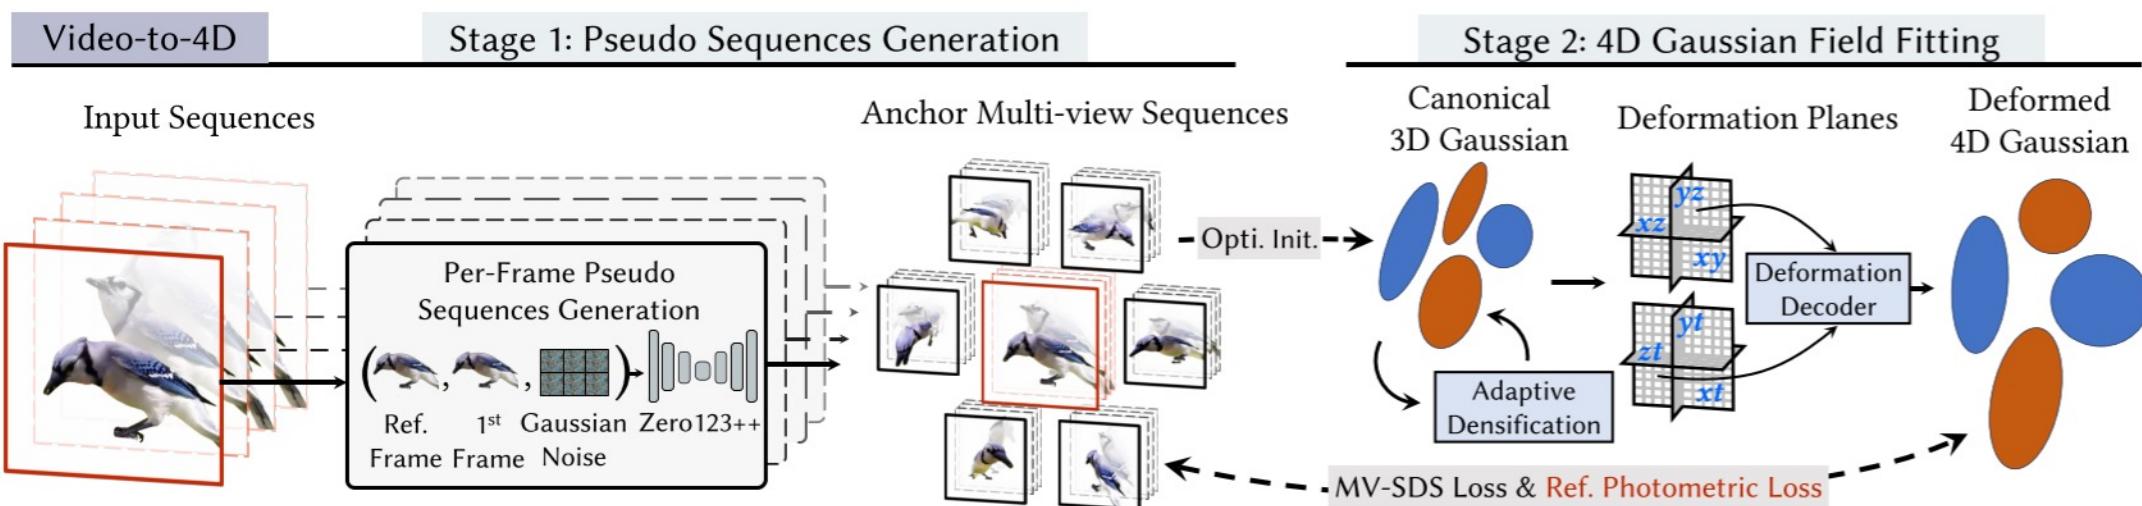
¹ Tsinghua University ² HKUST



STAG4D: Spatial-Temporal Anchored Generative 4D Gaussians

Yifei Zeng^{*1}, Yanqin Jiang^{*2}, Siyu Zhu³, Yuanxun Lu¹, Youtian Lin¹, Hao Zhu¹, Weiming Hu²,
Xun Cao¹, Yao Yao^{1✉}

¹Nanjing University, ²CASIA, ³Fudan University,



Priors | LVM | Video-to-4D Generation



SV4D: Dynamic 3D Content Generation with Multi-Frame and Multi-View Consistency



- 基于四维空间 (4D Representation) 的动态三维重建方法
- 基于形变场 (Deformation Field) 的动态三维重建方法
- 基于逐帧重建 (Per-frame Reconstruction) 的动态三维重建方法
- 基于模型先验 (Model Priors) 的动态三维重建方法

- More efficient 4D representations
- Large vision model for better 4D reconstruction
- Blurred boundaries between 4D reconstruction and generation



A differentiable point-based rendering library.

[Document](#) | [Paper \(Comming soon\)](#) | [DPTR Backend](#)

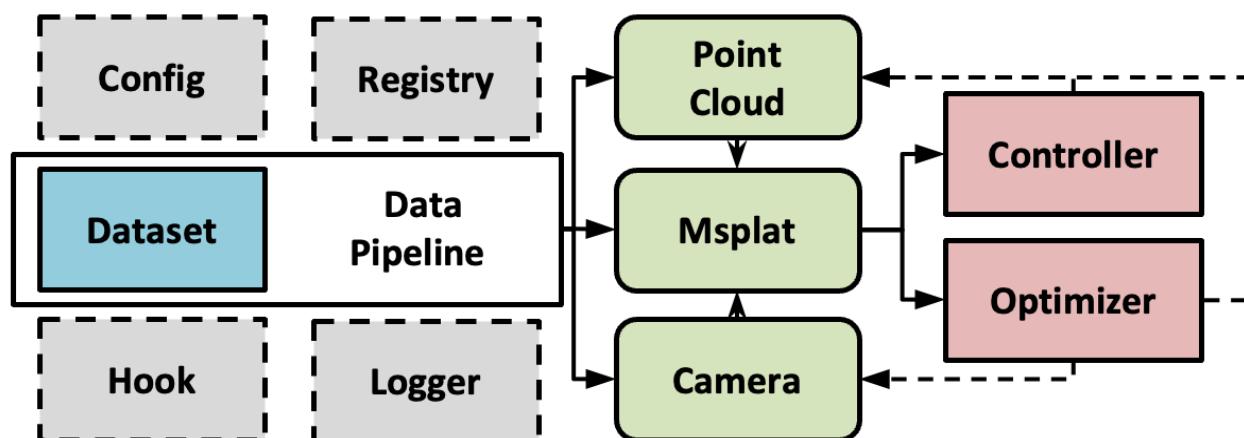
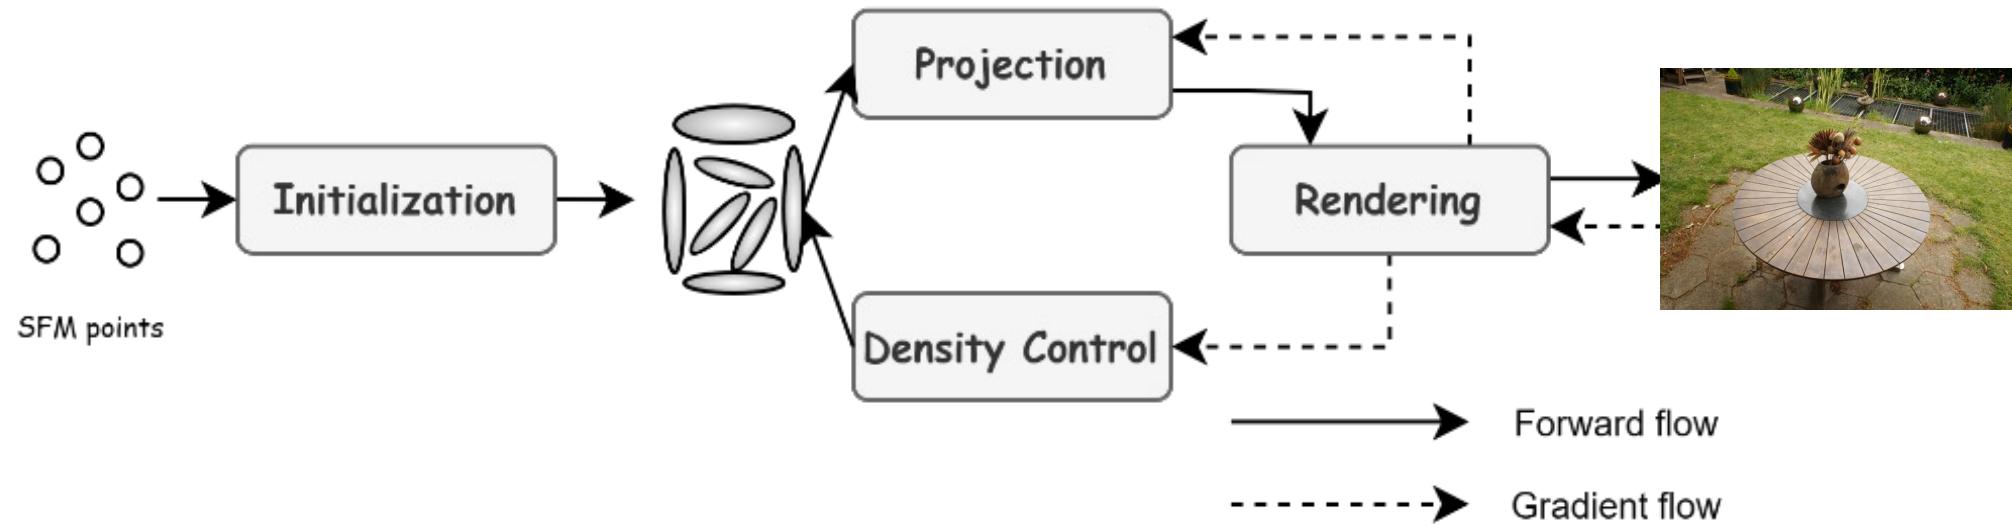
visitors 13 / 8380 stars repo not found Pointrix document

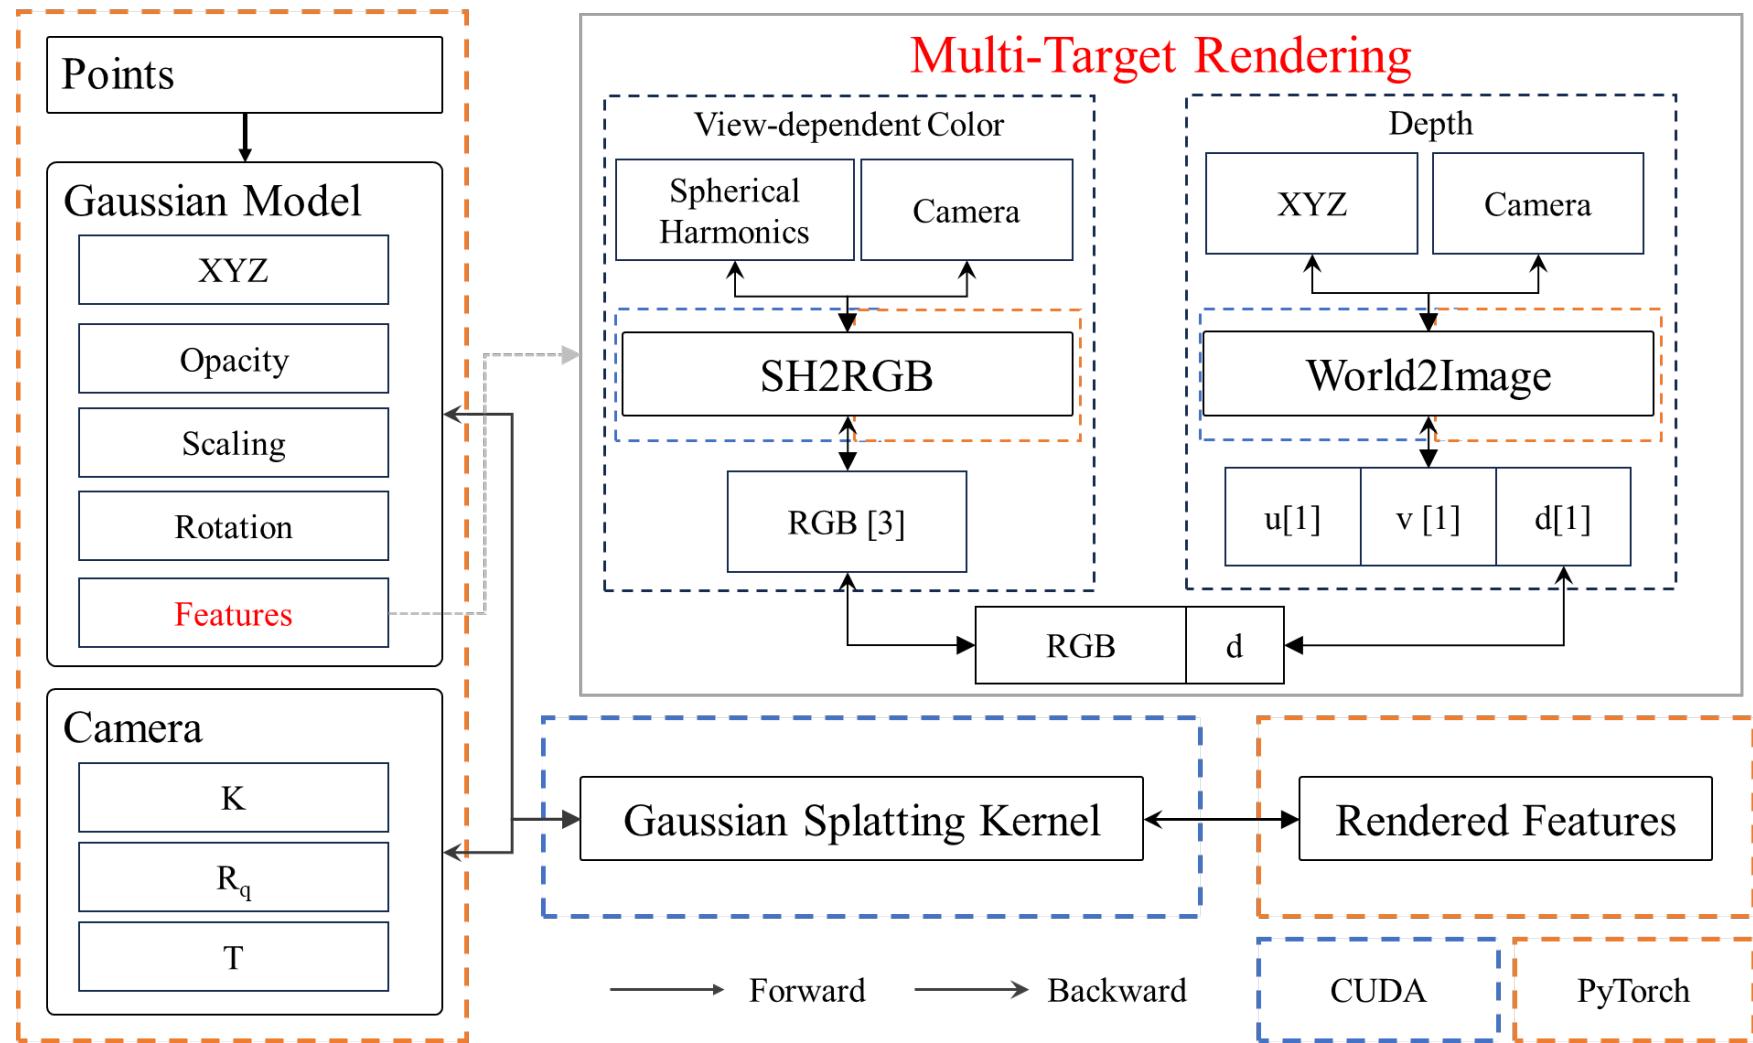
Pointrix is a differentiable point-based rendering library which has following properties:

- **Highly Extensible:**
 - Python API
 - Modular design for both researchers and beginners
 - Implementing your own method without touching CUDA
- **Powerful Backend:**
 - CUDA Backend
 - Forward Anything: rendering image, depth, normal, optical flow, etc.
 - Backward Anything: optimizing even intrinsics and extrinsics.
- **Rich Features:**
 - 3D Reconstruction: [Vanilla 3DGS](#) (Siggraph 2023), [Relightable 3DGS](#) (WIP)
 - 4D Reconstruction: [Deformable 3DGS](#) (CVPR 2024), [Gaussian-Flow](#) (CVPR 2024, WIP)
 - 3D Generation: [GSGen](#) (CVPR 2024, WIP), [DreamGaussian](#) (ICLR 2024, TBD), [LGM](#) (arXiv 2024, TBD)
 - 4D Generation: [STAG4D](#) (arXiv 2024, WIP), [DreamGaussian4D](#) (arXiv 2023, TBD)



Pointrix | 设计框架





```
1  @MODEL_REGISTRY.register()
2  class DeformGaussian(BaseModel):
3      def __init__(self, cfg, datapipeline, device="cuda"):
4          super().__init__(cfg, datapipeline, device)
5
6          # you can refer to projects/deformable_gaussian/model.py
7          # if you want to know the detail of DeformNetwork.
8          self.deform = DeformNetwork(is_blender=False).to(self.device)
9
10         # The gaussian point cloud is implemeted in BaseModel, we
11         # do not need to care about the detail here.
12
13     def forward(self, batch):
14         camera_fid = torch.Tensor([batch[0]['camera'].fid]).float().to(self.device)
15         position = self.point_cloud.get_position
16         time_input = camera_fid.unsqueeze(0).expand(position.shape[0], -1)
17         d_xyz, d_rotation, d_scaling = self.deform(position, time_input)
18
19         render_dict = {
20             "position": self.point_cloud.position + d_xyz,
21             "opacity": self.point_cloud.get_opacity,
22             "scaling": self.point_cloud.get_scaling + d_scaling,
23             "rotation": self.point_cloud.get_rotation + d_rotation,
24             "shs": self.point_cloud.get_shs,
25         }
26
27         return render_dict
28
```

Pointrix | Normal Prior 示例

Modify configuration to read Normal data automatically.

```
1 trainer:  
2   datapipeline:  
3     data_set: "ColmapDepthNormalDataset"  
4     shuffle: True  
5     batch_size: 1  
6     num_workers: 0  
7     dataset:  
8       data_path: "/home/linzhuo/gj/data/garden"  
9       cached_observed_data: ${trainer.training}  
10      scale: 0.25  
11      white_bg: False  
12      observed_data_dirs_dict: {"image": "images", "normal": "normals"}
```

new dataset
func name

variable name:
folder name

We highlight the modified part.

```
1 # Registry  
2 @DATA_SET_REGISTRY.register()  
3 class ColmapDepthNormalDataset(ColmapDataset):  
4   def _transform_observed_data(self, observed_data, split):  
5     cached_progress = ProgressLogger(description='transforming cached observed_data', suffix=  
6     cached_progress.add_task(f'Transforming', f'Transforming {split} cached observed_data', 1)  
7     with cached_progress.progress as progress:  
8       for i in range(len(observed_data)):  
9         # Transform Image  
10        image = observed_data[i]['image']  
11        w, h = image.size  
12        image = image.resize((int(w * self.scale), int(h * self.scale)))  
13        image = np.array(image) / 255.  
14        if image.shape[2] == 4:  
15          image = image[:, :, :3] * image[:, :, 3:4] + self.bg * (1 - image[:, :, 3:4])  
16        observed_data[i]['image'] = torch.from_numpy(np.array(image)).permute(2, 0, 1).float()  
17        cached_progress.update(f'Transforming', step=1)  
18  
19        # Transform Normal  
20        observed_data[i]['normal'] = \  
21          (torch.from_numpy(np.array(observed_data[i]['normal'])) \  
22            / 255.0).float().permute(2, 0, 1)  
23  
return observed_data
```

process normal
data

Pointrix | Normal Prior 示例

```
1 @MODEL_REGISTRY.register()
2 class NormalModel(BaseModel):
3     def forward(self, batch=None, training=True, render=True, iteration=None) -> dict:
4
5         if iteration is not None:
6             self.renderer.update_sh_degree(iteration)
7         frame_idx_list = [batch[i]["frame_idx"] for i in range(len(batch))]
8         extrinsic_matrix = self.training_camera_model.extrinsic_matrices(frame_idx_list) \
9             if training else self.validation_camera_model.extrinsic_matrices(frame_idx_list)
10        intrinsic_params = self.training_camera_model.intrinsic_params(frame_idx_list) \
11            if training else self.validation_camera_model.intrinsic_params(frame_idx_list)
12        camera_center = self.training_camera_model.camera_centers(frame_idx_list) \
13            if training else self.validation_camera_model.camera_centers(frame_idx_list)
14
15        point_normal = self.get_normals
16        projected_normal = self.process_normals(
17            point_normal, camera_center, extrinsic_matrix)
18
19        render_dict = {
20            "extrinsic_matrix": extrinsic_matrix,
21            "intrinsic_params": intrinsic_params,
22            "camera_center": camera_center,
23            "position": self.point_cloud.get_position(),
24            "opacity": self.point_cloud.get_opacity(),
25            "scaling": self.point_cloud.get_scaling(),
26            "rotation": self.point_cloud.get_rotation(),
27            "shs": self.point_cloud.get_shs(),
28            "normals": projected_normal
29        }
30        if render:
31            render_results = self.renderer.render_batch(render_dict, batch)
32            return render_results
33        return render_dict
34
35    def get_loss_dict(self, render_results, batch) -> dict:
36        loss = 0.0
37        gt_images = torch.stack(
38            [batch[i]["image"] for i in range(len(batch))],
39            dim=0
40        )
41        normal_images = torch.stack(
42            [batch[i]["normal"] for i in range(len(batch))],
43            dim=0
44        )
45        l1_loss = l1_loss(render_results['rgb'], gt_images)
46        ssim_loss = 1.0 - ssim(render_results['rgb'], gt_images)
47        loss += (1.0 - self.cfg.lambda_ssim) * l1_loss
48        loss += self.cfg.lambda_ssim * ssim_loss
49        # normal 监督的损失
50        normal_loss = 0.1 * l1_loss(render_results['normal'], normal_images)
51        loss += normal_loss
52
53        loss_dict = {"loss": loss,
54                    "l1_loss": l1_loss,
55                    "ssim_loss": ssim_loss,
56                    "normal_loss": normal_loss}
57
58        return loss_dict
```

```
1 @ RENDERER_REGISTRY.register()
2 class MsplatNormalRender(MsplatRender):
3     def render_iter(self, height, width, extrinsic_matrix, intrinsic_params, camera_center, posit
4                     scaling, rotation, shs, normals, **kwargs) -> dict:
5         direction = (position -
6                         camera_center.repeat(position.shape[0], 1))
7         direction = direction / direction.norm(dim=1, keepdim=True)
8         rgb = msplat.compute_sh(shs.permute(0, 2, 1), direction)
9         extrinsic_matrix = extrinsic_matrix[:, :, :]
10
11     ...
12
13     - get normal of points
14
15     ...
16
17
18     32     # sort
19     (gaussian_ids_sorted, tile_range) = msplat.sort_gaussian(
20         uv, depth, width, height, radius, tiles_touched
21     )
22
23     37     Render_Features = RenderFeatures(rgb=rgb, depth=depth, normal=normals) →
24     render_features = Render_Features.combine()
25
26
27     40     ndc = torch.zeros_like(uv, requires_grad=True)
28     try:
29         ndc.retain_grad()
30     except:
31         raise ValueError("ndc does not have grad")
32
33
34     46     # alpha blending
35     rendered_features = msplat.alpha_blending(
36         uv, conic, opacity, render_features,
37         gaussian_ids_sorted, tile_range, self.bg_color, width, height, ndc
38     )
39     rendered_features_split = Render_Features.split(rendered_features)
40
41
42     53     normals = rendered_features_split["normal"]
43
44
45     55     # convert normals from [-1,1] to [0,1]
46     normals_im = normals / normals.norm(dim=0, keepdim=True)
47     normals_im = (normals_im + 1) / 2
48
49
50     59     rendered_features_split["normal"] = normals_im
51
52
53
54
55
56
57
58
59
60
61     return {"rendered_features_split": rendered_features_split,
62             "uv_points": ndc,
63             "visibility": radius > 0,
64             "radii": radius
65         }
```

Add new feature for
alpha blending



get normal of points



Add new feature for
alpha blending



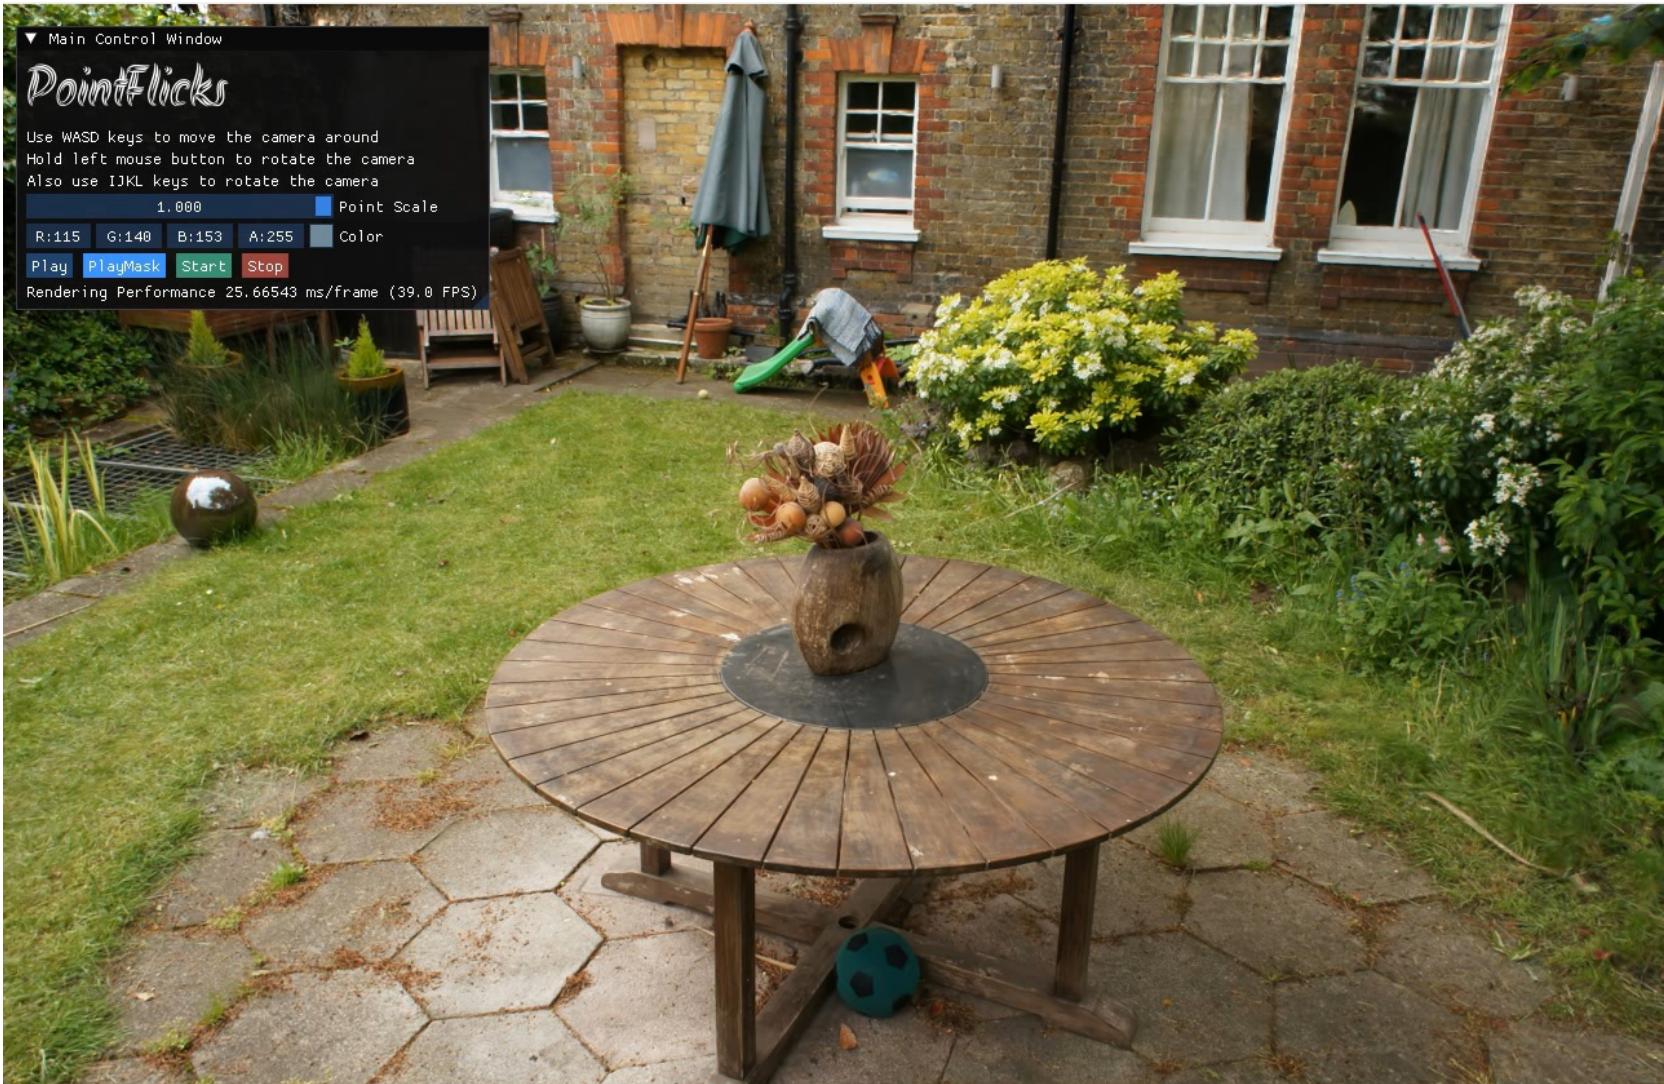
add normal loss



Pointrix | 运行界面



Pointrix | Interactive GUI





Search

⌘ + K

Get started (In 10 minutes)

Installation

Run Your First Model

Using Pointflicks (GUI)

Configuration Files in Pointrix

Tutorial

Adding Supervision for 3DGS

Framework

Overview

Data Pipeline

Model

Trainer

Hook

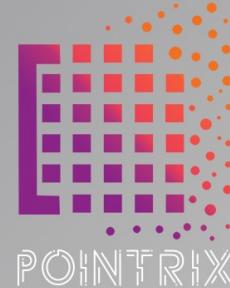
Advanced

Camera Optimization



A differentiable point-based rendering framework

Get Started



Overview

[Pointrix](#) is a **light-weight differentiable point-based rendering framework** which has following properties:

Highly Extensible

Pointrix adopts a modular design, with clear structure and easy extensibility.

Rich Feature

Pointrix supports the implementation of various types of tasks.

Powerful backend

Msplat which offer foundational functionalities for point rendering serves as the backend of Pointrix.

[Learn more »](#)

[Learn more »](#)

[Learn more »](#)



 Search ⌘ + K

开始 (10分钟内学习Pointrix)

安装

运行你的第一个模型

Pointrix 中的配置文件

使用 Pointflicks (GUI) 可视化

教程

为点云渲染添加监督

架构

总览

数据流水线

模型

训练器 (Trainer)

钩子函数

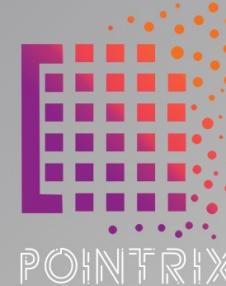
Advanced

相机优化



一个可微分点云渲染框架

Get Started



Overview

[Pointrix](#) 是一个 轻量可微分点云渲染框架 并具有以下几个特点:

高度可扩展性

Pointrix 采用模块化设计，具有清晰的结构和强大的扩展性。

丰富的特性

Pointrix 支持各种不同类型的工作实现，包括静态动态场景重建，生成以及PBR渲染。

强大的后端

MSplat 作为Pointrix的后端，支持各类点云渲染功能：包括各类特征的渲染，超高阶的sh阶数，以及所有输入的梯度反传。



Thanks!

Yao Yao

Nanjing University



Linzhuo Chen



Youtian Lin